

Consolidated Report: coreweave

Generated 2026-03-29 04:09 UTC from 3 analysis profile(s): company_profile, marketing_analysis, operational_metrics.

Note: The primary source of information in this report is the company's own website. The content reflects the company's self-reported claims and marketing materials and should not be treated as independently verified fact.

Executive Summary

Executive Summary: CoreWeave

Company Overview

CoreWeave is a purpose-built AI cloud infrastructure company founded in 2017 and headquartered in Livingston, New Jersey. Originally started by three co-founders operating GPUs out of a small setup, the company has scaled rapidly to become a Nasdaq-listed entity (CRWV, March 2025) with over 2,000 employees, 43 data centers, 250,000+ GPUs, and more than 3.1 GW of contracted power capacity. Branding itself as "The AI Hyperscaler™" and "The Essential Cloud for AI™," CoreWeave generated \$5 billion in revenue in fiscal year 2025—168% year-over-year growth—making it the fastest cloud platform in history to reach that milestone, with a \$66.8 billion revenue backlog.

Key Findings

Corporate Profile & Leadership. CoreWeave's leadership team combines founder-driven vision with deep hyperscaler expertise. CEO Michael Intrator and co-founders Brian Venturo, Brannin McBee, and Peter Salanki are complemented by senior hires from Google Cloud, AWS, Oracle, and IBM. The company has executed an aggressive acquisition strategy—Weights & Biases (~\$1.7B), Core Scientific (~\$9B), OpenPipe, Monolith, and Marimo—to build a vertically integrated "metal to model" platform spanning infrastructure, developer tools, MLOps, and domain-specific AI applications.

Market Positioning & Brand Strategy. CoreWeave positions itself explicitly against legacy hyperscalers (AWS, Azure, GCP) and emerging neoclouds, claiming up to 20% higher GPU cluster performance, 96% goodput versus a 90% industry average, and 10x faster inference spin-up times. Independent validation underpins these claims: CoreWeave is the sole recipient of SemiAnalysis's Platinum ClusterMAX™ rating (earned twice), holds MLPerf training records, and carries NVIDIA Exemplar Cloud certification. Its marquee customer roster—including OpenAI (contracts totaling ~\$22.4B), IBM, Mistral AI, Cohere, Jane Street, Perplexity, and Cursor—provides powerful peer proof. The company's pricing philosophy emphasizes transparency (zero egress and transfer fees) and performance-adjusted cost rather than sticker price.

Infrastructure & Operational Scale. CoreWeave operates across 43 data centers in 15+ U.S. states and multiple European locations (UK, Sweden, Scotland), with approximately 590

MW active and a clear path beyond 850 MW by year-end 2025. The Core Scientific acquisition adds ~1.3 GW of gross power with over 1 GW of incremental expansion potential. The company consistently achieves first-to-market deployment of NVIDIA's latest GPUs—GB200 NVL72, GB300 NVL72, HGX B300—reinforced by a \$2 billion strategic investment from NVIDIA. Storage, networking (InfiniBand at 3,200 Gbps per node), and rack densities exceeding 130 kW per rack are all designed ground-up for AI workloads.

Cross-Cutting Themes, Strengths & Risks

Strengths. CoreWeave's deepest competitive advantage is its symbiotic NVIDIA relationship, which creates a self-reinforcing cycle of preferential hardware access, performance validation, and customer acquisition. Vertical integration from physical facilities through developer tooling raises switching costs and platform stickiness. Third-party benchmarks provide credibility that pure marketing cannot replicate.

Risks. Concentration risk is significant: the OpenAI relationship alone represents an outsized share of backlog, and dependence on NVIDIA for hardware supply creates single-vendor exposure. Capital intensity is extraordinary—over \$12 billion raised, with multi-billion-dollar data center and acquisition commitments funded heavily by debt. Any slowdown in AI investment, GPU supply disruptions, or competitive pricing pressure from hyperscalers could strain this capital structure. Geographic and regulatory risk is also rising as CoreWeave pursues federal contracts (FedRAMP) and sovereign AI deployments across multiple jurisdictions.

Strategic Outlook

CoreWeave is executing an ambitious playbook to become the default infrastructure layer for the AI era—analogue to AWS's role in web computing. Its 2030 target of 5+ GW of AI factory capacity, planned deployment of NVIDIA's Vera Rubin platform in late 2026, and expansion into agentic AI and inference-optimized workloads position it at the center of the next wave of demand. Success will depend on sustaining the NVIDIA partnership, diversifying its customer base beyond a handful of large AI labs, managing leverage prudently, and converting its \$66.8 billion backlog into durable, profitable revenue as the AI infrastructure market matures and competition intensifies.

Company Profile

Source: coreweave.md

Marketing – Factual & Quantitative: <https://coreweave.com>

Extract concrete, verifiable facts from a company's website: corporate details, infrastructure specs, customer names, product listings, and quantitative claims.

Based on 328 stored pages analysed in 1 batch(es) with 1 AI calls (namespace: webscrape, model: claude-opus-4-6, profile: company_profile).

Headquarters

CoreWeave's corporate address is **290 W Mt Pleasant Ave, Suite 4100, Livingston, NJ 07039**. This address appears consistently across the website footer. The company also has a **European headquarters in London, UK** (announced in 2024), and a **Washington, D.C. office at 1001 Pennsylvania Avenue NW**. CEO Michael Intrator previously referenced "Roseland, N.J." in earlier press releases, though the current address is listed as Livingston, NJ.

Years In Business

CoreWeave was **founded/established in 2017**. Multiple references confirm this:

- "Founded in 2017, CoreWeave offers industry-leading, scalable, on-demand computing resources"
 - "Established in 2017, CoreWeave completed its public listing on Nasdaq (CRWV) in March 2025"
 - "In 2017, we bought and plugged in our first GPU"
 - CEO Mike Intrator: "Since 2017, CoreWeave has become the engine that powers the creation and delivery of the intelligence that drives innovation"
-

Company Executive Names And Titles

Co-founders:

- **Mike Intrator** – Chief Executive Officer, Co-founder, Chairman of the Board. Previously CEO of Hudson Ridge Asset Management LLC (2013–2018). B.A. Political Science from Binghamton University; M.P.A. from Columbia University.
- **Brian Ventura** – Chief Strategy Officer, Co-founder, Board member. Previously CTO (2017–2024). Previously Partner at Hudson Ridge Asset Management. B.A. Economics from Haverford College.
- **Brannin McBee** – Chief Development Officer, Co-founder. Previously Chief Strategy Officer (2017–2024). B.S. Finance from University of Colorado Boulder.
- **Peter Salanki** – Chief Technology Officer, Co-founder. Previously VP of Engineering (2022–2024) and Director of Engineering (2019–2022). Prior role: Director, Americas, Office of the CTO at Sandvine.

Other C-Suite / Senior Leaders:

- **Nitin Agrawal** – Chief Financial Officer (since March 2024). Previously VP Finance at Google Cloud; CFO at Mapbox; Finance Director at AWS. MBA from Duke University Fuqua School.
- **Chen Goldberg** – Executive Vice President, Product & Engineering (since August 2024). Previously GM & VP of Engineering at Google Cloud (Kubernetes and Serverless). B.A. and MBA from The Open University of Israel.
- **Sachin Jain** – Chief Operating Officer (since August 2024). Previously SVP at Oracle Cloud (AI infrastructure); prior roles at Google Cloud and Amazon.
- **Michelle O'Rourke** – Chief People Officer (since October 2024). Joined CoreWeave in 2021. B.A. Human Services from Villanova University.

- **Kristen McVeety** – General Counsel (since March 2022), Corporate Secretary (since December 2024). Previously founder/partner at Gabler & McVeety LLP. J.D. from Brooklyn Law School.
- **Max Hjelm** – Senior Vice President, Revenue (since January 2024). Joined CoreWeave in July 2020. B.S. Sociology from Haverford College.
- **Jean English** – Chief Marketing Officer. Previously CMO at Juniper Networks, Palo Alto Networks, NetApp, and IBM Asia Pacific.
- **Sandy Venugopal** – Chief Information Officer. Previously CIO at SentinelOne and Uber. Board member at The Adecco Group.
- **James Higgins** – Chief Information Security Officer. Previously Director of Product Security at Google; CISO at Block, Inc. (Square) and Snap, Inc.
- **Jeff Baker** – Chief Accounting Officer (since August 2024). Previously CAO at Rivian (instrumental in its IPO). B.S. Accountancy from University of Florida.
- **Jon Jones** – Chief Revenue Officer (since September 2025). Previously Global VP of Startups & VC and VP of Go-to-Market at AWS. MBA from Wharton.
- **Chetan Kapoor** – Chief Product Officer (since June 2024). Previously part of the Compute leadership team at AWS managing Hardware Accelerated Compute.
- **Karen Boone** – Independent Board Director and Audit Committee Chair (since January 2025). Previously Interim Co-CEO of Peloton; former President and CFO of Restoration Hardware.

Customer Names, Locations, Industry/Type Of Entity, And Other Relevant Details

Named customers with details:

- **OpenAI** – AI lab. Contract value up to ~\$22.4 billion total (initial \$11.9B, expanded \$4B, then \$6.5B). OpenAI invested \$350M in CoreWeave. CEO Sam Altman quoted: "CoreWeave has helped us build really large-scale computing clusters." Uses CoreWeave for training and serving advanced models.
- **IBM** – Enterprise technology. Uses NVIDIA GB200 NVL72 clusters for training Granite™ models. Achieved up to 80% higher performance. Built a supercomputer on CoreWeave twice the size of Blue Vela. IBM's Danny Barnett (VP Emerging Technology Engineering) and Hillary Hunter (CTO/GM Innovation, IBM Infrastructure) quoted. IBM Research collaborated on MLPerf Training v5.0.
- **Mistral AI** – Paris-based AI lab. Customer since 2023 (first contract for H100s). Achieved 2.5x faster training speeds on NVIDIA GB200 NVL72. Cut training time in half for open-source reasoning models. Founded by Arthur Mensch (CEO), Timothée Lacroix (CTO), and Guillaume Lample.
- **Cohere** – AI company. Deployed one of industry's first NVIDIA GB200 NVL72 clusters, unlocking 3× faster training. Powers North agentic AI platform. Autumn Moulder (VP of Engineering) quoted.
- **Jane Street** – Global quantitative trading firm. Customer and investor. Uses CoreWeave for ML and research workloads. Craig Falls (Head of Quant Research) quoted. Adam Canady (Research Contributor) appeared at GTC 2025 session.
- **Perplexity** – AI search company. Multi-year strategic partnership for inference workloads on dedicated NVIDIA GB200 NVL72 clusters. Dmitry Shevelenko (CBO) quoted.
- **Runway** – AI video generation company. Signed contract for next-generation AI video models. Will use NVIDIA GB300 NVL72 for training and inference, plus W&B Models and CoreWeave AI Object Storage. Anastasis Germanidis (Co-Founder and CTO) quoted.

- **Cursor** – AI-powered code editor. Already running production workloads on NVIDIA HGX B200, moving toward B300. Aman Sanger (Co-Founder) quoted.
- **Cognition** – AI company. Listed among diversified customer base.
- **Midjourney** – AI image generation. Listed among diversified customer base.
- **Mercado Libre** – E-commerce (Latin America). Listed among diversified customer base.
- **Aston Martin Aramco Formula One™ Team** – F1 racing team. Official AI Cloud Computing Partner. Multi-year partnership. CoreWeave supports team's first large-scale cloud computing facility; wind tunnel named "CoreWeave Wind Tunnel." Lawrence Stroll (Executive Chairman) quoted.
- **Poolside** – Foundation model company. Partnership for 40,000+ NVIDIA GPUs (GB300 NVL72). CoreWeave planned as anchor tenant for Poolside's Project Horizon 2GW AI campus in West Texas. Eiso Kant (Co-Founder/Co-CEO) quoted.
- **Cline** – Open and Secure Coding Agent. Integrates CoreWeave's W&B Inference. 5 million+ installations. Saoud Rizwan (Founder) quoted.
- **Trillion Labs** – AI startup creating foundational models for Asia/Korea. Uses H100 clusters. Jay Shin (CEO and Co-Founder) quoted.
- **General Intuition** – AI company (sibling of Medal.tv). Achieved 20–30% better performance in CoreWeave ARENA for world model training. Florian Laurent (Infra Lead) quoted.
- **NovelAI / Anlatan** – AI storytelling subscription service. Early H100 customer. 400K unique users/month, 40K registered users. Eren Doğan (Anlatan CEO) quoted.
- **Bit192, Inc.** – Japanese company working on open-source AI/NLP. CEO Yasu Seno. 400K unique users/month.
- **Decart / Decart AI** – Generative interactive video. Uses NVIDIA HGX B200 at production scale. Orian Leitersdorf (Chief Scientist) quoted. Dean Leitersdorf (CEO) quoted.
- **Moonvalley** – Generative filmmaking company. Naeem Talukdar (Co-Founder and CEO) quoted. CoreWeave Ventures portfolio company.
- **Tarteel AI** – AI-powered Quran app, 3M+ downloads. Migrated from AWS to CoreWeave via Zeet. 22% improvement in latency and ~56% cost reduction. Anas Abou Allaban (Co-Founder and CEO) quoted.
- **Redesign Science** – Computational chemistry/drug discovery. 84% cost savings on cloud spend. Haotian Li (CTO) quoted.
- **Product Insight** – Product design/engineering firm (Boston area). VFX rendering project for Mevion Medical Systems. Anthony Catacchio (VP of Engineering) quoted.
- **OthersideAI / HyperWrite** – AI writing assistant. Matt Brockman quoted.
- **Inflection AI** – AI lab. Used CoreWeave's H100 infrastructure (MLPerf submission).
- **Neuralift AI** – Marketing technology. Customer story featured.
- **inference.net** – AI model deployment. Ibrahim Ahmed (CTO) quoted regarding Spot instances.
- **Periodic Labs** – AI research. Xander Dunn (Member of Technical Staff) quoted.
- **Hologen** – AI company. Chris Goddard (Co-Founder & CTO) quoted.
- **Numerata** – AI-powered software development tools. Seed investment by Jane Street and CoreWeave Ventures. Kevin Wang (Founder) quoted.
- **Replicate** – ML model hosting. Morgan Fainberg (Principal Engineer) quoted.
- **SquadStack.ai** – AI-powered contact center platform.
- **QA Wolf** – Hybrid testing platform. Jon Perl (CEO) quoted.
- **LG AI Research and Riskfuel** – Use Weights & Biases tools.
- **GSK** – Pharmaceutical. Customer story: developed and managed production models with W&B.

Industries served: AI labs, quantitative trading/financial services, enterprise technology, automotive/F1, media & entertainment, healthcare/life sciences, e-commerce, gaming,

Partner Names, Locations, Industry/Type Of Entity, And Other Relevant Details

- **NVIDIA** – Primary technology partner. NVIDIA invested \$2 billion in CoreWeave (January 2026). CoreWeave is NVIDIA's first Elite Cloud Services Provider for Compute. Collaboration on GB200, GB300, HGX B300, Rubin platform, BlueField DPUs, Quantum InfiniBand networking, NIM microservices, AI Enterprise, NVCF, NVIDIA Exemplar Cloud validation. Joint MLPerf submissions. Speakers: Jensen Huang (CEO), Ian Buck (VP Hyperscaler and HPC Computing), Dion Harris (Sr. Director AI & HPC Infrastructure).
- **Weights & Biases** – Acquired by CoreWeave (completed May 2025) for ~\$1.7B. AI developer platform for MLOps/LLMOps. Founded by Lukas Biewald, Chris Van Pelt, and Shawn Lewis. Used by 1M+ AI engineers and 1,500+ organizations.
- **OpenPipe** – Acquired (September 2025). RL/fine-tuning platform. Founded by Kyle Corbitt (CEO). Agent Reinforcement Trainer (ART) open-source toolkit.
- **Monolith AI** – Acquired (October 2025). London-based, AI/ML for physics and engineering. Founded by Dr. Richard Ahlfeld. Customers include Nissan, BMW, Honeywell.
- **Marimo** – Acquired (October 2025). Open-source AI-native Python notebook. Co-founder Akshay Agrawal (CEO).
- **Core Scientific** – Acquisition announced (July 2025). Data center infrastructure provider. Austin, TX. ~1.3 GW of gross power. Adam Sullivan (President and CEO).
- **Dell Technologies** – Extended relationship for AI infrastructure at scale.
- **Intel** – CPU partner. 4th Gen Xeon Scalable (Sapphire Rapids), 5th Gen (Emerald Rapids) used in CoreWeave clusters.
- **Red Hat** – CoreWeave is founding member of Red Hat's llm-d open-source project (alongside IBM Research, Google, NVIDIA).
- **Run:ai** – GPU orchestration partner. Joint deployments including hundreds of H100s for life sciences AI lab. Yael Dor (VP Sales & Partnerships) quoted.
- **SkyPilot / Anyscale** – Multi-cloud AI orchestration. SkyPilot support officially added. Anyscale BYOC on CKS announced.
- **SchedMD** – Developer of Slurm. Collaboration on SUNK (Slurm on Kubernetes).
- **Loft Labs** – vCluster partner for virtual Kubernetes clusters.
- **Authentik Security** – Open-source identity provider. Used for hackathon provisioning.
- **UbiOps** – AI serving and orchestration platform (Amsterdam-based).
- **Zeet** – Integrated DevOps platform/deployment partner.
- **PureWeb** – Real-time 3D streaming platform.
- **TensorWorks** – Scalable Pixel Streaming for Unreal Engine.
- **Conductor** – Cloud-based rendering service.
- **Determined.AI** – AI training platform (integration).
- **VAST Data** – Storage partner (mentioned at SC23).
- **Equinix and Megaport** – Network POP/interconnect partners.
- **Switch** – Data center partner (Las Vegas Tier 5® facility; hosted GB200/GB300 deployments).
- **Digital Realty** – Data center partner (Crawley, UK).
- **Global Switch** – Data center partner (London Docklands, UK).
- **EcoDataCenter** – Data center partner in Sweden for NVIDIA Blackwell clusters.
- **DataVita** – Data center partner in Scotland for UK sovereign AI deployment.
- **Vertiv** – Data center infrastructure (cooling).

- **CrowdStrike** – Cybersecurity partner. "Falcon platform" to secure CoreWeave Cloud.
 - **Pinecone** – Managed vector database partner for RAG/inference.
 - **Milvus** and **Dragonfly** – Open-source vector databases supported on CKS.
 - **Meta** (Stanford Scaling Intelligence Lab) – Collaboration on torchforge RL framework at 512 H100 GPU scale.
 - **Cisco** – Investor in CoreWeave (\$650M secondary share sale, November 2024).
 - **Blackstone** – Led \$2.3B debt facility (August 2023) and \$7.5B debt facility (May 2024).
 - **Magnetar Capital** – Led Series B and co-led debt facilities.
 - **Coatue** – Led \$1.1B Series C (May 2024).
 - **Epic Games** – Partnership mentioned for Unreal Engine Pixel Streaming.
-

Products

Foundational Infrastructure:

- **GPU Compute** – Access to NVIDIA GPUs (Blackwell: GB300 NVL72, HGX B300, HGX B200, GB200 NVL72; Hopper: HGX H100, HGX H200, GH200; Ada Lovelace: L40, L40S, RTX PRO 6000 Blackwell Server Edition; older: A100, A40, A6000, etc.) on bare metal via Kubernetes.
 - **CPU Compute** – AMD EPYC Genoa and Intel Emerald Rapids processors.
 - **Bare Metal Servers** – Kubernetes directly on bare metal, no hypervisor layer. NVIDIA BlueField DPUs offload networking/storage.
 - **Networking Services** – NVIDIA Quantum InfiniBand (up to 3200Gbps per node, non-blocking); Virtual Private Cloud (VPC) powered by
-

Solutions

No relevant information found across any batch.

Services

No relevant information found across any batch.

Marketing Analysis

Source: coreweave.md

Marketing – Strategic & Interpretive: <https://coreweave.com>

Analyse a company's website for strategic positioning, brand identity, value propositions, and competitive strategy. These topics require interpreting tone, patterns, and implied intent rather than extracting explicit facts.

Positioning Statements

CoreWeave Positioning Statements — Consolidated Summary

Primary Taglines

- **"The Essential Cloud for AI™"** — used consistently across virtually every page and press release
- **"The AI Hyperscaler™"** — used extensively in press releases and announcements; also expressed as "the first true AI Hyperscaler"
- **"Ready for Anything, Ready for AI"** — the company's first integrated brand campaign (February 2026), featuring Chance the Rapper
- **"Built for pioneers, by pioneers"** — recurring brand tagline throughout the site

Hero & Brand Copy

- **"The force multiplier for AI. Trusted by the world's leading AI pioneers."** — hero copy; "force multiplier" repeated across multiple pages
- **"The world's #1 AI cloud platform, purpose-built for what's next"**
- **"CoreWeave Cloud is an AI-native platform purpose-built for AI"**
- **"pace, performance, and partnership"** — brand value framing

Category Positioning

- Explicitly positioned **against "general-purpose clouds," "legacy hyperscalers" (AWS, Azure, GCP), and "neoclouds"**
- **"The first cloud built from the ground up for AI"**
- **"We didn't start as a general-purpose cloud — we started by solving the hard problems in AI"**
- **"CoreWeave is a world-class new generation AI-Native cloud"** — Jensen Huang, NVIDIA CEO
- **"No other AI cloud provides this level of depth from metal to model"** — Peter Salanki, CTO
- Implicit positioning: the **purpose-built AI infrastructure alternative** to AWS, Azure, and GCP for GPU-intensive workloads

Competitive Differentiators

Ratings & Benchmarks

- **"The only AI cloud to earn a Platinum rating twice"** (SemiAnalysis ClusterMAX™ 1.0 and 2.0) — sole Platinum-rated cloud among 84 providers evaluated, including AWS, GCP, Azure, Crusoe, Nebius, Lambda

- "CoreWeave sets the bar for others to follow and is the only cloud to consistently command premium pricing" — SemiAnalysis quote
- Record-breaking **MLPerf** results: **29x faster** than the next best competitor (GPT-3 training); **more than 2x faster** training performance than NVIDIA Hopper-based systems

Performance Metrics

- **Up to 35x faster and 80% less expensive** than generalized public clouds
- **Up to 20% higher GPU cluster performance** (Model FLOPS Utilization / MFU) and **10x greater reliability** on thousand-GPU clusters
- **96% goodput** vs. the industry average of 90%
- **10x faster inference spin-up times**
- **50% fewer interruptions per day**
- General-purpose clouds "on average only unlock 35–45% of GPU cluster performance potential"

First-to-Market GPU Access

- First cloud provider to make **NVIDIA GB200 NVL72** instances generally available; first with **H200**, first with **H100 HGX**, first with **GB300 NVL72**

Business Milestone

- "Fastest cloud platform in history to surpass \$5 billion in annual revenue"

Positioning Against Specific Competitors

Against legacy hyperscalers (AWS, Azure, GCP):

- "Generalized clouds operated by traditional hyperscalers were not built to serve the specific requirements of AI"
- "CoreWeave's infrastructure...is built from the ground up and hyper-optimized for AI workloads, unlike solutions from traditional cloud providers that were designed for web-scale and are encumbered by a legacy technical architecture"

Against neoclouds:

- "Neoclouds often look cheaper up front but surprise you with hidden costs" — positions CoreWeave as full-stack vs. fragmented neocloud alternatives

Third-Party Endorsements Used in Positioning

- **Jensen Huang (NVIDIA CEO):** "CoreWeave is a world-class new generation AI-Native cloud"
- **Chen Goldberg (SVP Engineering):** "CoreWeave is light years ahead and delivering what others cannot"
- **SemiAnalysis:** Only provider to earn Platinum; "sets the bar for others to follow"

Overall Brand Identity

Overall Brand Identity — Consolidated Summary

Mission

"To power the creation and delivery of the intelligence that drives innovation."

This statement appears consistently across the site, sometimes expanded to reference "the next wave of AI innovation." A complementary tagline reinforces purpose: **"We make AI better, so pioneers can make the world better."**

Brand Values (Explicitly Stated)

The About page lists five formal values, supported by additional implicit values surfaced throughout the site:

#	Value	Key Language
1	Be curious at your core	"Curiosity drives innovation... challenge the status quo"
2	Empower employees	"Trust each other to make decisions, communicate transparently"
3	Act like an owner	"Take full responsibility... step up, innovate, and drive progress"
4	Deliver best-in-class experiences	"Customer-obsessed... exceed expectations at every step"
5	Achieve more together	"Collaboration is at the heart of our success"

Additional recurring value themes:

- **Pioneer spirit** — "Built for pioneers by pioneers"; the word *pioneers* describes both CoreWeave and its customers throughout the site.
- **Speed / urgency** — "Move at the pace of innovation"; "Getting to market quickly isn't just a matter of growth. It's a matter of survival."
- **Transparency & openness** — "You see what we see"; "Open standards, no hidden fees"; complete visibility into performance and cost.
- **Open-source commitment** — Founding member of Red Hat's llm-d project; Tensorizer open-source; commitment to maintaining Marimo's open-source project; emphasis on interoperability (e.g., Weights & Biases).
- **Technical / engineering excellence** — Heavy use of benchmarks, MLPerf results, specifications.

Tone of Voice

- **Confident and authoritative:** "We're built for this"; "This is where CoreWeave stands apart"; "We're the #1 AI cloud."
 - **Technical credibility blended with aspiration:** Benchmarks and specs sit alongside language like "Bring us your biggest breakthroughs. Your moonshots. Your brightest ideas." and "Pioneering the Possible."
 - **Urgent, action-oriented:** "Bleeding-edge," "at the epicenter," "hypergrowth," "unprecedented pace."
 - **Customer-centric:** "Customer obsession" is a phrase used by leadership and marketing alike.
 - **Founder-driven authenticity:** CEO Michael Intrator writes personally about the journey; co-founder Brannin McBee frames the origin story around "audacity, humility, and the support of other believers."
-

Culture & Ethos

- **Entrepreneurial and mission-driven:** Leadership consistently describes a "high-performing work environment" with a startup mentality (Michelle O'Rourke, Michael Intrator).
 - **Rapid growth narrative:** Founded 2017 by "3 guys in a garage," now 800+ team members worldwide and Nasdaq-listed (March 2025). This origin-to-scale arc is a central identity motif.
 - **Engineering-excellence driven:** Fast-moving, fearless problem-solving; first-to-market positioning emphasized repeatedly.
 - **Inclusive and learning-oriented:** Intrator describes a "vibrant and inclusive culture that promotes learning, growth, and innovation."
 - **Strategic partnership identity:** Brand is strongly tied to the NVIDIA partnership and a broader posture of "transformative partnership" with customers and ecosystem players, offering "deep consultative expertise and direct-to-expert technical support."
-

Key Takeaway

CoreWeave's brand identity is built on a **pioneer archetype**: technically elite, fast-moving, and founder-authentic, with a clear promise that it exists to serve other pioneers. The identity balances **hard technical credibility** (benchmarks, open-source contributions, NVIDIA alignment) with **aspirational, urgency-laden storytelling** — consistently positioning speed, transparency, and engineering excellence as the traits that set it apart in the AI cloud market.

Value Proposition Statements

Consolidated Summary: CoreWeave Value Proposition Statements

Core Positioning

- **"CoreWeave Cloud is an AI-native platform purpose-built for AI. It combines next-generation infrastructure, intelligent tools, and expert support to power the world's most complex AI workloads—faster, more efficiently, and with complete transparency."**
- **"We bridge the gap between AI ambition and execution by delivering high-value performance, scale, and expertise with the infrastructure AI needs today and in the future."**
- **"Get your next big breakthrough to market faster with CoreWeave."**
- CFO-oriented TCO messaging centers on **"performance-per-dollar, not the sticker price."**

Quantified Performance Claims

GPU Cluster Performance & Reliability

- **Up to 20% higher GPU cluster performance** than alternative solutions (via CoreWeave Kubernetes Service / CKS)
- **20% higher MFU (Model FLOPs Utilization) — exceeding 50% on NVIDIA Hopper GPUs** vs. industry typical 35–45%
- **96% cluster goodput** (vs. ~90% industry average) — production-ready on Day 1
- **50% fewer interruptions per day**
- **3.66 days MTTF** at 1,024 GPUs (10x improvement over 0.33-day baseline)
- **97.5% ETTR (Effective Training Time Ratio)**

Training Throughput

- **2.5x faster training throughput** vs. neoclouds
- **Llama 3.1 405B trained in 27.3 minutes** (2x faster than Hopper at same scale)
- **Up to 4x faster LLM training** compared to previous GPU generations
- **8x faster checkpointing** via async Tensorizer implementation

Inference Performance

- **10x faster inference spin-up times; 5x faster model download speeds**
- **800 TPS on Llama 3.1 405B** (2.86x per-chip improvement over H200)
- **33,000 TPS on Llama 2 70B** (40% improvement over H100)
- **Up to 30x faster real-time LLM inference** vs. previous generations
- **6.5x performance improvement on DeepSeek R1 inference** (GB300 vs. H100)
- **3.42x higher token generation on Kimi K2.5** (HGX B300 vs. HGX H200)

Next-Generation GPU Leadership

- **Consistently first to market** deploying NVIDIA GB200 NVL72, GB300 NVL72, HGX B300, and RTX PRO 6000

Cost & TCO Advantages

- **Up to 25x lower TCO and 25x less energy** for real-time inference (next-gen GPUs)
- **More than 75% lower storage costs** with CoreWeave AI Object Storage

- **Zero egress, ingress, and request/transaction fees** on storage
- **40% lower cost than self-managed** for Serverless RL
- **Zero Egress Migration program** saves customers **up to \$1 million +** for typical data transfers

Storage Performance

- **Up to 7 GB/s per GPU throughput** (LOTA technology, regardless of data location)
- **99.9% uptime** and **11 nines durability** for AI Object Storage

Speed to Market

- "From final fit-out to fully provisioned capacity, we accelerate delivery to weeks versus months."
- W&B acquisition: first joint product launched **after just six weeks**

Named Customer Outcomes

Customer	Outcome
Mistral AI	2.5x faster training on GB200 NVL72; "cut training time in half"
Cohere	3x faster training on GB200 NVL72 for 100B-parameter models
IBM	80% performance boost; supercomputer "twice the size of Blue Vela"
Tarteel	22% latency improvement, ~56% cost reduction (migrated from AWS)
NovelAI	Served requests 3x faster
General Intuition	20–30% better performance on CoreWeave ARENA
Cursor	"Proven operating model gives us confidence to focus on building... rather than infrastructure risk"
Perplexity	Inference partnership with "predictable cost characteristics"
Poolside	40,000+ GPU partnership agreement
OpenAI	Total contract value up to ~\$22.4 billion

"We Help X Do Y by Z" Pattern

- "Maximize your potential with CoreWeave and get to market faster, focusing on

what you do best: building, training, and serving AI models that are transforming the world."

- "CoreWeave simplifies this, collaborating closely with NVIDIA to deliver optimized solutions that enable customers to fully capture the value of their AI investments."
-

Ideal Customer Profile

[Synthesis error: Server error '500 Internal Server Error' for url 'http://127.0.0.1:7012/generate' For more information check: <https://developer.mozilla.org/en-US/docs/Web/HTTP/Status/500>]

Implied Marketing Strategy

Consolidated Implied Marketing Strategy

CoreWeave's overarching strategy is to **define and own the "AI-native cloud" category** — positioning itself as the purpose-built, vertically-integrated alternative to general-purpose hyperscalers (AWS, Azure, GCP) and the dominant platform for a new computing paradigm, analogous to what AWS became for web infrastructure.

1. Hardware Moat via Deep NVIDIA Partnership

- CoreWeave is NVIDIA's closest cloud partner: first Elite CSP, NVIDIA Exemplar Cloud validation, and a **\$2 billion NVIDIA investment** (January 2026) with Jensen Huang's personal endorsement.
- **Consistently first-to-market** with every NVIDIA GPU generation — H100, H200, GB200 NVL72, GB300 NVL72, HGX B300, RTX PRO 6000, with planned Vera Rubin NVL72.
- Partnership extends beyond GPUs to NVIDIA CPUs (Vera), storage (BlueField DPUs), networking (Quantum InfiniBand), and NVIDIA AI Enterprise software support.
- Creates a **self-reinforcing cycle**: NVIDIA steers early GPU allocations to CoreWeave → CoreWeave delivers first-to-market access → customers choose CoreWeave → NVIDIA deepens the relationship.
- Directly addresses the market's top concern: **GPU scarcity and capacity constraints** (250,000+ GPUs, \$66.8B revenue backlog).

2. Vertical Integration — "Metal to Model"

CoreWeave is building depth at every layer of the stack to create compounding performance advantages that cannot be replicated by assembling third-party components:

- **Physical infrastructure**: Purpose-built data centers with liquid cooling, 130kW + racks; ownership deepened via **Core Scientific acquisition** (~\$9B, 1.3 GW of power).
- **Networking**: InfiniBand at scale via NVIDIA Quantum.
- **Storage**: Proprietary CAIOS/LOTA storage, AI Object Storage with zero egress fees.

- **Orchestration & compute:** SUNK, CKS, BlueField DPU integration for zero-trust security.
- **Observability & management:** Mission Control for production-grade reliability, security, and compliance.
- **Developer tools & MLOps:** Acquired **Weights & Biases** (~\$1.4–1.7B) for experiment tracking, model management, evaluation, and monitoring.
- **RL & agentic AI:** Acquired **OpenPipe** (September 2025) for reinforcement learning tools.
- **Developer experience:** Acquired **Marimo** (October 2025), open-source Python/AI notebook.
- **Enterprise verticals:** Acquired **Monolith** (October 2025) for AI in industrial/physics engineering.

This creates an end-to-end platform spanning "**compute → model management → AI application evaluation and monitoring**," increasing switching costs and platform stickiness.

3. Performance-as-Proof — Third-Party Validation Stack

- **SemiAnalysis ClusterMAX™ Platinum** — sole recipient, twice — serves as independent, credible proof for technical and enterprise buyers.
- **MLPerf benchmark records** demonstrate measurable performance leadership.
- **NVIDIA Exemplar Cloud** validation reinforces the NVIDIA partnership narrative.
- Each new benchmark and validation compounds the previous ones, building a **defensible "proof stack"** that is difficult for competitors to replicate quickly.

4. Land-and-Expand Go-to-Market Motion

- **Initial beachhead:** AI-native labs (OpenAI, IBM, Mistral, Cohere, Jane Street, Cursor) validate technical credibility and generate peer proof for subsequent buyers.
- **CoreWeave ARENA:** A production-ready AI lab offering structured "try before you buy" — customers run real workloads on purpose-built infrastructure to validate performance, scaling, and cost before committing. Lowers evaluation friction and drives evidence-based conversion.
- **Enterprise expansion:** Products like Mission Control (observability, compliance), CoreWeave Federal (government), and Monolith (industrial AI) broaden the addressable market beyond AI labs.
- **CoreWeave Ventures:** Compute-for-equity and capital investments in startups create a pipeline of future customers and ecosystem network effects.

5. Pricing Strategy — Reframing the Value Conversation

- **Transparent pricing:** No hidden fees, zero egress, zero request fees — directly attacking hyperscaler lock-in and opacity.
- **Flexible Capacity Plans:** Reservations, Flex Reservations, Spot, On-Demand — address the binary choice problem and accommodate variable demand patterns (especially for agentic AI).
- **Zero Egress Migration** eliminates switching costs from hyperscalers.
- **"Performance-adjusted cost"** framing shifts the conversation from \$/GPU/hour sticker price to total value delivered, neutralizing cheaper neoclouds and more expensive hyperscalers simultaneously.

- **TCO consultation** positions CoreWeave as a transparent partner vs. opaque competitors.

6. Market Dynamics Being Exploited

- **GPU scarcity:** CoreWeave's NVIDIA relationship provides preferential access and capacity guarantees in a supply-constrained market.
- **Hyperscaler legacy debt:** AWS/Azure/GCP were built for web workloads; their architectures are positioned as fundamentally misaligned for AI's compute, networking, and cooling demands.
- **AI moving from experimentation to production:** "AI has passed the experimentation phase and is now firmly in production" (Davos 2026). CoreWeave emphasizes production-grade reliability, observability, and TCO.
- **Inference demand explosion:** "The future growth will be fueled by inference." Heavy investment in Dedicated Inference, llm-d, NVIDIA Blackwell/HGX B300 for inference workloads.
- **Agentic AI emergence:** Multiple product launches — agent development tools, Serverless RL, Flex Reservations for variable demand, W&B Inference — designed for agentic AI workload patterns.
- **Data center redesign megatrend:** Legacy facilities cannot support modern AI power/cooling demands; CoreWeave's purpose-built facilities are designed from the ground up.

7. Geographic Expansion & Sovereign AI

- **US:** 28+ data centers across 15+ states; \$6B+ Pennsylvania investment with political support (Senator McCormick, Governor Shapiro, President Trump); Core Scientific acquisition deepens footprint.
- **US Federal:** CoreWeave Federal entity, D.C. office, FedRAMP pursuit, DOE Genesis Mission, Senate testimony, Pentagon/defense positioning.
- **UK:** £2.5B investment, London HQ, renewable energy commitment, UK government partnership.
- **Europe:** Sweden (EcoDataCenter), Nordic presence at Techarena, \$3.5B European investment commitment, sustainability messaging.
- Addresses **data sovereignty and residency requirements** and follows demand patterns while building geographic redundancy.

8. Brand and Awareness Strategy

- **Technical credibility first:** Named customer logos, executive quotes, customer stories, and ARENA case studies serve as "peer proof" for technical buyers.
- **Broadening awareness:** Aston Martin F1 partnership provides consumer brand visibility beyond technical audiences.
- **"Ready for Anything, Ready for AI"** brand campaign featuring Chance the Rapper signals a deliberate move from B2B technical marketing to broader awareness — preparing for enterprise and potentially public-market brand building.
- **Open-source ecosystem alignment:** Marimo, OpenPipe ART, W&B interoperability, and llm-d contributions signal developer-friendly values and reduce perceived lock-in risk.

9. Competitive Moat Summary

CoreWeave is constructing a **multi-layered, compounding moat**:

1. **Privileged NVIDIA hardware access** (supply advantage + credibility)
2. **Purpose-built software stack** from bare metal to developer tools (vertical integration)
3. **Owned data center infrastructure** via Core Scientific (physical control + cost structure)
4. **Third-party performance validation** (SemiAnalysis, MLPerf, NVIDIA) that compounds over time
5. **Ecosystem of AI-native customers, startups, and developer tools** that creates network effects and switching costs

The implied ambition: **become to AI infrastructure what AWS was to web infrastructure — the default platform for a new computing paradigm.**

Operational Metrics

Source: coreweave.md

Marketing – Factual & Quantitative: <https://coreweave.com>

Extract concrete, verifiable facts from a company's website: corporate details, infrastructure specs, customer names, product listings, and quantitative claims.

Based on 328 stored pages analysed in 1 batch(es) with 1 AI calls (namespace: webscrape, model: claude-opus-4-6, profile: operational_metrics).

Number Of Data Centers, Their Locations, Gpus, And Power Capacity

Data Centers:

- **43 state-of-the-art data centers** (from infrastructure page)
- **42 active data centers** by end of 2025 (from CEO year-end review): "Since the beginning of the year, we brought 9 new data centers online across the United States and Europe for a total of 42 active data centers in our network."
- **28 data centers** by end of 2024; anticipated doubling globally
- Data centers located in **North America and Europe**, across **over 33 data centers** in those regions (infrastructure page)
- **28 of which are strategically located across the US in the most important markets across 15 states**
- In 2023, CoreWeave had 14 data centers; by end of 2024, 28 data centers

Specific Locations Mentioned:

- Las Vegas, NV (Switch Tier 5® data center)
- LGA1 – New York City (at the mouth of the Lincoln Tunnel, Hudson River), ISO 27001

certified

- Crawley, UK (with Digital Realty) – operational October 2024
- London Docklands, UK (with Global Switch) – operational December 2024
- London – European headquarters
- Lancaster, Pennsylvania – new data center, initial 100 MW with potential to expand to 300 MW, \$6 billion commitment
- Plano, Texas – \$1.6 billion data center announced 2023
- Livingston, NJ (corporate HQ: 290 W Mt Pleasant Ave Suite 4100, Livingston, NJ 07039)
- Washington, D.C. – office at 1001 Pennsylvania Avenue NW, expanded to 22,590 sq ft
- Scotland (partnership with DataVita for sovereign AI infrastructure)
- Sweden (EcoDataCenter partnership for NVIDIA Blackwell clusters)
- New Jersey – first greenfield build
- United Kingdom – multiple sites, total investment £2.5 billion

GPUs:

- **250,000+ high-performance GPUs**
- GPU types offered: NVIDIA GB300 NVL72, GB200 NVL72, HGX B300, HGX B200, RTX PRO 6000 Blackwell Server Edition, HGX H100, HGX H200, GH200, L40, L40S, A100, A40, RTX A6000, RTX A5000, RTX A4000, Quadro RTX 5000, Quadro RTX 4000, H100 PCIe
- Mega clusters with **100k+ GPUs**; superclusters of 100,000+ GPUs interconnected to create **300k+ GPU megaclusters**
- Clusters up to **42,000 H200 GPUs** deployed
- NVIDIA HGX H100 clusters scaling up to **16,384 H100 SXM5 GPUs** under the same InfiniBand fabric
- Poolside partnership: **more than 40,000 GPUs** (GB300 NVL72 systems)
- IBM deployment: "several thousand GB200s" enabling a supercomputer "twice the size of their previous supercomputer, Blue Vela"
- MLPerf Training v5.0: deployed **2,496 NVIDIA Blackwell GPUs** across 39 racks
- HGX B200 deployment: "more than 10,000 NVIDIA Blackwell GPUs"

Power Capacity:

- **3.1 GW+ of contracted power capacity** (infrastructure page)
- **2.9 GW of contracted capacity** (CEO year-end 2025 review)
- **~590 megawatts of active power** by end of Q3 2025, with "a clear path toward exceeding 850 megawatts by year's end"
- Goal: **more than 5 gigawatts of AI factories by 2030**
- Lancaster, PA facility: up to **300 MW**
- Core Scientific acquisition: approximately **1.3 GW of gross power** across Core Scientific's footprint, with **incremental 1 GW+ potential**
- UK investment: total **£2.5 billion**
- Europe: commitment to invest **\$3.5 billion across Europe**
- Rack density: up to **130 kW per rack**
- NVIDIA GB200 NVL72 racks: **~130 kW per rack**

Active And Contracted Megawatts, Terawatts, And Other Measures Of Power

- **3.1 GW + of contracted power capacity** (infrastructure page claim)
 - **2.9 GW of contracted capacity** (Q3 2025 financial disclosure)
 - **~590 MW of active power** by end of Q3 2025
 - **Exceeding 850 MW** targeted by end of 2025
 - **More than 260 megawatts online in a single quarter** (from defining year blog)
 - **1.6 GW + of contracted power capacity** (from May 2025 Senate testimony – earlier figure)
 - Lancaster, PA data center: initial **100 MW**, expandable to **300 MW**
 - Core Scientific acquisition: **~1.3 GW gross power** plus **1 GW + incremental expansion potential**
 - Poolside Project Horizon: **2 GW AI campus** in West Texas, CoreWeave anchor tenant for first phase of **250 MW gross power** with option for additional **500 MW**
 - McKinsey cited: U.S. data center demand will balloon to **80 gigawatts by 2030**
 - UK data centers: powered entirely by **renewable energy** (Crawley and London Docklands)
 - Scotland deployment: backed entirely by **renewable energy** with closed-loop cooling
 - No specific PPA details disclosed; renewable energy used across European data centers
-

Installed, Active And Future Number Of Gpus

- **250,000 + high-performance GPUs** currently (infrastructure page)
 - **Over 250,000 GPUs** in more than 30 data centers (Senate testimony, May 2025)
 - Mega clusters: **100k + GPUs** in single clusters; **300k + GPU megaclusters** via interconnected superclusters
 - MLPerf Training v5.0 submission: **2,496 NVIDIA Blackwell GPUs** (largest-ever Blackwell cluster submission by a CSP), **39x larger** than the next CSP submission
 - Previous MLPerf (v3.0/3.1): **3,500 + NVIDIA H100 GPUs** (record-breaking GPT-3 training)
 - H200 clusters: deployed in clusters with up to **42,000 GPUs**
 - HGX B200 initial deployment: **more than 10,000 NVIDIA Blackwell GPUs**
 - Poolside partnership: **more than 40,000 GPUs** (GB300 NVL72)
 - IBM partnership: "several thousand GB200s"
 - NVIDIA HGX H100 fabric: scalable up to **16,384 H100 SXM5 GPUs**
 - GPU models available: GB300 NVL72, GB200 NVL72, HGX B300, HGX B200, RTX PRO 6000 Blackwell Server Edition, HGX H100, HGX H200, GH200, L40, L40S, A100, plus legacy (A40, A6000, A5000, A4000, RTX 5000, RTX 4000, V100)
 - First to market with: NVIDIA H100, H200, GB200 NVL72 (GA), GB300 NVL72, HGX B200 (GA), HGX B300 (GA), RTX PRO 6000 Blackwell Server Edition
 - Future: NVIDIA Vera Rubin NVL72 platform expected in second half of 2026; planning **more than 5 GW** of AI factories by 2030
-

Teraflops And Exaflops

- **NVIDIA GB200 NVL72**: delivers up to **1.4 exaFLOPS** of AI compute power per rack (mentioned multiple times)
- **Model FLOPS Utilization (MFU)**: CoreWeave achieves **exceeding 50% MFU** on NVIDIA Hopper GPUs; industry average is **35-45%**; up to **20% higher** than alternatives
- Specific MFU benchmarks (NVIDIA Exemplar Cloud):

○ DeepSeek v3 (BF16): **1.9% greater MFU** than NVIDIA reference target on 512

- GPUs
 - Grok-1 314B (BF16): **4.7% greater MFU** than NVIDIA reference target on 512 GPUs
 - Llama 3.1 405B (FP8): **3.8% greater MFU** on 512 GPUs
 - Llama 3 70B (FP8): **2.4% greater MFU** on 512 GPUs
 - H100 benchmark study: achieved **51-52% MFU** on NVIDIA H100 GPUs (vs. 35-45% typically reported)
 - NVIDIA L40S: **733 TFLOPs**
 - NVIDIA InfiniBand bidirectional throughput: up to **3.2 Tbps on 8NDR** per node
 - NVIDIA HGX H100: **3.2Tbps of bandwidth per node**
 - RTX PRO 6000 Blackwell: over **5x AI performance at FP4 precision** vs. L40S; up to **5.3x faster LLM inference**
-

Other Capacity And Performance

Revenue and Financial:

- **\$5 billion in annual revenue** in fiscal year 2025 (fastest cloud platform in history to surpass this)
- **168% year-over-year revenue growth** in 2025
- **\$66.8 billion revenue backlog**
- OpenAI total contract value: **~\$22.4 billion** (\$11.9B + \$4B + \$6.5B)
- **\$12 billion + raised** from equity and debt investors (as of mid-2024)
- IPO priced at **\$40/share**, offering **37.5 million shares** on Nasdaq under "CRWV"
- NVIDIA invested **\$2 billion** in CoreWeave at \$87.20/share
- \$7.5 billion debt facility (led by Blackstone, 2024)
- \$2.3 billion debt facility (2023)
- \$1.1 billion Series C (May 2024)
- Cisco \$650 million secondary share sale investment
- Core Scientific acquisition: implied equity value of **~\$9.0 billion**
- Weights & Biases acquisition: **~\$1.7 billion**
- UK total investment: **£2.5 billion**
- Pennsylvania: **\$6 billion** commitment
- Europe overall: **\$3.5 billion** investment commitment

Performance Metrics:

- **96% cluster goodput** (vs. 90% industry average)
- **50% fewer interruptions per day**
- **10x faster inference spin-up times**
- **2.5x faster training throughput** and **20% higher MFU** (vs. neoclouds)
- **Up to 20% higher GPU cluster performance** than alternative solutions
- MLPerf Training v5.0: Llama 3.1 405B trained in **27.3 minutes** (more than 2x faster than Hopper systems)
- MLPerf Inference v5.0: **800 tokens per second** on Llama 3.1 405B (NVIDIA GB200); **33,000 TPS** on Llama 2 70B (H200, 40% improvement over H100)
- H100 benchmark: **3.66 days MTTF** at 1,024 GPUs (10x improvement over 0.33 days baseline); **97.5% ETTR**
- NVIDIA HGX B300: **3.42x higher token generation** on Kimi K2.5 vs. HGX H200; **4.93x faster end-to-end request latency** on DeepSeek-R1 vs. HGX H200
- GB300 NVL72: **6.5x performance improvement** on DeepSeek R1 inference; **50x**

increase in output for reasoning model inference vs. Hopper

- Cohere on GB200: **3x faster training** for 100B parameter models
- Mistral AI: **2.5x faster training** on GB200 NVL72
- IBM: **up to 80% faster** performance on CoreWeave

Storage Performance:

- CoreWeave AI Object Storage: up to **7 GB/s per GPU** throughput (on Blackwell Ultra); **2 GB/s per GPU** on H200
- **Exabyte-scale** capacity, **S3-compatible**
- **Eleven 9's of durability** (99.999999999%), **99.9% uptime**
- **Zero egress, ingress, or request fees**
- Storage costs reduced by **more than 75%** with automated usage-based billing
- Data migration: typically exceeding **1 PB per day**

Networking:

- NVIDIA Quantum-2 InfiniBand: up to **3200 Gbps per node** (400Gb/s per GPU)
- NVIDIA Quantum-X800 InfiniBand: **800 Gb/s per GPU** (for B300/GB300)
- Non-blocking architecture supporting **100k+ GPUs** in a single cluster
- Direct Connect: port speeds up to **400 Gbps**
- **Zero data transfer or egress fees** for networking

Workforce:

- Grew from ~900 to **more than 2,000 employees** worldwide in 2025
- Lancaster PA facility: over **600 local construction jobs**, at least **75 operational jobs** (scaling to 175+)

Key Customers Named:

- OpenAI, IBM, Mistral AI, Cohere, Jane Street, Cognition, Cursor, Mercado Libre, Midjourney, Runway, Perplexity, Moonvalley, Decart AI, Poolside, Cline, NovelAI, Bit192, Trillion Labs, General Intuition, Aston Martin Aramco F1 Team, Redesign Science, Tarteel, Product Insight, Meta, LG AI Research, Riskfuel, Inflection AI, SquadStack.ai, QA Wolf, inference.net, Periodic Labs, Hologen, Replicate, EleutherAI, Anlatan/NovelAI

Certifications and Ratings:

- Only AI cloud to earn **SemiAnalysis Platinum ClusterMAX™** rating twice (both 1.0 and 2.0)
 - **NVIDIA Exemplar Cloud** for both training and inference on GB200 NVL72
 - **NVIDIA's first Elite Cloud Services Provider** for Compute
 - **SOC 2 Type II** certified (Bare Metal and CKS)
 - Pursuing **FedRAMP** and other federal authorizations
 - **ISO 27001** certified (LGA1 data center)
-
-