Consolidated Report: iren

Generated 2026-03-29 04:08 UTC from 3 analysis profile(s): company_profile, marketing_analysis, operational_metrics.

> **Note:** The primary source of information in this report is the company's own website. The content reflects the company's self-reported claims and marketing materials and should not be treated as independently verified fact.

Executive Summary

Executive Summary: IREN Limited

Company Overview

IREN Limited (NASDAQ: IREN), headquartered in Melbourne, Australia with operations spanning North America, is a vertically integrated AI cloud service provider and next-generation data center operator. Originally founded as Iris Energy — a Bitcoin mining company — IREN has undergone a strategic transformation to become a large-scale provider of GPU-as-a-service, colocation, and build-to-suit data center solutions purpose-built for artificial intelligence and high-performance computing workloads. The company is led by co-founders and co-CEOs Daniel and Will Roberts, and its stock was added to the MSCI USA Index in February 2026.

Key Findings

Infrastructure & Scale: IREN has secured over 4.5 GW of power capacity across six data center campuses totaling 4,500 acres in British Columbia, Texas, and Oklahoma. Of this, 810 MW is operational, 2,100 MW is under construction, and 1,600 MW is in development. The company's GPU fleet has expanded dramatically to over 150,000 NVIDIA GPUs as of March 2026, spanning H100, H200, B200, B300, and GB300 NVL72 architectures. Its Sweetwater, Texas mega-campus alone can support over 700,000 liquid-cooled Blackwell GPUs at full buildout. Canadian facilities achieve a best-in-class PUE of 1.1.

Commercial Traction: IREN's most significant commercial validation is a $9.7 billion, five-year agreement with Microsoft to deliver GPU cloud infrastructure at its 750 MW Childress, Texas campus, expected to generate approximately $1.94 billion in annualized run-rate revenue when fully commissioned. This deal was supported by a $5.8 billion equipment purchase from Dell Technologies and a $2.3 billion convertible notes offering. The company operates as a Preferred NVIDIA Cloud Partner, building all AI cloud infrastructure on NVIDIA reference architecture.

Strategic Positioning & Brand: IREN positions itself as the vertically integrated alternative to legacy cloud providers, owning the full stack from land and power substations through to GPU hardware. Its go-to-market strategy emphasizes speed-to-deployment, cost transparency (including no data ingress/egress fees and ~$0.05/kWh all-in power costs in West Texas), and 100% renewable energy sourcing. The company targets hyperscalers, AI-native companies, and large enterprises through a sales-led model with no self-serve pricing, reflecting the high-value, contract-based nature of its customer relationships.

Cross-Cutting Themes, Strengths, and Risks

Strengths: IREN's primary competitive moat is its early lock-up of gigawatt-scale, grid-connected power — the single most constrained resource in AI infrastructure today. Its pivot from Bitcoin mining gave it a structural head start in securing low-cost energy and operational expertise. The Microsoft anchor contract provides both revenue visibility and enterprise credibility, while the NVIDIA partnership ensures preferential technology access and market legitimacy.

Risks: IREN's concentration risk is notable — a single $9.7 billion contract with Microsoft represents transformative but heavily concentrated revenue. The company's rapid capital deployment ($5.8 billion in equipment, $2.3 billion in convertible debt) creates significant financial leverage and execution risk across simultaneous multi-gigawatt construction projects. IREN operates purely at the infrastructure layer, offering no software, MLOps, or developer tooling, which limits its addressable market to infrastructure buyers and leaves it dependent on sustained demand for bare-metal GPU compute. Additionally, its renewable energy claims rely partly on REC purchases rather than exclusively on-site generation, which may face increasing ESG scrutiny.

Strategic Outlook

IREN is executing an ambitious land-and-power land-grab strategy at a moment of unprecedented AI infrastructure demand. If the company successfully delivers on its construction pipeline — energizing approximately 2 GW between 2026 and 2027, with an additional 1.6 GW by 2028 — it will control one of the largest independent AI-ready power portfolios in North America. The ability to attract additional hyperscaler customers beyond Microsoft will be the critical test of whether IREN can evolve from a concentrated, project-driven business into a diversified AI infrastructure platform. Its trajectory over the next 18–24 months will be defined by construction execution, customer diversification, and prudent capital management amid a rapidly scaling balance sheet.

Company Profile

Source: iren.md

Marketing – Factual & Quantitative: https://iren.com

Extract concrete, verifiable facts from a company's website: corporate details, infrastructure specs, customer names, product listings, and quantitative claims.

Based on 103 stored pages analysed in 1 batch(es) with 1 AI calls (namespace: webscrape, model: claude-opus-4-6, profile: company_profile).

Headquarters

IREN Limited's registered address is **Level 13, 664 Collins Street, Docklands VIC 3008** (Melbourne, Australia). The company also has offices in **Sydney, Australia** (corporate functions supporting global operations) and **Vancouver, Canada**. Teams operate across the United States, Australia, and Canada.

Years In Business

The Community Grants Program has "funded local initiatives since **2022**." Earlier investor presentations reference the company under its former name "**Iris Energy Limited**" (e.g., "Iris Energy 2Q FY23 Results Conference Call" in Feb 2023; "Iris Energy 2022 Annual Meeting of Shareholders" in Dec 2022). The Supplier Code of Conduct refers to "IREN Limited (formerly known as Iris Energy Limited)." No explicit founding year is stated on the website, but the company's history as Iris Energy dates back to at least 2022 based on SEC filings and presentations.

Company Executive Names And Titles

Board of Directors:

- **David Bartholomew** – Independent Chair
- **Mike Alfred** – Independent Non-Executive Director
- **Chris Guzowski** – Independent Non-Executive Director
- **Sunita Parasuraman** – Independent Non-Executive Director
- **Daniel Roberts** – Co-Founder and Co-CEO (also a board member)
- **Will Roberts** – Co-Founder and Co-CEO (also a board member)

Executive Leadership:

- **Anthony Lewis** – Chief Financial Officer (appointed September 2025)
- **David Shaw** – Chief Operating Officer
- **Kent Draper** – Chief Commercial Officer
- **Denis Skrinnikoff** – Chief Technology Officer
- **Cesilia Kim** – Chief Legal Officer
- **John Gross** – Chief Innovation Officer (appointed February 2026)

Other named personnel:

- **Tim Delcourt** – Commercial Director, AI & HPC
- **Heather Miller** – Vice President of People, Culture and Community
- **Noah Clark** – Senior Operations Manager for HPC Data Centers
- **Trey Lee** – Operations Manager, Childress, TX
- **Pamela Teshereau-Viau** – Project Manager, Prince George, BC
- **Shaylene Sager** – Site Manager, Canal Flats, BC
- **Alexandra Ayton** – Executive Assistant, Sydney, NSW

Customer Names, Locations, Industry/Type Of Entity, And

Other Relevant Details

- **Microsoft –** Multi-year agreement valued at approximately **$9.7 billion** to deliver GPU cloud infrastructure powered by NVIDIA GB300 GPUs. Deployment at IREN's 750 MW Childress, Texas campus. GPUs deployed in 4 phases through 2026 (Horizon 1–4) providing 200MW of critical IT load. Five-year contract with 20% prepayment, expected ~$1.94 billion annualized run-rate revenue once fully commissioned. **Jonathan Tinter**, President of Business Development and Ventures at Microsoft, is quoted. Industry: hyperscaler / cloud computing.
- The website also mentions IREN has expanded its AI Cloud capacity to **150,000 GPUs** (March 2026 news release) and was "Added to MSCI USA Index" (Feb 2026).

Partner Names, Locations, Industry/Type Of Entity, And Other Relevant Details

- **NVIDIA –** IREN is a **Preferred NVIDIA Cloud Partner**. AI Cloud is "built on NVIDIA reference architecture." GPU models deployed include NVIDIA H100, H200, B200, B300, and GB300 NVL72. Networking uses NVIDIA Quantum-2 400Gb/s InfiniBand with 3.2TB/s bandwidth.
- **Dell Technologies –** IREN entered into an agreement with Dell to purchase GPUs and ancillary equipment for approximately **$5.8 billion** (in conjunction with the Microsoft partnership). Also featured in an "IREN and Dell Success Story" video.
- **WEKA.IO –** Storage partner; WEKA.IO storage is offered across GPU clusters.
- **Canaccord Genuity –** Investment bank; hosted fireside chats with IREN.
- **Cantor Fitzgerald –** Investment bank; hosted fireside chat with IREN (Feb 2024).
- **Jefferies –** Investment bank; hosted fireside chat with IREN (Apr 2025).
- **BC Hydro –** Power grid provider for Canadian sites (Canal Flats, Mackenzie, Prince George connected to BC Hydro grid).
- **ERCOT –** Power grid for Texas sites (Childress, Sweetwater connected to ERCOT grid).
- The website references "Tier 1 engineering firms and vendor partners" and "trusted delivery partners" but does not name them beyond the above.

Stock analyst coverage firms: Bernstein (Gautam Chhugani), B Riley (Nick Giles), BTIG (Gregory Lewis), Canaccord Genuity (Joseph Vafi), Cantor Fitzgerald (Brett Knoblauch), Citizens (Gregory Miller), Compass Point (Michael Donovan), H.C. Wainwright & Co (Mike Colonnese), Jones Research (Stephen Glagola), JPMorgan (Reginald L. Smith), Macquarie (Paul Golding), Needham (John Todaro), Roth Capital Partners (Darren Aftahi).

Products

- **IREN Cloud™ (AI Cloud) –** GPU-as-a-service offering bare-metal NVIDIA GPU clusters for AI training and inference. Available GPU configurations:
    - **NVIDIA H100**: 80GB GPU memory, 2,048 GB RAM, 224 vCPUs, 30TB NVMe storage
    - **NVIDIA H200**: 141GB GPU memory, 2,048 GB RAM, 224 vCPUs, 30TB NVMe storage
    - **NVIDIA B200**: 180GB GPU memory, 2,048 GB RAM, 224 vCPUs, 15.4TB NVMe storage

    - **NVIDIA B300**: 288GB GPU memory, 3,096 GB RAM, 256 vCPUs, 30TB NVMe storage
    - **NVIDIA GB300 NVL72**: 288GB GPU memory, 20,736 GB RAM, 2,592 vCPUs, 276TB NVMe storage
    - All with NVIDIA Quantum-2 400Gb/s InfiniBand networking (3.2TB/s bandwidth), no data ingress or egress fees, WEKA.IO storage.
    - Total fleet: ~22,200 NVIDIA GPUs (including 7.1k B300, 4.2k B200), expanded to 150,000 GPUs (March 2026). Prince George alone has >23,000 GPUs installed or on-order.

- **GPU Cloud Buyer's Guide** – A downloadable guide on cost, networking, SLAs, and scaling for GPU cloud procurement (2026 edition).

Solutions

- **AI Cloud** – On-demand / reserved GPU clusters for AI training and inference workloads. Designed for AI workloads including fine-tuning LLMs (Llama, Mistral, Falcon, DeepSeek, StabilityAI), deep learning, recommender systems, anomaly detection, digital twin simulations, and multimodal/diffusion/transformer model training. Supports frameworks like TensorFlow, PyTorch, JAX, and containers (Docker, Apptainer).
- **Colocation** – Turnkey colocation data center solutions. Rack-ready, AI-optimized space with air and liquid cooling, 24/7 on-site support, multi-layered physical security, backup generators/battery backup, and multi-path network connectivity. Available at four operational locations. Both retail and wholesale colocation.
- **Build-To-Suit** – Custom-designed and built data centers on IREN-owned land with grid-connected power. 2.75 GW of locked-in, grid-connected power across Texas campuses. Includes site master planning, Tier 1 engineering partnerships, modular expansion capability (1 MW to 50+ MW), and leverages existing IREN substations and high-voltage infrastructure. ~US$0.05/kWh all-in power cost in West Texas (Source: Aurora Energy Research, Q4 2024).

Services

- **24/7 Customer Support** – On-site, in-house team for critical workload support across all solutions (AI Cloud, Colocation, Build-To-Suit).
- **Proactive Support** – Performance monitoring and tracking for GPU cloud workloads.
- **On-site Engineering Access** – Direct access to engineers experienced with AI infrastructure.
- **Remote Operations Command Center** – Centralized monitoring with real-time response.
- **Community Grants Program** – Annual grants of up to $100,000 per community (Mackenzie, Prince George, Childress, Fisher County), with $500k+ USD total granted in 2024 across 100+ recipients. Maximum $10,000 per individual application. The program has been running since 2022.
- **Data Center Design, Development & Construction Services** (Build-To-Suit) – End-to-end project delivery including expert advice, Tier 1 partner coordination, mobilized local labor force, custom design, and on-site support.
- **IREN Cloud™ Reserved Instances** – Dedicated H100, H200, or Blackwell GPU nodes.
- **Investor Relations** – IR contact at ir@iren.com; media enquiries at media@iren.com;

general info at info@iren.com.

Marketing Analysis

Source: iren.md

Marketing – Strategic & Interpretive: https://iren.com

Analyse a company's website for strategic positioning, brand identity, value propositions, and competitive strategy. These topics require interpreting tone, patterns, and implied intent rather than extracting explicit facts.

Based on 103 stored pages analysed in 1 batch(es) with 1 AI calls (namespace: webscrape, model: claude-opus-4-6, profile: marketing_analysis).

Positioning Statements

- **Primary tagline/hero copy:** "AI Cloud for the next big thing"
- **Sub-tagline:** "Next-Gen Data Centers for AI, HPC & Sustainable Compute" (site title)
- **Repeated positioning copy:** "Launch sooner. Scale faster." and "Get your AI to market faster with IREN Cloud"
- **Key differentiator statement:** "The IREN difference — Infrastructure engineered not just for what AI needs now, but for what AI will demand next."
- **"We are X" statements:**
    - "IREN (NASDAQ: IREN) is a leading AI Cloud Service Provider, delivering large-scale GPU clusters for AI training and inference."
    - "IREN is a leading next-generation data center business powering the future with 100% renewable energy."
- **Competitive differentiators explicitly stated:**
    - "Vertical integration" — repeatedly emphasized: "IREN's vertical integration, expertise and flexibility empowers you to scale faster, accelerate time to market and gain your competitive edge."
    - "Owned and Operated" — "With our end-to-end control, you gain operational efficiency and reliability, backed by 24/7 dedicated support."
    - "Trusted AI Infrastructure" — "Purpose-built for demanding AI training and inference by the experts in high-performance infrastructure."
    - "Tomorrow-Ready" — "Future-proof your growth with infrastructure located at large-scale sites in emerging AI hubs."
    - 100% renewable energy (*from clean or renewable energy sources or through the purchase of RECs)
    - No data ingress or egress fees
    - NVIDIA Preferred Cloud Partner status
    - Built to NVIDIA reference architecture

- **Implicit category positioning:** Positioned against general-purpose/legacy cloud providers and hyperscalers: "Legacy cloud infrastructure was built for web apps, SaaS, and distributed storage — not multi-GPU training, high-throughput inference, or massive model checkpoints."
- **Fourth Industrial Revolution framing (Daniel Roberts quote):** "We're living through the Fourth Industrial Revolution... Inside these new AI factories, the 'workers' aren't people on assembly lines – they're GPUs, clustered together, never clocking off, never tiring, able to scale to millions."

Overall Brand Identity

- **Core brand values:**
    - Sustainability: 100% renewable energy commitment, ESG front and center, community grants program ($500k+ USD distributed, 100+ recipients in 2024)
    - Community responsibility: "We believe that human progress is invaluable, but it should be done in the right way – responsibly, sustainably, and positively"
    - Innovation and future-readiness: "Tomorrow-Ready" and "next-generation" are recurring themes
    - Transparency and trust: Emphasized in pricing ("exclusive pricing," "no data ingress or egress fees"), operations, and customer relationships
    - Speed and execution: "moving fast with intention"
- **Culture and ethos:**
    - "Innovation starts with people" (careers tagline)
    - "We work together with curiosity and purpose, and we give people the space to do meaningful work that creates real impact."
    - "A culture built on learning, responsibility, and purpose"
    - Employee testimonials reinforce themes of growth, community impact, collaboration, and accountability
    - Co-founded by Daniel Roberts and Will Roberts (Co-CEOs), suggesting a founder-led culture
- **Tone of voice patterns:**
    - Confident but not aggressive — authoritative technical language balanced with community warmth
    - Forward-looking and aspirational ("the next big thing," "powering the future")
    - Engineering-centric credibility (detailed specs, PUE ratings, power architectures)
    - Sustainability as a genuine thread, not an afterthought
- **Brand heritage note:** Formerly "Iris Energy" — transitioned to IREN, signaling a rebrand aligned with the pivot from Bitcoin mining toward AI cloud services. The Bitcoin mining history is acknowledged but de-emphasized in current positioning ("IREN pausing Bitcoin expansion to focus on AI data center build-out").

Value Proposition Statements

- **Core value proposition:** "Whether you need immediate access to GPUs, world-class turnkey data centers for your own hardware, or a partner to design and build your AI-ready facility, IREN provides the expertise, infrastructure and flexibility to accelerate your AI journey."

- **For AI Cloud specifically:**
  - "Get your AI to market faster with IREN Cloud — Launch sooner. Scale faster."
  - "Train powerful AI models and run inference at scale, fast. Owned and operated by IREN, our purpose-built data centers and GPU cloud solutions deliver the high-performance compute your business demands, with the rapid scalability and control you need."
  - "Deal directly with IREN's fully integrated data centers to access exclusive pricing and customer service tailored to your GPU cloud needs."
  - "No data ingress or egress fees" — explicit cost advantage claim
  - "IREN Cloud™ is built to NVIDIA reference architecture with full non-blocking 3.2TB/s InfiniBand connectivity."
- **Quantified scale/outcomes:**
  - 4.5 GW of power secured across 6 locations
  - 810 MW operational, 2,100 MW under construction, 1,600 MW in development
  - 4,500 acres total area across North America
  - $9.7 billion agreement with Microsoft for AI cloud infrastructure deployment
  - ~$1.94 billion in annualized run-rate revenue expected from Microsoft contract
  - 200MW of critical IT load in 4 phases through 2026 (Microsoft deal)
  - Over 23,000 GPUs installed or on-order (Prince George)
  - Expanded fleet to over 150,000 GPUs (Mar 2026 news)
  - Sweetwater: capacity for more than 700,000 liquid-cooled Blackwell GPUs
  - PUE (avg.) of 1.1 at Canadian facilities
  - AI projected to contribute $15.7 trillion to global economy by 2030 (cited)
  - ~US$0.05/kWh all-in power cost in West Texas (Source: Aurora Energy Research, Q4 2024)
  - 20% prepayment included in Microsoft contract
  - $5.8 billion Dell Technologies GPU/equipment purchase
  - $2.3 billion convertible notes offering closed (Dec 2025)
  - Oklahoma site: <6ms latency to network hubs
  - Sweetwater: <10ms latency to nearest hyperscaler regions
- **For Colocation:** "Access our world-class data center campuses — four premier locations with immediate, scalable capacity... reduces capital expenditures, speeds time-to-market, and eliminates the complexities of managing an in-house data center."
- **For Build-To-Suit:** "2.75GW of locked-in, grid-connected power in West Texas" with "Low-Cost Power Markets ~US$0.05/kWh all-in power cost."
- **GPU Cloud Buyer's Guide value props (key takeaways):**
  - The Provider Ecosystem: comparing specialized vs general-purpose
  - Networking Architecture: why high-throughput NVIDIA Quantum InfiniBand is critical
  - Cost Optimization: evaluating pricing and avoiding hidden fees
  - Owner-Operator Advantage: benefits of direct data center ownership
  - Support Ecosystems: 24/7 specialist access

Ideal Customer Profile

Explicit customer segments from the website:

1. **AI Cloud / GPU Cloud customers:** Organizations needing GPU clusters for AI training and inference — companies building LLMs, fine-tuning models (Llama, Mistral, Falcon, DeepSeek), deep learning, digital twins, recommender systems, anomaly detection. Use cases explicitly listed include: "Fine-Tuning Models," "Training and Fine-Tuning Various Models" (language, action, multimodal, diffusion, transformer, foundation, frontier/world models), "Recommender Systems," "Anomaly Detection," "Digital Twin Technology."

2. **Colocation customers:** Businesses that want to deploy their own hardware in IREN's AI-optimized facilities. "Ideal for cloud providers, financial services and enterprise platforms" (wholesale colocation). Contact form options: GPU Cloud, AI Data Centers | Colocation, AI Data Centers | Build-To-Suit.

3. **Build-To-Suit customers:** Large enterprises/hyperscalers needing custom-built data centers at massive scale. Industries cited: "AI companies, government agencies and financial services." The $9.7B Microsoft deal validates hyperscaler customers.

4. **Hyperscalers as customers:** The Microsoft partnership explicitly demonstrates IREN targeting global hyperscalers. Daniel Roberts stated it "opens access to a new customer segment among global hyperscalers."

Match to provided personas:

- **Carter (CTO/Technical Buyer) — STRONG MATCH:** IREN's messaging around GPU capacity, performance optimization, infrastructure scaling, guaranteed capacity, and NVIDIA reference architecture directly addresses Carter's top buying triggers (performance optimization, GPU capacity) and concerns (vendor lock-in mitigated by open frameworks like PyTorch/TensorFlow/JAX, GPU cost volatility addressed by "exclusive pricing" and "no egress fees"). The technical depth of GPU specs (H100, H200, B200, B300, GB300 NVL72) with detailed specs speaks to his evaluation process.

- **Alice (AI Platform Leader) — STRONG MATCH:** IREN addresses her core challenges of compute/GPU resource scarcity, cost management, and scaling. The vertical integration, 24/7 support, and "owned and operated" messaging directly target her need for operational efficiency and reliability. The emphasis on "no data ingress or egress fees" and "exclusive pricing" addresses cost efficiency with visibility. The NVIDIA Preferred Partner status and reference architecture alignment address her evaluation of new tech vendors.

- **Cedric (CEO) — MODERATE-STRONG MATCH:** The $9.7B Microsoft deal signals enterprise credibility and board-ready validation. IREN's sustainability positioning (100% renewable energy) and community grants address ESG requirements. The scale messaging (>4.5 GW, 6 locations) provides the infrastructure certainty CEOs need for mission-critical launches.

- **Diana (AI Leader) — MODERATE MATCH:** IREN addresses her compute capacity constraints and cost predictability concerns. The bare metal GPU access with root access and framework flexibility (TensorFlow, PyTorch, JAX) supports her R&D needs. However, IREN doesn't offer software-layer tools (MLOps, monitoring, governance) that

Diana also evaluates.

- **Malik (AI Researcher) — MODERATE MATCH:** IREN's infrastructure supports his compute-intensive training needs and large-scale experimentation. The support for "bring your own data, models, and AI/ML frameworks" aligns with his workflow. However, IREN is pure infrastructure — it lacks the developer tooling layer (experiment tracking, model lineage) Malik values.

- **Paul (Platform Engineer) — MODERATE MATCH:** IREN addresses his cost management challenges and scalability constraints. The bare metal access, InfiniBand networking, and detailed infrastructure specs (PUE, rack density) speak to his infrastructure concerns. Open-source compatibility (Docker, Apptainer, PyTorch, TensorFlow, JAX) addresses his vendor lock-in concerns.

- **Sonia (AI Engineer) — WEAK-MODERATE MATCH:** IREN provides the underlying compute infrastructure but doesn't directly address Sonia's application-layer needs (prompt engineering, agent flows, RAG pipelines, monitoring). She would be a downstream user of IREN infrastructure rather than a direct buyer.

Geographic focus: North America — specifically British Columbia (Canada) and Texas/Oklahoma (USA). Corporate presence in Sydney, Australia. All data centers in North America.

Inferred company size of target customers: Large enterprises and hyperscalers primarily (evidenced by the Microsoft deal, wholesale colocation focus, build-to-suit offerings starting at multi-MW scale). AI-native companies also targeted through GPU Cloud offering (smaller entry point with reserved instances).

Implied Marketing Strategy

How IREN intends to win:

IREN's implied strategy is to become the dominant **vertically integrated AI infrastructure provider** at the physical layer — owning the full stack from land and power to data centers to GPUs — and competing on **guaranteed capacity, speed-to-deployment, cost transparency, and sustainability** rather than on software/platform capabilities.

Competitive moat being built:

1. **Power & land lock-up:** IREN has secured >4.5 GW of grid-connected power across 4,500 acres at 6 locations. This is an enormous physical moat — power capacity is the single most constrained resource in AI infrastructure today. By locking in grid connections, substations, and freehold land ownership early, IREN creates a barrier that takes competitors years to replicate.

2. **Vertical integration:** IREN repeatedly emphasizes it "builds, owns, and operates" everything — land, substations, data centers, and GPU hardware. This allows them to control costs, quality, and speed in ways that companies leasing third-party facilities cannot. This is explicitly contrasted against competitors: "Is the provider using 3rd party data center facilities and 3rd party support staff?"

3. **Renewable energy positioning in low-cost power markets:** By locating in renewable-rich regions (BC Hydro in Canada, ERCOT wind/solar corridor in West

Texas), IREN secures ~$0.05/kWh all-in power costs while maintaining a 100% renewable energy claim. This serves dual purposes: cost advantage and ESG compliance for enterprise customers.

4. **NVIDIA strategic partnership:** Being an NVIDIA Preferred Cloud Partner with deployments across H100, H200, B200, B300, and GB300 NVL72 creates a technology credibility moat and likely preferential access to hardware allocation.

Market dynamics being exploited:

1. **Power scarcity as the bottleneck:** IREN is fundamentally betting that access to large-scale, grid-connected power is the primary constraint in AI infrastructure growth. Their strategy of securing GW-scale power in advance of demand is a land-grab play.

2. **Shift from Bitcoin mining to AI:** IREN's origin as "Iris Energy" (Bitcoin mining) gave them early mover advantage in securing cheap power and building data centers. They're now pivoting that infrastructure to the higher-margin AI compute market — a classic asset arbitrage strategy.

3. **Hyperscaler overflow:** Major cloud providers (AWS, Azure, GCP) face their own capacity constraints. IREN is positioning as a supplier *to* hyperscalers (evidenced by the $9.7B Microsoft deal) rather than competing *against* them.

4. **Build-to-suit as a wedge:** Offering BTS data centers at massive scale with pre-secured power creates sticky, long-term contracts with hyperscalers and large enterprises.

Trends being ridden:

1. **Explosive AI compute demand:** Explicitly cited ($15.7T economic contribution by 2030, 92% of companies investing in AI).
2. **GPU infrastructure scarcity:** Addressed through massive fleet expansion (150,000+ GPUs) and GW-scale power.
3. **Liquid cooling transition:** IREN is investing heavily in direct-to-chip liquid cooling for NVIDIA Blackwell/Blackwell Ultra, positioning for the density requirements of next-gen GPUs.
4. **Enterprise sustainability mandates:** 100% renewable energy claim provides ESG-ready infrastructure.
5. **Cloud-first enterprise AI adoption:** Citing Google's finding that 74% of organizations prefer hybrid cloud.
6. **AI factory-scale computing:** Aligning with NVIDIA's vision of Gigawatt-class AI factories.

Marketing approach pattern:

- **Top of funnel:** Thought leadership content (blog posts on Diffusion LLMs, Llama 4 requirements, AI infrastructure trends), industry event presence (NVIDIA GTC, All-In Summit, World Summit AI, Yotta, RAISE Summit), executive media appearances (Daniel Roberts on theCUBE, Financial Post, TIME)
- **Mid-funnel:** GPU Cloud Buyer's Guide (gated content), data center tours, technical blog posts with infrastructure requirements
- **Bottom of funnel:** "Contact Sales" / "Talk to us" CTAs throughout, sales-led motion (no self-serve pricing published — all "Contact sales")
- **Investor relations as marketing:** Extensive IR presence (NASDAQ: IREN, 13 analyst firms covering the stock, $41.94 share price, added to MSCI USA Index) signals financial stability and legitimacy to enterprise buyers

- **Partnership marketing:** Microsoft ($9.7B deal), Dell Technologies (success story), NVIDIA Preferred Partner status — all serve as enterprise social proof
- **Community/ESG marketing:** Community grants program serves dual purpose of local license-to-operate and corporate brand building

Pricing strategy implied: Enterprise/sales-led with no published pricing. Emphasis on "exclusive pricing," "transparent pricing," "no data ingress or egress fees," and "cost-effective predictable pricing" suggests IREN competes on total cost of ownership transparency rather than unit price competitiveness with hyperscalers. The "contact sales" model for all GPU tiers (H100 through GB300 NVL72) indicates reserved instance/contract-based pricing rather than on-demand.

Operational Metrics

Source: iren.md

Marketing – Factual & Quantitative: https://iren.com

Extract concrete, verifiable facts from a company's website: corporate details, infrastructure specs, customer names, product listings, and quantitative claims.

Based on 103 stored pages analysed in 1 batch(es) with 1 AI calls (namespace: webscrape, model: claude-opus-4-6, profile: operational_metrics).

Number Of Data Centers, Their Locations, Gpus, And Power Capacity

IREN operates **6 data center locations across North America** with a total secured power of >**4.5 GW** across **4,500 acres**:

Operational/Existing:

- **Mackenzie** – British Columbia, Canada | 80 MW | 11 acres | Air-cooled | PUE avg. 1.1 | Max rack density 70 kW | Connected to BC Hydro grid
- **Canal Flats** – British Columbia, Canada | 30 MW | 10 acres | Air-cooled | PUE avg. 1.1 | Max rack density 70 kW | Connected to BC Hydro grid via 30MW on-site substation
- **Prince George** – British Columbia, Canada | 50 MW | 12 acres | Air-cooled | PUE avg. 1.1 | Max rack density 70 kW | >23,000 GPUs installed or on-order | Connected to BC Hydro grid
- **Childress** – Texas, USA | 750 MW | 576 acres | Air + Liquid Cooling | Connected to ERCOT grid

Under Construction:

- **Sweetwater** – Texas, USA | 2,000 MW total | >1,800 acres | Air + Liquid Cooling | Connected to ERCOT grid

- **Sweetwater 1**: 1,400 MW | 1,300 acres | Substation energization 2026
  - **Sweetwater 2**: 600 MW | 500 acres | Substation energization 2027
  - Capable of supporting **over 700,000 liquid-cooled Blackwell GPUs**

In Development:

- **Oklahoma** – Oklahoma, USA | 1,600 MW | 2,000 acres | Substation energization 2028 | <6ms latency to network hubs

GPU types deployed: NVIDIA H100, H200, B200, B300, and GB300 NVL72 systems. Total NVIDIA GPU fleet approximately **22,200** as of the Blackwell deployment blog post, with capacity expanded to **over 150,000 GPUs** (per March 4, 2026 news release: "IREN Expands AI Cloud Capacity to 150,000 GPUs"). At Prince George, a new **10 MW (IT load) liquid-cooled data center** is under construction to support NVIDIA GB300 NVL72 deployments capable of handling **more than 4,500 GPUs**.

Status summary:

- 810 MW Operational
- 2,100 MW Under Construction
- 1,600 MW In Development

Active And Contracted Megawatts, Terawatts, And Other Measures Of Power

- **Total secured power:** >**4.5 GW** (also referenced as "nearly 3 GW of grid-connected power" in some contexts and "3GW secured power portfolio" in others)
- **810 MW Operational**
- **2,100 MW Under Construction**
- **1,600 MW In Development**
- **Over 600 MW of data centers energized** (referenced in the 50 EH/s milestone, achieved in 18 months)
- **510 MW of operating data centers** (referenced at 31 EH/s milestone)
- **Build-To-Suit: 2.75 GW** of locked-in, grid-connected power in West Texas; "750MW energized today and ~2GW energized from 2026-2027"
- **Microsoft contract: 200 MW of critical IT load** across 4 phases (Horizon 1-4) at the 750 MW Childress campus
- **Energy: 100% renewable** (from clean or renewable energy sources or through the purchase of RECs)
- **Power cost: ~US$0.05/kWh all-in** in West Texas (Source: Aurora Energy Research, Q4 2024)
- Sweetwater 1: 1,400 MW bulk substation (345kV/138kV) under construction
- Sweetwater 2: 600 MW, substation energization 2027
- Oklahoma: 1,600 MW, substation energization 2028

Installed, Active And Future Number Of Gpus

- **Total NVIDIA GPU fleet: approximately 22,200** spanning NVIDIA H100, H200, B200, B300, and GB300 NVL72 systems (as of September 2025 blog post)

- Most recent expansion included **7,100 NVIDIA HGX B300 GPUs** and **4,200 NVIDIA HGX B200 GPUs**
- >**23,000 GPUs installed or on-order** at Prince George specifically
- **AI Cloud capacity expanded to 150,000 GPUs** (per March 4, 2026 investor news release)
- Prince George: new 10 MW liquid-cooled data center under construction capable of housing >**4,500 GB300 NVL72 GPUs**, with the campus capable of housing **over 20,000 NVIDIA Blackwell and Blackwell Ultra GPUs**
- Sweetwater: capacity for **over 700,000 liquid-cooled Blackwell GPUs**
- Microsoft agreement: GPU clusters deployed in **4 phases through 2026** (Horizon 1-4) using **NVIDIA GB300 GPUs**, providing 200 MW of critical IT load at Childress
- Dell Technologies agreement: approximately **$5.8 billion** for GPUs and ancillary equipment related to the Microsoft deployment

---

Teraflops And Exaflops

- **IREN achieved 50 EH/s** (exahashes per second) of self-mining Bitcoin capacity — "In just 18 months, we've energized over 600MW of data centers and scaled from 5.6 to 50 EH/s"
- **31 EH/s** previously achieved at 15 J/TH efficiency with 510 MW of operating data centers
- **NVIDIA InfiniBand networking: 3.2 TB/s bandwidth** (NVIDIA Quantum-2 400Gb/s InfiniBand)
- No specific teraflops or exaflops compute performance figures (TFLOPS/PFLOPS/EFLOPS) were found on the website.

Other Capacity And Performance

- **PUE (Power Usage Effectiveness):** Average 1.1 at Mackenzie, Canal Flats, and Prince George
- **Max rack density:** 70 kW at Mackenzie, Canal Flats, and Prince George; power draw per rack can exceed 100 kW for AI workloads
- **Networking:** NVIDIA Quantum-2 400Gb/s InfiniBand; full non-blocking 3.2 TB/s InfiniBand connectivity; NVLink interconnect up to 3.2 TB/s throughput
- **Latency:** Oklahoma site <6ms to network hubs; Sweetwater <10ms to nearest hyperscaler regions
- **Microsoft deal value:** approximately **$9.7 billion**, five-year contract with 20% prepayment; expected **~$1.94 billion in annualized run-rate revenue** once fully commissioned
- **$2.3 billion convertible notes offering** closed (December 2025); **$2 billion convertible notes** priced (December 2025)
- **GPU Cloud specs offered:**
  - H100: 80GB GPU memory, 2,048 GB RAM, 224 vCPUs, 30TB NVMe
  - H200: 141GB GPU memory, 2,048 GB RAM, 224 vCPUs, 30TB NVMe
  - B200: 180GB GPU memory, 2,048 GB RAM, 224 vCPUs, 15.4TB NVMe
  - B300: 288GB GPU memory, 3,096 GB RAM, 256 vCPUs, 30TB NVMe
  - GB300 NVL72: 288GB GPU memory, 20,736 GB RAM, 2,592 vCPUs, 276TB NVMe

- **Bitcoin mining efficiency:** 15 J/TH
- **IREN is a Preferred NVIDIA Cloud Partner**
- **Cooling types:** Air-cooled (BC sites), Air + Liquid cooling (Texas sites), direct-to-chip liquid cooling (Horizon 1 at Childress)
- **Stock:** NASDAQ: IREN, added to MSCI USA Index (Feb 2026)
- **Fiber:** Dual, physically diverse fiber paths at all facilities
- **Total area:** 4,500 acres across all 6 locations
- **25 employees** at Mackenzie site specifically mentioned