

Consolidated Report: lambda

Generated 2026-03-29 04:11 UTC from 3 analysis profile(s): company_profile, marketing_analysis, operational_metrics.

Note: The primary source of information in this report is the company's own website. The content reflects the company's self-reported claims and marketing materials and should not be treated as independently verified fact.

Executive Summary

Executive Summary: Lambda, Inc.

Company Overview

Lambda, Inc. is a San Francisco-based AI infrastructure company founded in 2012 by published machine learning engineers Stephen Balaban (CEO) and Michael Balaban (CTO). Originally born out of Noisebridge, an anarchist hackerspace in the Mission District, Lambda has evolved over 14 years into what leadership describes as an "emerging hyperscaler" — a pure-play, AI-only cloud platform providing GPU compute infrastructure for training and inference workloads. The company brands itself as "The Superintelligence Cloud" and serves over 200,000 AI developers, 97% of top U.S. research universities, and marquee customers including Microsoft, which has signed a multibillion-dollar infrastructure agreement. Lambda has raised over \$2.3 billion across funding rounds, including a \$1.5B+ Series E in November 2025 led by TWG Global and the US Innovative Technology Fund.

Key Findings

Company Profile & Product Portfolio. Lambda offers a three-tiered product ladder — On-Demand Instances (1–8 GPUs), 1-Click Clusters (16–2,000+ GPUs), and Superclusters (4,000–165,000+ GPUs) — creating a natural land-and-expand revenue model spanning individual researchers to hyperscalers. The company is an early launch partner for every recent NVIDIA GPU generation, including GB300 NVL72, HGX B300, and the forthcoming Vera Rubin platform. Notably, Lambda exited its on-premise hardware business in August 2025, signaling a decisive strategic pivot to higher-margin, recurring-revenue cloud services. Its software ecosystem (Lambda Stack, Cloud API, Filesystem S3 Adapter) and embedded co-engineering services differentiate it from self-service-only competitors.

Strategic Positioning & Brand Identity. Lambda occupies a distinctive niche between hyperscalers (AWS, Azure, GCP) and commodity GPU clouds, competing on AI-native specialization, transparent pricing (no ingress/egress fees), single-tenant security architecture, and deep NVIDIA alignment — underscored by seven consecutive NVIDIA Partner of the Year awards. The anti-hyperscaler positioning is deliberate and systematic: Lambda attacks incumbents on cost opacity, vendor lock-in, shared tenancy, and generic infrastructure. Its research credibility — publishing at NeurIPS, ICML, and other top venues — generates demand among the practitioner communities that influence enterprise purchasing decisions.

Operational Scale & Infrastructure. Lambda is executing an aggressive physical buildout toward a stated vision of 3GW+ of AI data center capacity and over one million GPUs by decade's end. Current and planned facilities span Texas, California, Ohio, Illinois, Georgia, Missouri, and Virginia, with the company's megawatt footprint expected to quadruple between 2025 and 2026. Performance metrics are impressive: Model FLOPs Utilization exceeding 60% (versus a 35–45% industry norm), MLPerf Training results showing 27% gains on GB300 versus GB200, and customer-reported outcomes including 3× faster training cycles and 2× R&D turnaround improvements.

Cross-Cutting Themes, Strengths, and Risks

Lambda's greatest strength is the mutually reinforcing relationship between its NVIDIA partnership, technical credibility, and scale ambitions — creating supply-side advantages that are difficult to replicate. The founder-led, engineering-first culture and transparent pricing philosophy resonate strongly with technically sophisticated buyers. However, several risks warrant attention: the company's heavy dependence on NVIDIA creates concentration risk; the capital-intensive data center buildout requires sustained access to financing; competition from CoreWeave, hyperscaler AI-specific offerings, and emerging GPU clouds is intensifying; and the pivot away from hardware narrows revenue diversification. Additionally, Lambda's predominantly U.S.-centric footprint may limit international opportunity capture as sovereign AI compute demand grows globally.

Strategic Outlook

Lambda is well-positioned to capitalize on secular tailwinds — GPU scarcity, enterprise AI adoption, open-source model proliferation, and growing hyperscaler fatigue — but execution risk is significant given the magnitude of its infrastructure ambitions. The Microsoft partnership and \$2.3B+ in cumulative funding provide meaningful runway, and the company's progression from hackerspace startup to emerging hyperscaler demonstrates strategic conviction. The critical near-term question is whether Lambda can scale its operational capacity fast enough to meet contracted demand while maintaining the performance, reliability, and engineering culture that differentiate it. If it can, Lambda is on a credible trajectory to become a top-tier AI infrastructure provider; if buildout timelines slip or GPU supply dynamics shift, the capital-intensive model could prove unforgiving.

Company Profile

Source: lambda.md

Marketing – Factual & Quantitative: <https://lambda.ai>

Extract concrete, verifiable facts from a company's website: corporate details, infrastructure specs, customer names, product listings, and quantitative claims.

Based on 371 stored pages analysed in 1 batch(es) with 1 AI calls (namespace: webscrape, model: claude-opus-4-6, profile: company_profile).

Headquarters

Lambda was founded in San Francisco in 2012. The About page states: "Lambda was founded in San Francisco in 2012. Our first 'office' was Noisebridge, an anarchist hackerspace in the Mission District." The company is incorporated as "Lambda, Inc." No specific current corporate address is provided beyond San Francisco.

Years In Business

Lambda was founded in **2012**. Multiple pages confirm this: "Lambda is an AI-only company founded by ML engineers in 2012," "Founded in 2012 by published AI engineers," and "Lambda: enabling the AI community & innovating since 2012." As of the website content (dated 2026), that represents approximately 14 years in business.

Company Executive Names And Titles

Founders:

- **Stephen Balaban** — Co-founder and CEO
- **Michael Balaban** — Co-founder and CTO

Executive Team:

- **Robert Brooks IV** — Chief Commercial Officer (also described as "Founding Team & VP, Revenue" in earlier content)
- **David Connolly** — Chief Legal Officer
- **Larry Getman** — EVP Operations
- **Jerry Hunter** — Vice Chairman, Compute Delivery and Special Advisor to the Board. Former AWS infrastructure leader and Snap COO; brings 30 years of hyperscale expertise.
- **Ariel Nissan** — General Counsel
- **Ken Patchett** — VP Infrastructure (VP of Data Center Infrastructure)
- **Collin Roe-Raymond** — Chief Design Officer
- **Leonard Speiser** — Chief Operating Officer. Former co-founder and CEO of Clover (processes over \$300B in credit card transactions/year).
- **Matthew Worby** — Chief People Officer
- **Paul Zhao** — Head of Product
- **Charles Fisher** — Chief Financial Officer. Previously CFO at Turo; EVP of Corporate Finance & Development at Charter Communications.

Board/Other Senior Leaders:

- **Heather Planishek** — Previously appointed CFO (December 2025), then succeeded by Charles Fisher. Previously Board member and Audit Chair. Former Chief Operating and Financial Officer at Tines; former Chief Accounting Officer at Palantir Technologies.
- **Stacey Finerman** — VP, Investor Relations. Previously VP IR at Marqeta, PagerDuty, Square (Block), Zayo Group.
- **Mitesh Agrawal** — Chief Operating Officer and Head of Cloud (earlier role; described as COO in some content)

- **Peter Seibold** — Chief Financial Officer (mentioned in October 2025 context)
- **David Hall** — VP of NVIDIA Solutions
- **Maxx Garrison** — Product Manager for Cloud Infrastructure
- **Kenneth "Ken" Patchett** — VP of Data Center Infrastructure (built Microsoft's first data center in 1988 as an ironworker; previously worked at Microsoft, Google, Oracle, Amazon)

Blog Authors/Team Members mentioned: Chuan Li (Chief Science Officer/ML Research), Khushboo Goel, Anket Sah, Thomas Bordes, Amir Zadeh, Sam Khosroshahi, Lea Alcantara, Jessica Nicholson, Cody Brownstein, Kathy Bui, Nick Harvey, Corey Lowman, Xi Tian, Zach Mueller, and others.

Customer Names, Locations, Industry/Type Of Entity, And Other Relevant Details

Named Customers:

- **Microsoft** — Multibillion-dollar agreement to deploy AI infrastructure powered by tens of thousands of NVIDIA GPUs under a multi-year contract. Lambda has worked with Microsoft for more than 8 years.
- **Iambic Therapeutics** — Clinical-stage biotech/drug discovery. Uses Lambda for AI-driven drug discovery, training NeuralPLexer3 and Enchant models. Selected Lambda NVIDIA HGX B200 cluster. Quote: "The scale and stability of the Lambda cluster, accelerated by NVIDIA, helped enable us to train NeuralPLexer3." — Matt Welborn, PhD., VP of Machine Learning.
- **Genesis Therapeutics** — Drug discovery, targeting "undruggable" proteins. Uses Lambda 1-Click Clusters. Achieved 3x faster model training. Quote: "This field is so complex that AI might understand it better than humans." — Wojtek Swiderski, Principal Software Engineer.
- **Pika** — AI video generation startup. Uses Lambda Reserved Clusters with NVIDIA H100 GPUs.
- **fal** — Generative AI, trained AuraFlow model on Lambda 1-Click Clusters. Quote: "I saw 1CC... and I was like, this is an amazing idea." — Batuhan Taskaya, Founding Engineer.
- **Meshy** — 3D AI generative model company (#1 3D AI tool by traffic). Uses Lambda 1-Click Clusters for training Meshy-4 and Meshy-5. 2x R&D turnaround time improvement.
- **Synlico** — Biotech startup (founded 2021). Won Lambda Golden Ticket. Focuses on causal discovery in cell biology for oncology. Has 57M cell single-cell transcriptomics database.
- **ServiceNow** — Enterprise AI. Trained and deployed Apriel 5B model on Lambda.
- **Nous Research** — AI research. Trained Hermes 3 (first full-parameter fine-tune of Llama 3.1 405B) on Lambda 8-node 1-Click Cluster. Co-founders: Jeffrey Quesnelle, Teknum.
- **Allen Institute for AI (Ai2)** — Trained Olmo Hybrid 7B on 512 NVIDIA Blackwell GPUs on Lambda infrastructure (64 NVIDIA HGX B200 systems). Processed ~3 trillion tokens in 6.19 days.
- **Voltron Data** — Data company (Apache Arrow ecosystem). Used Lambda Reserved Cloud Clusters.
- **Abacus.ai** — Recorded nearly 2x improvement in latency using NVIDIA GH200 on Lambda.

- **Leonardo.ai** — Achieved over 3x throughput improvements using Lambda's NVIDIA GH200 Superchip cluster.
- **Ceramic AI** — Rethinking AI training infrastructure; benchmarked on Lambda's NVIDIA HGX B200 clusters. Founded by Anna Patterson (20+ years at Google and Stanford).

Customer Categories (unnamed):

- "AI Hyperscaler" — Built a cluster of thousands of GPUs in 90 days.
 - Hyperscalers, enterprises deploying AI in regulated industries, frontier labs building trillion-parameter foundation models.
 - Previous customer list mentioned: Sony, Intel, Microsoft, Amazon Research, Raytheon, Intuitive Surgical, MIT, Stanford, Harvard, GE, Department of Defense, Google, Tencent, Kaiser Permanente, Caltech, Anyscale, Rakuten, The AI Institute.
 - Over 200,000 AI developers and teams worldwide; 150,000+ Cloud users.
 - 97% of the top research universities in the U.S.
 - Lambda serves healthcare, finance, manufacturing, retail, logistics industries.
-

Partner Names, Locations, Industry/Type Of Entity, And Other Relevant Details

Major Technology Partners:

- **NVIDIA** — Core technology partner. Lambda is a seven-time NVIDIA Partner of the Year (2020-2025 consecutively), NVIDIA Exemplar Cloud, NVIDIA Cloud Partner, Platinum sponsor at GTC. Awards include: 2020 NPN Solution Integration Provider of the Year, 2021 Solutions Integration Partner of the Year, 2023 Solution Integration Partner of the Year, 2024 AI Excellence Partner of the Year, 2025 Healthcare Partner of the Year. Early launch partner for NVIDIA Vera Rubin, GB300 NVL72, HGX B200, GH200, H200, H100, Blackwell Ultra, Vera CPUs, BlueField-4 STX.
- **Supermicro** (NASDAQ: SMCI) — San Jose, CA. Hardware partner for AI-optimized servers including NVIDIA HGX B200 and GB300 NVL72 systems. Charles Liang, President and CEO.
- **Dell Technologies** — Listed as ecosystem partner for Superclusters.

Data Center Partners:

- **Cologix** — Denver, CO. Network-neutral interconnection and hyperscale edge data center company. Deployed NVIDIA HGX B200 clusters at COL4 in Columbus, Ohio. Chris Heinrich, Chief Revenue Officer.
- **EdgeConneX** — Herndon, VA. Building 30+ MW AI-enabled data center infrastructure in Chicago and Atlanta. Don MacNeil, Chief Revenue Officer.
- **Prime Data Centers** — International data center developer. Deployed at LAX01 in Vernon, CA (first AI-ready data center in Vernon). 33 MW critical power, 242,000 sq ft. Lambda initially leasing 21 MW.
- **ECL** — Mountain View, CA. Zero-water, zero-emissions off-grid hydrogen-powered data center. Lambda deployed first hydrogen-powered NVIDIA GB300 NVL72 systems. Yuval Bachar, Founder and CEO.
- **Equinix** — Mentioned as collaboration partner for AI factory expansion.

Software/Platform Partners:

- **dstack** — Open-source container orchestration for ML teams, natively integrated with Lambda.
- **SkyPilot** — Open-source orchestration tool for ML job deployment on Lambda Cloud.
- **Oumi** — End-to-end custom model development partner. Enables custom AI in hours, not months.
- **Baseten** — Partner for inference performance testing.
- **Ceramic AI** — Training infrastructure partner.
- **Hugging Face** — Partnered for Whisper Fine-Tuning Event.
- **Scale (Scale Nucleus)** — Data management partnership.
- **Razer** — Hardware partnership for Tensorbook laptop; also investor.
- **Safebase by DRATA** — Customer Trust Portal partner.
- **Open Compute Project (OCP)** — Lambda joined OCP Advisory Board.
- **MLCommons** — MLPerf benchmark participation.

Investors:

- **TWG Global** (led by Thomas Tull and Mark Walter) — Led \$1.5B+ Series E
 - **US Innovative Technology Fund (USIT)** — Led Series C (\$320M), participated in Series E
 - **Andra Capital** and **SGW** — Co-led Series D (\$480M)
 - Other investors: ARK Invest, NVIDIA, Supermicro, Pegatron, Wistron, Wiwynn, B Capital, SK Telecom, T. Rowe Price, In-Q-Tel (IQT), G Squared, 1517 Fund, Gradient Ventures, Bloomberg Beta, Crescent Cove, Mercato Partners, Andrej Karpathy, Greg Brockman, Adam D'Angelo, Garry Tan, Razer, Georges Harik, and others.
-

Products

Cloud Compute Products:

1. **Superclusters** — Single-tenant, shared-nothing AI cloud from 4,000 to 165,000+ NVIDIA GPUs. Purpose-built for large-scale training and inference. NVIDIA GB300 NVL72 and HGX B300 systems with NVIDIA Quantum-2 InfiniBand. Contract terms: 3-5 years. Features: physically isolated clusters, caged enclosures, SOC 2 Type II certified, managed Kubernetes or Slurm orchestration, Prometheus/Grafana observability.
2. **1-Click Clusters™** — Self-serve production-ready NVIDIA HGX B200 and H100 GPU clusters from 16 to 2,000+ GPUs. Fully optimized for distributed AI workloads with NVIDIA Quantum-2 InfiniBand networking (3.2 Tb/s aggregate). Commitment: 2 weeks to 3 years.
 - **Pricing:** NVIDIA HGX B200 On-Demand: \$4.62/GPU/hr; NVIDIA H100 On-Demand: \$2.76/GPU/hr. Reserved pricing (1-3 year) available via sales.
3. **Instances (On-Demand Cloud)** — 1 to 8 NVIDIA GPU virtual machine instances. Available GPUs: NVIDIA B200, H100 SXM, H100 PCIe, GH200, A100 SXM, A100 PCIe, A10, A6000, Tesla V100, Quadro RTX 6000.
 - **Pricing examples:** 8x B200 SXM6: \$5.74/GPU/hr; 1x GH200: \$1.99/hr; 8x H100 SXM: \$3.44/GPU/hr; 1x A10: \$0.86/hr; 1x RTX 6000: \$0.58/hr.
4. **Lambda Bare Metal Instances** — Announced for NVIDIA Vera Rubin NVL72 and GB300 NVL72. One-to-one mapping between instances and physical hosts, no hypervisor overhead.

Software Products:

1. **Lambda Stack** — Free deep learning software suite for Ubuntu. Pre-installed on all Lambda products. Includes: PyTorch 2.7.0, TensorFlow 2.19.0, Keras 3.10.0, JAX 0.6.0, CUDA 12.8.93, cuDNN, NCCL 2.26.2, NVIDIA Container Toolkit. Upgradable via apt.
2. **Lambda Cloud API** — API for managing instances, creating/stopping/restarting resources programmatically.
3. **Lambda Chat** — Chat interface hosting open-source models like DeepSeek-R1 and others.
4. **Inference API** (winding down) — Serverless inference API for open-source models including DeepSeek-R1-0528, DeepSeek V3-0324, Qwen3-32B. No rate limits. Pricing example: Qwen3-32B at \$0.10/M input tokens, \$0.30/M output tokens.
5. **Filesystem S3 Adapter** — S3-compatible API for Lambda Filesystem storage (GetObject, PutObject, DeleteObject, list). Eliminates need for VM to move data.
6. **Cloud Metrics Dashboard** — Real-time GPU/VRAM/CPU utilization monitoring powered by Guest Agent.
7. **LLM Index** — Catalog of model cards with architecture breakdowns, hardware requirements, deployment guides, and throughput benchmarks.
8. **GPU Benchmarks** — Interactive benchmark tool for GPU training/inference speeds using PyTorch/TensorFlow.
9. **LLM Performance Benchmarks Leaderboard** — Standardized comparison of leading LLMs.

Legacy Hardware Products (discontinued August 29, 2025):

- **Vector One** — Single-GPU desktop PC (NVIDIA GeForce RTX 4090/5090)
 - **Vector** — Dual-GPU desktop (2x NVIDIA GeForce RTX 5090, liquid-cooled)
 - **Vector Pro** — Multi-GPU workstation (up to 4x NVIDIA RTX PRO Blackwell GPUs, up to 96 GB GDDR7 ECC)
 - **Scalar** — GPU server (1-8x GPUs, PCIe connected)
 - **Hyperplane** — GPU server (4-16x NVIDIA GPUs with NVLink/NVSwitch, InfiniBand support)
 - **Tensorbook** — Deep learning laptop (collaboration with Razer)
 - **Echelon** — Turn-key GPU cluster
-

Solutions

By Mission/Vertical:

- **Superintelligence** — Purpose-built AI factories for frontier labs building trillion-parameter foundation models. Superclusters up to 165,000+ NVIDIA GPUs.
- **Enterprise** — "Scale fast. Stay secure." Accelerate AI training and inference with enterprise-grade security, SOC 2 Type II compliance, managed orchestration, predictable pricing, no ingress/egress fees.
- **Government** — CRADA enablement, AI professional services, secure commercial

clusters. Air-gapped Supercluster deployments for sensitive missions. SOC 2 Type II certified, U.S.-based data centers and personnel.

- **Startups and Researchers** — On-demand access from 1 GPU to clusters. Research grants up to \$5,000 in cloud credits. NVIDIA Inception/Connect programs with \$7,500 cloud credit offer.
- **Foundations/AI Model Builders** — Infrastructure for training and serving foundation models at global scale.

AI Infrastructure Solutions:

- **AI Factories** — Modular, interoperable data centers integrating high-density power, liquid cooling, and NVIDIA GPUs. Contiguous AI factories (75MW+ campuses) and distributed AI factories. Direct-to-chip liquid cooling, precision air cooling, N+1 to N+N redundancy.

Services

No relevant information found across any batch.

Marketing Analysis

Source: lambda.md

Marketing – Strategic & Interpretive: <https://lambda.ai>

Analyse a company's website for strategic positioning, brand identity, value propositions, and competitive strategy. These topics require interpreting tone, patterns, and implied intent rather than extracting explicit facts.

Based on 371 stored pages analysed in 2 batch(es) with 7 AI calls (namespace: webscrape, model: claude-opus-4-6, profile: marketing analysis).

Positioning Statements

Positioning Statements — Consolidated Summary

Primary Tagline & Hero Copy

- **"The Superintelligence Cloud"** — Lambda's core brand descriptor, used consistently across all pages and at flagship events (e.g., Platinum sponsor at GTC 2026)
- **"Supercomputers for training and inference" / "Supercomputers for**

- **superintelligence** — primary hero headings (slight variation across pages)
- **"Built for AI. Ready for superintelligence."** — recurring tagline across multiple pages
- **"Leadership for the age of superintelligence"** — leadership page framing
- **"Breakthroughs on demand"** — Instances page hero copy

Mission & "We Are" Statements

- **Stephen Balaban (CEO):** "Lambda's mission is to build the #1 AI compute platform in the world"
- **Ken Patchett (VP Data Center Infrastructure):** "I like to think about Lambda more as an emerging hyperscaler" — positions Lambda as a nascent hyperscaler, not a niche cloud
- **Brand statement:** "Lambda builds infrastructure for training and inference at superintelligence scale."
- **"Join the race to superintelligence"** — aspirational CTA framing Lambda as essential to the AI arms race

Category Creation Language

- **"AI factories"** — Lambda consistently frames its data centers/GPU clusters as purpose-built manufacturing facilities for intelligence, not commodity cloud instances or traditional data centers
 - "Yesterday's data centers were built for servers. Tomorrow's must be built for GPUs."
 - **"Gigawatt-scale AI factories"** — scale framing with "a vision to get to 3GW+ of AI data center space"

Competitive Differentiators

Pure-Play AI Focus

- "Lambda is an AI-only company founded by ML engineers in 2012. One hundred percent of our engineering, operations, and support is dedicated to AI workloads and infrastructure."
- **"Purpose-built AI cloud"** — differentiating explicitly from general-purpose providers (AWS, Azure, GCP)
- "Unlike general-purpose clouds, our systems are engineered for mission-critical AI"

Infrastructure Layer Positioning

- **"You bring models. We bring the compute."** — positions Lambda as infrastructure, not a model provider
- **"Your supercomputer. Your rules."** — autonomy/control positioning
- **"Single-tenant, shared-nothing AI cloud"** — positioning against multi-tenant hyperscaler environments
- **"Designed for builders"** — user identity positioning

Scalability

- "From one GPU to hundreds of thousands — performance that keeps pace."

- Superclusters: **"from 4,000 to 165,000 + NVIDIA GPUs"**
- **"Self-serve supercomputers"** (1-Click Clusters) — democratizing access to supercomputer-class infrastructure

Engineering Efficiency

- **"No FLOPS are left behind"** — engineering efficiency as a philosophical commitment
- **"Do more with every watt"** — efficiency positioning
- **MFU whitepaper:** "60%+ Peak MFU achieved vs. 35–45% industry norm" and "2.11x MFU uplift for Llama 70B on 16x NVIDIA HGX B200" — quantified performance superiority

Enterprise Readiness

- **"Secure by design. Mission-critical by default."**

Competitive Price-Performance Positioning

Pricing Philosophy

- **"Pay flat rates — no lock-in, no ingress/egress fees"** — directly attacking hyperscaler pricing complexity
- **"Predictable pricing"** — repeated across multiple pages

Specific Price Points & Comparisons

- 1-Click Cluster NVIDIA HGX B200: **\$4.62/hr** on-demand; H100: **\$2.76/hr**
- V100 on-prem vs. AWS: Lambda Hyperplane TCO **\$69,441 less** than AWS p3dn over 3 years
- **"2x better than comparable on-demand"** cloud instances for V100 8-GPU
- Lambda Cloud RTX 6000: **"roughly 4x the performance per dollar"** of a comparable AWS p3.8xlarge instance

NVIDIA Partnership Signal

- **Platinum sponsor at GTC 2026** — associating with NVIDIA at the highest tier, reinforcing credibility as a top-tier AI infrastructure provider
-

Overall Brand Identity

Overall Brand Identity — Consolidated Summary

Mission Statement

- **"Our mission is to make compute as ubiquitous as electricity and give everyone in America the power of superintelligence."**

- Democratization mantra: **"One person, one GPU."**
- The mission appears consistently across public-facing pages (About, Careers, Trust framework).

Brand Values

- **Customer obsession:** "Obsess over our customers' success. Work backward from the customer experience." Quotes Henry Ford (1922): "He is best served by the highest quality at the lowest price."
- **Engineering excellence / rigor:** Technical depth pervades all content — MFU benchmarks, interconnect architecture details, GPU specs, research publications. Tone is precise, quantitative, and confident without hyperbole.
- **Transparency:** "Clear, straightforward pricing," "Security you can verify," published benchmark methodologies on GitHub, no egress fees, transparent Trust Portal.
- **Co-engineering partnership:** "Experts included." "Lambda's ML, HPC, and infrastructure engineers embed with your team." Positioned as a hands-on collaborator, not just a vendor.
- **First principles thinking:** "Try to operate from timeless first principles."
- **Team brilliance:** "Build a team of brilliant people with an unstoppable culture." "Real recognize real — great people want to work with great people." "Every problem is a people problem. Every solution is a people solution."
- **Written communication:** "If you can't write your idea down into a one- or two-page document, you haven't thought about it enough."
- **Care deeply:** "Care about the company, care about the customer, care about the work you do."
- **Hacker culture:** Founded at Noisebridge hackerspace in San Francisco's Mission District. Ethos of "Do-ocracy" and "Be Excellent to Each Other."
- **Low ego:** "Focused on outcomes, not on titles. We value anonymous feedback. Own your mistakes."
- **Move fast:** "Usually, the only advantage that a small company has over a large company is speed."
- **Trust / Security:** Zero-trust security architecture, SOC 2 Type II certification. "Trusted by hyperscalers, global enterprises, and frontier labs operating in regulated environments." "Over a decade of experience."
- **Research commitment:** Funding research grants (up to \$5,000 in cloud credits); publishing at ICML, NeurIPS, ICCV, CoRL, EMNLP; CEO presenting "Neural Software" vision.
- **No politics, no crypto:** "We don't discuss cryptocurrencies or politics at work."

Tone of Voice

- **Technical authority paired with accessible, engineer-to-engineer directness.** Developer-native language throughout (CLI, API references, MFU metrics, code tutorials, deployment guides).
- **Confident but not arrogant** — described internally as "quiet confidence." Avoids marketing fluff; credibility established through practitioner-level detail and published benchmarks.
- **Aspirational framing** around "superintelligence" and "the race," balanced by grounded, quantitative evidence.
- **Cultural references** to hacker heritage and pop culture (Hackers 1992, Bill & Ted) alongside classic business thinking (Henry Ford, Bill Belichick).
- **AI-native aesthetic:** "Lambda Agent Terminal" JSON-LD manifests on product pages

create an intentionally futuristic feel, signaling the brand is built for an AI-first audience.

Culture & Origin

- **Founder-led, technical DNA:** Founded in 2012 by Stephen Balaban (CEO) and Michael Balaban (CTO), both published AI engineers (NeurIPS, ICCV, SPIE). "AI is in our DNA."
- **San Francisco origin story:** Started at Noisebridge hackerspace in the Mission District.
- **Decisive strategic conviction:** Ended on-premise hardware business (August 29, 2025) to go all-in on cloud.
- **Leadership style cues:** Titles use code-style formatting (e.g., `chief.executive.officer()`), reinforcing the engineering-first identity.

Ecosystem Positioning

- **Key partners prominently featured:** NVIDIA (deepest alignment — reference architectures, Vera Rubin readiness, Platinum GTC sponsor), Supermicro, Dell Technologies, VAST, Oumi, dstack, SkyPilot, SchedMD.
 - Positioned as the **engineer's cloud** — a technically credible, transparent, co-engineering partner for frontier AI labs, hyperscalers, and enterprises in regulated environments.
-

Value Proposition Statements

Consolidated Summary: Lambda Value Proposition Statements

Core Value Proposition

Lambda provides **purpose-built, production-ready GPU cloud infrastructure (AI factories)** that enables organizations to train, fine-tune, and serve AI models at scale — from single instances to **165,000+ GPU superclusters** — with transparent pricing, enterprise-grade security, and embedded AI engineering expertise.

Key taglines:

- "Get complete AI factories integrating high-density power, liquid cooling, and NVIDIA GPUs into one system designed for peak AI performance."
 - "Run large-scale AI workloads without the operational burden. We manage your clusters so you can focus on innovation."
 - "Co-engineer your workloads with the very people building the infrastructure behind the world's most advanced models."
 - "Accelerate AI training and inference on scalable infrastructure with enterprise-grade security, proven reliability, and predictable costs."
-

Quantified Performance Claims

Training & Inference Benchmarks

- **NVIDIA HGX B200:** Up to $3\times$ **faster training** and $15\times$ **faster inference** vs. prior generation (H100); $2\times$ the VRAM and FLOPS of H100 GPUs
- **NVIDIA GB300 NVL72:** Outperforms GB200 NVL72 by **27%** (MLPerf Training v5.1); $72\times$ Blackwell Ultra GPUs / $36\times$ Grace CPUs per rack, 37 TB fast memory / 130 TB/s NVLink Switch bandwidth
- **NVIDIA HGX B300:** 72 PF FP8 training / 144 PF FP4 inference, 2.1 TB HBM3e memory
- **MLPerf v5.1:** Up to **15.4% performance gains** over previous round's best results (Llama 3.1 405B Server scenario)
- **MFU (Model FLOPs Utilization):** Above **60%**, a **25%+ improvement over industry baseline** with no architectural changes; $2.11\times$ **MFU uplift** for Llama 70B on $16\times$ NVIDIA HGX B200
- **GH200 vs. H100 inference:** $7.6\times$ **better throughput** and $8\times$ **reduction in cost per token**
- **8B–405B parameter range** benchmarked on Llama 3.1 models with no architecture changes required

Customer/Partner Outcomes

- **Iambic Therapeutics:** $3\times$ **faster model training** — models that previously took months now complete in a week or less
- **fal / Meshy:** 1-Click Clusters sped R&D turnaround time by $2\times$, cutting months-long experimentation in half
- **Lambda + Oumi partnership:** Enterprises can build and deploy custom models $100\times$ **faster**, with $10\times$ **better cost efficiency**, and superior accuracy
- **Olmo Hybrid training:** **97–99% active training time**; median recovery time under **4 minutes**
- **FlashAttention-2 / GPT3-175B replication:** 242,400 GPU hours (H100 80GB SXM5) at \$458,136 using three-year reserved cluster — **90% cost savings**
- **Hyperscale deployment:** Thousands of GPUs delivered as a private cloud within **90 days**

Cost & Pricing Transparency

- "No hidden fees. No charges for data ingress or egress. Just clear, predictable pricing you can trust."
 - "Pay flat rates — no lock-in, no ingress/egress fees"
 - Up to **57% savings vs. AWS** for NVIDIA A100 bare metal hosted servers
 - **Specific published pricing:**
 - 1-Click Clusters: HGX B200 On-Demand **\$4.62/GPU/hr**; H100 On-Demand **\$2.76/GPU/hr**
 - Instances: B200 SXM6 $8\times$ from **\$5.74/GPU/hr**; H100 SXM $8\times$ from **\$3.44/GPU/hr**; A100 SXM 80GB from **\$2.06/GPU/hr**; Quadro RTX 6000 at **\$0.58/GPU/hr**
-

Key Benefit Statements

- **Speed:** "Accelerate time-to-market by removing infrastructure bottlenecks"; "Launch in minutes — no lengthy setup, no driver installs"; "Faster time to results"
 - **Security & Compliance:** "De-risk deployments with secure, isolated environments and built-in compliance"
 - **Efficiency:** "Optimize GPU utilization for peak performance and cost efficiency"; intelligent workload scheduling to minimize idle time
 - **Reliability:** "Ship with enterprise-grade reliability and managed services built for production AI"; "Ship with confidence from POCs to global rollout"
 - **Observability:** "Catch performance issues before they start"
 - **Cost Optimization:** "Reduce infrastructure costs by maximizing resource utilization"; pay-by-the-minute billing
-

Consultation & Partnership Value Props

- **Oumi partnership:** "Higher model accuracy, lower latency, better performance" via customized AI model consultation
 - **GTC workshop framing:** "Map a clear technical path from idea to ROI in 30 minutes"; "See how custom models deliver higher accuracy while keeping your data in your environment"; "Learn how dedicated GPU access turns iteration speed into competitive advantage"
 - **Healthcare vertical:** "Lambda Honored to Accelerate AI Innovation in Healthcare with NVIDIA"
-

Ideal Customer Profile

Consolidated Ideal Customer Profile for Lambda

Explicitly Named Target Segments

Lambda's navigation ("For every mission") defines five distinct segments:

1. **Superintelligence** — Frontier AI labs building trillion-parameter foundation models at extreme scale
 2. **Enterprise** — Large organizations deploying AI in production, particularly in regulated industries
 3. **Government** — U.S. government agencies (Air Force, Navy, Department of Energy); CRADA enablement; sovereign/regulated compute needs
 4. **Startups and Researchers** — Individual developers, academic researchers, and AI startups
 5. **Foundations/Hyperscalers** — Hyperscalers and foundation model companies racing to build global AI infrastructure
-

Customer Examples by Segment

- **Hyperscalers:** Microsoft (multibillion-dollar agreement), unnamed AI hyperscaler (built cluster in 90 days)
 - **AI Natives/Frontier Labs:** Pika (video generation), fal (open-source generative AI), Meshy (3D AI), Nous Research (Hermes 3), Moonshot AI (Kimi K2)
 - **Enterprise/Biotech/Healthcare:** Iambic Therapeutics (drug discovery), Genesis Therapeutics (drug discovery), Synlico (cell biology), ServiceNow (Apriel 5B)
 - **Government:** Air Force, Navy, Department of Energy (via CRADAs)
 - **Research/Academic:** MIT, Stanford, Harvard, Caltech (97% of top U.S. research universities)
 - **Fortune 500:** Sony, Intel, Microsoft, Amazon Research, Raytheon, GE, Kaiser Permanente, Tencent
 - **Channel partners:** VARs, referral partners, and ISVs — indicating mid-market and enterprise channel reach
-

Company Size & Scale Signals

Lambda serves a remarkably wide range, from solo researchers to the world's largest AI organizations:

- **Superclusters:** 4,000 to 165,000+ NVIDIA GPUs; 3–5 year contracts
 - **1-Click Clusters:** 16 to 2,000+ GPUs; 1 week to 3 years
 - **Instances:** 1–8 GPUs; pay-as-you-go (starting at \$0.58/hr)
 - **Research grants:** Up to \$5,000 in cloud credits for academic/independent researchers
 - **Positioning:** *"From one GPU to hundreds of thousands"*
 - *"Trusted by hyperscalers, global enterprises, and frontier labs operating in regulated environments"*
-

Workload Types Targeted

- Foundation model training (8B to trillion-parameter models)
 - Large-scale distributed training
 - Production inference at scale
 - Fine-tuning and experimentation
 - Drug discovery / healthcare AI
 - Generative AI (video, 3D, image)
 - Agentic AI and reasoning workloads (GB300 NVL72 "optimized for AI reasoning")
-

Technical Maturity of Target Buyers

- **High** — assumes familiarity with Kubernetes, Slurm, CUDA, PyTorch, InfiniBand, MFU metrics, CLI/API orchestration
 - **Dual path:** Self-serve for technically sophisticated users; co-engineering for enterprise teams needing deeper support
-

Geographic Focus

- **Primarily U.S.-centric** — mission framed as *"give everyone in America the power of superintelligence"*
 - **Data centers:** Texas, California, Ohio, Illinois, Missouri, Virginia/DC; partnership with Cologix in Columbus
 - **Some Canada presence**
 - **Hybrid/multi-cloud connectivity:** Networking to AWS, Azure, GCP, OCI, and on-prem
 - **GTC presence** in San Jose, CA
 - **Government pages** reference U.S. agencies specifically
-

Persona Alignment

Lambda's positioning maps to the following buyer personas, in order of strength:

- **Alice (AI Platform Leader) — Primary:** Directly addressed by orchestration flexibility (Kubernetes, Slurm, dstack, SkyPilot), cost transparency/predictable pricing ("flat rates —no lock-in, no ingress/egress fees"), SOC 2 Type II security, single-tenant architecture, observability stack, and vendor flexibility. Lambda's messaging directly counters her fear of being trapped in proprietary stacks.
 - **Carter (CTO/Technical Buyer) — Primary:** "Reserve your AI factory" CTA and co-engineering approach appeal to strategic build-vs-buy evaluation. Guaranteed GPU capacity (dedicated bare-metal, up to 165,000+ GPUs), performance optimization (MLPerf results), security/compliance, transparent pricing, and SLAs address his top buying triggers. Strong NVIDIA partnership and early access to latest GPUs address his focus on guaranteed high-performance capacity.
 - **Paul (Platform Engineer) — Strong:** Lambda Stack (pre-installed PyTorch, CUDA), managed Kubernetes/Slurm, S3-compatible storage, API/CLI automation, cloud metrics dashboard, built-in observability, and open-source tool compatibility (dstack, SkyPilot, Ray) align with daily workflow needs for cost transparency, observability, and open-source support.
 - **Malik (AI Researcher) — Strong:** Research page, grants program, LLM benchmarks leaderboard, MFU optimization whitepapers, NeurIPS/ICML publications, extensive GPU benchmarks, technical blog tutorials, fine-tuning guides, and open-source model support (Hugging Face, vLLM, MLflow) address needs for compute availability and cutting-edge features.
 - **Cedric (CEO) — Secondary:** Investor page, \$1.5B+ fundraising, Microsoft partnership, and "one person, one GPU" mission address CEO-level concerns about vendor stability and strategic vision.
 - **Diana (AI Leader) — Secondary:** Relevant through Lambda's enterprise AI deployment positioning and regulated-industry focus.
-

Implied Marketing Strategy

Implied Marketing Strategy: Consolidated Summary

Lambda is executing a **focused vertical cloud strategy** to become the **default AI infrastructure layer** — positioned between hyperscalers (AWS/Azure/GCP) above and commodity GPU clouds below — by being the only pure-play, AI-native cloud operating at scale. The strategy rests on several deeply interlocking components:

1. NVIDIA Alignment as the Deepest Competitive Moat

- Seven-time NVIDIA Partner of the Year, NVIDIA Exemplar Cloud (one of the first), Platinum GTC sponsor, and NVIDIA investor in Lambda's cap table.
- Early/day-one launch partner for every new GPU generation (GB300 NVL72, HGX B300, Blackwell, Vera Rubin), with products built on NVIDIA reference architectures (NVLink, InfiniBand, SHARP).
- Creates a **supply-side advantage** — guaranteed GPU capacity when GPU scarcity is the #1 buying trigger — and a **trust signal** — enterprises treat NVIDIA validation as a proxy for infrastructure quality.
- This "fast-follower to NVIDIA innovation" positioning is extremely difficult for competitors to replicate.

2. "Emerging Hyperscaler" Scale with AI-Native Focus

- **Category creation:** Branding products as "AI factories" rather than cloud instances elevates the conversation from price-per-GPU-hour to industrial-scale production capability.
- Building toward 2–3GW+ of AI data center capacity and 1M+ GPUs by decade's end, supported by massive fundraising (\$320M Series C → \$480M Series D → \$1.5B+ Series E) and multibillion-dollar agreements with Microsoft.
- Recruiting hyperscaler veterans (Jerry Hunter from AWS, Leonard Speiser from Clover) to execute at hyperscaler scale while retaining the agility and specialization of a focused cloud.
- **"100% of our engineering, operations, and support is dedicated to AI workloads"** — the core differentiator vs. AWS/Azure/GCP, which serve all workloads. This translates into better performance, better support (ML engineers on support, not generalists), and better tooling.

3. Systematic Anti-Hyperscaler Positioning

Lambda attacks hyperscaler weaknesses across every dimension its personas care about:

- **Vendor lock-in:** "No lock-in," open-source orchestration options (dstack, SkyPilot), multi-cloud networking
- **Cost opacity:** "No hidden fees. No charges for data ingress or egress." Flat rates, pay-by-the-minute, publicly posted pricing tables — directly countering cloud bill shock
- **Shared tenancy risks:** "Single-tenant, shared-nothing" architecture addresses security concerns and guarantees performance consistency
- **Generic infrastructure:** "Purpose-built for AI" — every optimization decision made for

4. Land-and-Expand Product Ladder with Flexible Commitment Models

Three tiers capture different buyer needs and create a natural expansion path:

Tier	Scale	Commitment	Target Buyer
Instances	1 + GPU, self-serve	Pay-by-minute	Startups experimenting, individual researchers
1-Click Clusters	16–2,000 + GPUs	1 week minimum	Teams running hero training runs; fills gap between expensive reserved and unreliable spot instances
Superclusters	4,000–165,000 + GPUs	1–5 year reserved	Enterprises, hyperscalers, frontier labs building production systems

This creates a classic land-and-expand motion: researchers start with instances, graduate to clusters, and enterprises commit to superclusters with locked-in recurring revenue.

5. Co-Engineering and Embedded ML Expertise as Switching Cost

- "Experts included" — Lambda's ML, HPC, and infrastructure engineers embed with customer teams.
- Differentiates from self-service-only competitors (CoreWeave, commodity GPU clouds) and addresses enterprise need for a "co-innovation partner."
- Creates deep integration and **high switching costs** that deepen customer relationships over time.

6. Security, Compliance, and Government as Enterprise Unlocks

- SOC 2 Type II, single-tenant architecture, hardware-level isolation, Customer Trust Portal — table stakes for enterprise/government that many GPU clouds lack.
- Explicit government segment (CRADAs, Air Force/Navy/DOE) signals credibility for regulated industries and potential for sovereign/classified compute.

7. Research Credibility as Demand Generation

- Publishing at top ML venues (NeurIPS, ICML, ICCV, CoRL, EMNLP), sponsoring the ARC Prize, offering research grants, maintaining LLM benchmarks leaderboard, and publishing MFU optimization whitepapers.

- Targets researchers and AI leaders who are primary influencers/recommenders in buying decisions, reaching them through exactly the channels they trust: peer communities, publications, and conferences.

8. Content Marketing as Technical SEO and Thought Leadership Engine

- ~200+ blog posts spanning GPU benchmarks, deployment tutorials, model serving guides, conference coverage, fundraising announcements, and partnership news.
- **LLM Index** (model cards with deployment guides and benchmarks on Lambda hardware) is a particularly effective demand-gen asset capturing engineers researching model deployment.
- Content targets practitioners directly with practical, reproducible technical content while executive announcements target decision-makers.

9. Strategic Focus Signal: Exiting On-Prem Hardware

- Ending the on-premise hardware business (August 2025) signals Lambda is going all-in on cloud/managed infrastructure — concentrating resources on higher-margin, recurring-revenue offerings and eliminating distraction from hardware sales.

Market Dynamics Being Exploited

- **GPU scarcity:** Explosive demand outstripping traditional data center capacity; Lambda's guaranteed capacity at scale (165,000+ GPUs) is a scarce asset
- **Hyperscaler fatigue:** Growing dissatisfaction with opaque pricing, vendor lock-in, and multi-tenant performance variability
- **Shift from training to inference dominance:** Reasoning models (o1, DeepSeek-R1) multiplying inference compute requirements
- **Open-source model proliferation:** DeepSeek, Llama, Qwen, OLMo — all require self-hosted inference infrastructure
- **Enterprise AI crossing the chasm:** Requiring compliance, security, and dedicated support
- **Multi-cloud/hybrid-cloud strategies** among enterprises seeking alternatives to single-provider dependency

Trends Being Ridden

- **Superintelligence/AGI narrative:** Creates urgency, FOMO, and justification for massive infrastructure investment; Lambda's "joining the race to superintelligence" language deliberately leverages this
 - **NVIDIA ecosystem dominance:** Deep NVIDIA embedding creates compounding ecosystem effects
 - **Liquid cooling and high-density data center design** becoming essential for next-gen GPU deployments
 - **"AI factory" paradigm** replacing traditional data center mental models
 - **Geographic expansion:** Data centers across US and Canada with multi-cloud connectivity; 3GW+ vision implies continued massive physical buildout
-

Operational Metrics

Source: *lambda.md*

Marketing – Factual & Quantitative: <https://lambda.ai>

Extract concrete, verifiable facts from a company's website: corporate details, infrastructure specs, customer names, product listings, and quantitative claims.

Based on 371 stored pages analysed in 1 batch(es) with 1 AI calls (namespace: webscrape, model: claude-opus-4-6, profile: operational_metrics).

Number Of Data Centers, Their Locations, Gpus, And Power Capacity

Data center locations and expansion:

- Lambda operates data centers across the **United States and Canada** (shown on infrastructure map)
- Specific locations mentioned:
 - **Texas** (persistent storage region since 2022)
 - **California** (headquarters in San Francisco; ECL facility in **Mountain View**)
 - **Columbus, Ohio** — Cologix COL4 ScalelogixSM data center; NVIDIA HGX B200 clusters deployed (June 2025). Cologix has four data centers spanning 500,000 sq ft and 80 MW in Columbus.
 - **Chicago, Illinois** — EdgeConneX facility; **23MW** single-tenant data center planned RFS 2026; also existing air-cooled sites. **Over 30 MW** planned for Chicago and Atlanta combined.
 - **Atlanta, Georgia** — EdgeConneX air-cooled site for Lambda
 - **Kansas City, Missouri** — AI Factory planned launch early 2026 with **24MW initial capacity**, potential to scale to **more than 100MW**. More than **10,000 NVIDIA GPUs** planned.
 - **Southern California (Vernon/LAX01)** — Prime Data Centers facility; Lambda initially leasing **21MW** of a 33MW, 242,000 sq ft facility
 - **Washington, DC area (US-East-2, US-East-3)** — S3 Adapter supported regions
 - **Utah** — mentioned as a region
 - **Virginia** — mentioned as a region

GPU types available/planned:

- **NVIDIA GB300 NVL72** (Blackwell Ultra) — 72 GPUs per rack, 36 Grace CPUs; 142 kW per system; operational (first stood up October 2025; hydrogen-powered units at ECL Mountain View)
- **NVIDIA Vera Rubin NVL72** — announced January 2026, planned availability second half of 2026; 72 Rubin GPUs, 36 Vera CPUs per rack
- **NVIDIA HGX B300** — announced for Superclusters

- **NVIDIA HGX B200 (SXM6)** — 180 GB VRAM per GPU; available in 1-Click Clusters and On-Demand Instances
- **NVIDIA H200** — 141 GB HBM3e
- **NVIDIA H100 SXM** — 80 GB; available in clusters and instances
- **NVIDIA H100 PCIe** — 80 GB
- **NVIDIA GH200** Grace Hopper Superchip — 96 GB HBM3; 72-core Grace CPU + H100 GPU
- **NVIDIA A100 SXM** — 80 GB and 40 GB variants
- **NVIDIA A100 PCIe** — 40 GB
- **NVIDIA A6000** — 48 GB
- **NVIDIA A10** — 24 GB
- **NVIDIA Tesla V100** — 16 GB
- **NVIDIA Quadro RTX 6000** — 24 GB

Power capacity:

- Lambda's stated vision: "**gigawatt-scale AI factories**" and "**3GW + of AI data center space**"
- Contiguous AI factories described as **75MW + campuses**
- Kansas City: **24MW** initial, scalable to **100MW +**
- Vernon/LAX01: **21MW** leased of 33MW campus
- Chicago EdgeConneX: **23MW** planned
- Ken Patchett quote: "Our megawatt footprint, we're multiplying by 4 between 2025 and 2026"
- Rack power densities: **130 to 240 kW per rack** mentioned; roadmap toward **1 MW rack-scale designs**
- GB300 NVL72 systems: **142 kW** per system

Scale claims:

- Superclusters: **4,000 to 165,000 + NVIDIA GPUs**
- 1-Click Clusters: **16 to 2,000 + NVIDIA GPUs**
- Instances: 1 to 8 GPUs
- Over **200,000 AI developers/users** served
- Over **5,000 hardware/private cloud customers**
- Microsoft agreement: "**tens of thousands of NVIDIA GPUs**" including GB300 NVL72

Active And Contracted Megawatts, Terawatts, And Other Measures Of Power

- **3GW + vision** for AI data center space (stated on Superclusters page)
- **Kansas City, MO:** 24MW initial capacity, potential to scale to **more than 100MW**
- **Vernon/LAX01 (Prime Data Centers):** Lambda leasing **21MW** of 33MW facility
- **Chicago + Atlanta (EdgeConneX):** Over **30 MW** planned
- **Columbus, Ohio (Cologix):** Portfolio of **80 MW** across four data centers (Cologix total, not solely Lambda)
- **ECL Mountain View:** Lambda **doubled its footprint** from 50% to 100% of the facility
- **Megawatt footprint multiplying by 4x between 2025 and 2026** (Ken Patchett)
- Contiguous AI factories: **75MW + campuses** for large clusters
- Rack densities: **130-240 kW per rack**; roadmap toward **1 MW rack-scale**
- GB300 NVL72: **142 kW** per system/rack

- NVIDIA HGX B200 TDP: up to **700W per GPU**
- PUE improvement: Direct-to-chip liquid cooling improves PUE by **over 15%**

Renewable/alternative energy:

- **Hydrogen-powered** NVIDIA GB300 NVL72 systems at ECL Mountain View — zero-water, zero-emissions, off-grid modular data center using hydrogen fuel cells
 - On-site power generation strategies mentioned: natural gas, battery backup, solar, wind connectivity
 - Ken Patchett: "really important now to start thinking about...on-site power generation"
-

Installed, Active And Future Number Of Gpus

- **Superclusters:** 4,000 to **165,000 + NVIDIA GPUs** per deployment
- **1-Click Clusters:** 16 to **2,000 + NVIDIA GPUs**
- **Instances:** 1 to 8 GPUs per instance
- **Kansas City AI Factory:** More than **10,000 NVIDIA Blackwell Ultra GPUs** at launch, expected to double over time
- **Microsoft agreement:** "tens of thousands of NVIDIA GPUs" including GB300 NVL72
- **One customer (AI Hyperscaler) case study:** "thousands of GPUs" deployed in 90 days
- Lambda has "**deployed tens of thousands of GPUs**" for AI startups, hyperscalers, and Fortune 500 companies (Considerations blog)
- Vision: "**over a million GPUs deployed by the end of the decade**" (EdgeConneX blog)
- "From one GPU to hundreds of thousands" (homepage tagline)
- Ken Patchett mentioned trajectory toward **128K-GPU clusters**

Specific GPU models deployed:

- NVIDIA GB300 NVL72 — operational (72 GPUs per rack)
- NVIDIA HGX B200 — available on-demand and in 1-Click Clusters
- NVIDIA H100 SXM/PCIe — widely deployed in cloud and clusters
- NVIDIA H200 — deployed in clusters
- NVIDIA GH200 — available on-demand; 10x demo clusters offered
- NVIDIA A100 (80GB & 40GB) — available
- NVIDIA A6000, A10, V100, RTX 6000 — available in instances

Future GPU deployments:

- **NVIDIA Vera Rubin NVL72** — planned for second half of 2026
 - **NVIDIA HGX B300** — announced for Superclusters
 - **NVIDIA GB300 Desktop Superchip** — coming to Vector Pro later in the year
 - Lambda is an early partner for **NVIDIA Vera CPUs** and **NVIDIA BlueField-4 STX**
-

Teraflops And Exaflops

- **NVIDIA GB200 NVL72:** **1.4 exaFLOPS** of AI performance (from Lambda blog on Blackwell announcement)
- **NVIDIA HGX B300:** **72 PF FP8 training / 144 PF FP4 inference** per system; **2.1 TB HBM3e memory** (Superclusters page)
- **NVIDIA GB300 NVL72:** " + 1.5 × higher dense FP4 performance & 2 × faster attention

operations" vs GB200 NVL72

- **NVIDIA H100 SXM:** 3,958 TFLOPS FP8 with sparsity (GPU guide)
 - **NVIDIA H100:** FP8/FP16 TFLOPS: 3,958 (SXM5) at 700W = 5.7 FP8 TFLOPS/W; FP16 Tensor Core: 312 TFLOPS; FP32: 19.5 TFLOPS
 - **NVIDIA A100 SXM4:** FP16 Tensor Core: 312 TFLOPS (with sparsity 624); TF32: 156 TFLOPS; FP32: 19.5 TFLOPS
 - **NVIDIA V100:** FP16 Tensor Core: 125 TFLOPS; FP32: 15.7 TFLOPS
 - **Hyperplane-16 cluster (16x V100):** Cluster SGEMM: **720 TFLOPS FP32** (3-node cluster example from TCO blog)
 - GPT-3 175B training estimated at **3.14E23 FLOPS** (from GPT-3 blog)
 - "Theoretical 28 TFLOPS for V100" mentioned in GPT-3 cost analysis
 - FlashAttention-2 on H100: Extrapolated GPT3-175B training replication with **242,400 GPU hours** (H100 80GB SXM5), translating to **\$458,136** on Lambda Cloud
-

Other Capacity And Performance

Networking:

- **NVIDIA Quantum-2 InfiniBand:** 400 Gb/s; aggregate **3.2 Tb/s** per 1-Click Cluster node
- **NVIDIA Quantum-X800 InfiniBand:** 800 Gb/s per GPU for multi-rack clusters; **6.4 Tb/s** bandwidth between nodes
- **NVIDIA NVLink 5 (GB300 NVL72):** **1.8 TB/s** bandwidth per GPU; **130 TB/s NVLink Switch bandwidth** per rack
- **NVLink-C2C (GH200):** **900 GB/s** bidirectional bandwidth
- 1-Click Clusters: **3200 Gb/s** node-to-node compute between worker nodes
- Co-Packaged Optics (CPO): NVIDIA cites **3.5x power efficiency improvement, 10x higher resilience, 5x longer AI application runtime, 1.3x faster time to operation**
- SHARP acceleration: **+50% to +63% allReduce bandwidth improvement** on Lambda clusters (16-1536 GPUs)
- Instances: up to **10 Gbps** internet bandwidth

Memory:

- **NVIDIA GB300 NVL72:** **37 TB fast memory** (20 TB HBM3e) per rack; **3.84 TB NVMe cache per GPU** (276 TB per NVL72)
- **NVIDIA Vera Rubin NVL72:** Up to **20.7 TB HBM4** per rack; bandwidth targets up to **22 TB/s per GPU**
- **NVIDIA HGX B300:** **2.3 TB HBM3e** (NVL16) or **20 TB HBM3e** (GB300 NVL72)
- **NVIDIA B200:** **180 GB HBM3e** per GPU
- **NVIDIA H200:** **141 GB HBM3e**; **4.8 TB/s** memory bandwidth
- **NVIDIA H100 SXM:** **80 GB HBM3**; **3.35 TB/s** memory bandwidth
- **NVIDIA GH200:** Up to **576 GB** total (96 GB HBM3 + 480 GB LPDDR5X)

Benchmark performance results:

- **MLPerf Training v5.1:** Lambda's GB300 NVL72 outperforms GB200 NVL72 by **27%** (Llama 2-70B LoRA: 1.26 min vs 1.598 min)
- **MLPerf Inference v5.1:** Up to **15.4%** performance gains on NVIDIA HGX B200 vs prior round; Llama 3.1 405B Server: **+15.4%**
- **Ceramic AI on HGX B200:** MFU up to **85%** at 32K context (8B model on 8 Blackwell GPUs)

- **MFU optimization whitepaper:** Peak MFU above **60%** vs 35-45% industry norm; **2.11x MFU uplift** for Llama 70B on 16x HGX B200
- **Olmo Hybrid training on 512 B200 GPUs:** 3 trillion tokens in **7 days**; **97% active training time** (99% excluding development); median recovery time **under 4 minutes**
- **FlashAttention-2 (H100):** 3x+ speedup over baseline; H100 2x faster than A100
- **DeepChat RLHF (H100 vs A100):** H100 **2.5x-3.1x** faster
- **H100 vs A100:** Up to **3x throughput** on Tensor Core (same data types); **6x** with FP8 + Transformer Engine

Pricing (on-demand per GPU per hour):

- NVIDIA B200 SXM6 (8x): **\$5.74/GPU/hr**
- NVIDIA H100 SXM (8x): **\$3.44/GPU/hr**
- NVIDIA GH200: **\$1.99/hr**
- NVIDIA A100 SXM 80GB (8x): **\$2.06/GPU/hr**
- NVIDIA A10: **\$0.86/hr**
- NVIDIA RTX 6000: **\$0.58/hr**
- 1-Click Cluster HGX B200 on-demand: **\$4.62/GPU/hr**
- 1-Click Cluster H100 on-demand: **\$2.76/GPU/hr**
- Storage: **\$0.20/GB/month**; no ingress/egress fees

Fundraising:

- Series E: **Over \$1.5B** (November 2025, led by TWG Global and USIT)
- Series D: **\$480M** (February 2025)
- Series C: **\$320M** (February 2024)
- Series B: **\$44M** (March 2023)
- Series A: **\$15M** equity + \$9.5M debt (July 2021)
- Microsoft: "**Multibillion-dollar agreement**" for AI infrastructure

Company details:

- Founded: **2012** in San Francisco
 - Founded by: Published AI/ML engineers (Stephen Balaban, Michael Balaban)
 - Users: **Over 200,000** cloud users; **over 150,000** mentioned elsewhere
 - Certifications: **SOC 2 Type II**
 - **Seven-time NVIDIA Partner of the Year** (2020-2025, including Healthcare 2025)
 - **NVIDIA Exemplar Cloud** status achieved (October 2025)
 - Member of **Open Compute Project (OCP) Advisory Board** (January 2026)
 - Customers include: Microsoft, Air Force, Navy, Department of Energy, Intel, Google, Amazon Research, Tencent, Kaiser Permanente, MIT, Stanford, Harvard, Caltech, Department of Defense, Genesis Therapeutics, Iambic Therapeutics, Pika, fal, Meshy, ServiceNow, Nous Research, Synlico, Voltron Data, Leonardo.ai, Abacus.ai
-
-