

Consolidated Report: nebius

Generated 2026-03-29 04:10 UTC from 3 analysis profile(s): company_profile, marketing_analysis, operational_metrics.

Note: The primary source of information in this report is the company's own website. The content reflects the company's self-reported claims and marketing materials and should not be treated as independently verified fact.

Executive Summary

Executive Summary: Nebius

Company Overview

Nebius (Nasdaq: NBIS) is an AI-native cloud infrastructure company headquartered in Amsterdam, Netherlands. While its corporate roots trace back 25 years to the Yandex era (founded 2000, IPO 2011), the Nebius AI Cloud platform was publicly launched in November 2023, making it roughly 1.5–2 years old as a standalone brand. Led by founder and CEO Arkady Volozh, the company has rapidly scaled into a full-stack cloud platform purpose-built for AI workloads—spanning GPU compute, model training, inference, and managed services—with operations across Europe, North America, and Israel, and expansion planned into Asia-Pacific.

Key Findings

Corporate Profile & Customer Base. Nebius has assembled a remarkably diverse and high-profile customer portfolio in a short time. Anchor contracts with **Meta** (up to ~\$27 billion over five years) and **Microsoft** (multi-billion dollar agreement) provide revenue visibility and credibility. Beyond hyperscalers, the company serves 50+ named customers across healthcare, fintech, robotics, media, e-commerce, and AI research—including Revolut, Brave Software, Shopify, Recraft, and Stanford's CRISPR-GPT project. The leadership team blends deep technical expertise with operational experience, supported by a board that includes NVIDIA-aligned investors (Accel's Matthew Weigand) and seasoned technology directors.

Strategic Positioning & Brand Identity. Nebius positions itself squarely between hyperscalers and bare-metal GPU providers, claiming the advantages of both: "supercomputer performance with hyperscaler simplicity." The brand identity is builder-centric, technically authoritative, and grounded in verifiable benchmarks rather than abstract claims. Key differentiators include full vertical integration (custom-designed servers, proprietary software stack, owned data centers), an exceptionally deep NVIDIA partnership (\$2 billion investment, Exemplar Cloud Partner status, early access to each GPU generation), and aggressive cost competitiveness—with published GPU pricing (H100 at \$2.00/hr, B200 at \$3.00/hr) and independently verified inference costs up to $3\times$ more efficient than competitors.

Operational Scale & Performance. Nebius operates data centers in Finland (expanding to 75 MW), France, Iceland, Kansas City, and the UK, with a 300 MW facility in New Jersey and

a gigawatt-scale campus in Independence, Missouri under development. The company targets 5+ GW of NVIDIA systems deployed by 2030. Its ISEG2 supercomputer ranks #13 on the Top500. Operational metrics are strong: 167,000 GPU-hours MTBF on a 3,000-GPU cluster, 12-minute average recovery time for multi-host training, 99.9% inference SLA, and 100% customer satisfaction since July 2025. Financially, the company reports 479% year-over-year revenue growth, a \$20+ billion contract backlog, \$3.7 billion in cash, and a target annualized run-rate of \$7–9 billion by end of 2026.

Cross-Cutting Themes, Strengths & Risks

Strengths. The convergence of vertical integration, NVIDIA strategic alignment, and anchor-tenant contracts creates a compounding competitive moat that is difficult to replicate. The dual-product strategy (AI Cloud for training, Token Factory for inference) captures the full model lifecycle and builds natural stickiness. Open-source ecosystem alignment and compliance breadth (SOC 2, HIPAA, ISO 27001, GDPR, NIS2) broaden the addressable market from startups to regulated enterprises.

Risks. Nebius's growth depends heavily on continued NVIDIA supply priority and the execution of massive capital-intensive data center buildouts. Concentration risk is significant: the Meta and Microsoft contracts dominate the revenue outlook. As a young brand competing against hyperscalers with entrenched enterprise relationships, customer acquisition beyond the AI-native segment will require sustained investment. Geopolitical perception risk—stemming from the Yandex legacy—could affect enterprise adoption in sensitive verticals or markets.

Strategic Outlook

Nebius is executing an ambitious but coherent strategy to become the dominant AI-specialized cloud provider globally. If the company delivers on its infrastructure roadmap and converts its contract backlog into recurring revenue, it will transition from a high-growth neocloud into a structurally significant player in AI compute. The critical milestones to watch are the successful ramp of US data center capacity, diversification of the enterprise customer base beyond anchor tenants, and sustained differentiation as hyperscalers invest aggressively in their own AI infrastructure offerings. The next 18–24 months will determine whether Nebius can cement its position as the infrastructure backbone for the open-source AI era.

Company Profile

Source: nebius.md

Marketing – Factual & Quantitative: <https://nebius.com>

Extract concrete, verifiable facts from a company's website: corporate details, infrastructure specs, customer names, product listings, and quantitative claims.

Based on 748 stored pages analysed in 2 batch(es) with 7 AI calls (namespace: webscrape, model: claude-opus-4-6, profile: company_profile).

Headquarters

Headquarters – Consolidated Summary

Primary Headquarters:

- Nebius is headquartered in **Amsterdam, Netherlands**.
- Office address: **Tripolis Park, Burgerweeshuispad 101, 1076 ES Amsterdam, Netherlands**.
- Corporate entity: **Nebius B.V.** (also referenced as **Nebius Group N.V.**).
- Listed on **Nasdaq** under ticker **NBIS**.

U.S. Offices:

- **San Francisco** – Ferry Building
- **Dallas, Texas** – Dallas Parkway
- **New York**
- **Legal/DMCA contact address:** Nebius Legal, 10 State St., Newburyport, MA 01950

R&D Hubs:

- Located across **Europe, North America, and Israel**.

Both batches were fully consistent with no contradictions.

Years In Business

Nebius – Years in Business (Consolidated Summary)

Corporate Origins (Yandex era)

- **Founded / earliest executive tenure: 2000** – Arkady Volozh served as CEO and Executive Director from 2000 to 2022 (resuming in 2024). John Boynton has been a Non-Executive Director since 2000 and Chairman since 2016.
- **Initial Public Offering: 2011** – Charles Ryan became a Non-Executive Director at the time of the IPO.
- The group therefore has **~25 years of corporate history** tracing back to its Yandex roots.

Nebius AI Cloud Launch

- **Publicly launched: November 2023** as a European AI cloud platform, representing a strategic rebrand and transformation away from the legacy Yandex business.
- The company **celebrated its 1-year anniversary** of the public launch in **October 2024**, and was again marked by Nasdaq in Times Square in **July 2025**.
- As a standalone AI cloud brand, Nebius has been operating for roughly **1.5–2 years** (as

of mid-2025).

Other Notable Timelines

- **2002** – Slurm workload manager, referenced as a key technology used by Nebius, began development.
- **2015–2016** – First data center heat-recovery project in Finland was phased in.

Bottom Line

Nebius Group N.V. has **deep institutional roots going back to 2000** (and an IPO in 2011), but the **Nebius AI Cloud platform itself is a young venture, publicly launched in November 2023**. Depending on the frame of reference, the company can credibly cite either ~25 years of group-level history or ~1.5 years as a dedicated AI cloud business.

Company Executive Names And Titles

Nebius Group – Consolidated Executive Names and Titles

C-Suite / Senior Leadership

- **Arkady Volozh** – Founder, Chief Executive Officer, Non-independent Executive Director. Background in computer science; degree in applied mathematics from Gubkin Institute of Oil and Gas. Co-founded CompTek International and InfiNet Wireless. Early investor in Face.com (sold to Facebook 2012) and Getir; served on board of Neurosteer.
- **Ophir Nave** – Chief Operating Officer, Non-independent Executive Director (appointed May 2024). Doctor of Juridical Science from Harvard Law School; LL.B. from Tel Aviv University; B.Sc. in Computer Engineering from Technion. Previously Lead Partner at Arnon, Tadmor-Levy law firm.
- **Roman Chernin** – Co-founder, Chief Business Officer (CBO).
- **Andrey Korolenko** – Co-founder, Chief Infrastructure and Product Officer (also referenced as Chief Product and Infrastructure Officer / Head of Infrastructure).
- **Dado Alonso** – Chief Financial Officer.
- **Marc Boroditsky** – Chief Revenue Officer.
- **Danila Shtan** – Chief Technology Officer (CTO).
- **Boaz Tal** – General Counsel.
- **Daniel Bounds** – Chief Marketing Officer.
- **Tom Blackwell** – Chief Communications Officer.

*Note: **Danila Pavlov** was referenced as CFO in an earlier webinar context, suggesting he may have been a former/interim CFO prior to Dado Alonso.*

Board of Directors

Name	Role	Key Background
------	------	----------------

(appointed March 2026)

- **John Haarer** – General Manager for Asia-Pacific and Japan (appointed March 2026)
- **Rashid Ivaev** – General Manager, EMEA

Directors & Heads of Function

- **Boris Yangel** – Head of AI R&D team (decade+ experience leading complex AI projects)
- **Fedor Zhdanov** – Head of Applied AI
- **Gleb Kholodov** – Head of Infrastructure-as-a-Service (IaaS) Cloud Platform
- **Oleg Fedorov** – Head of Hardware R&D
- **Aleksandr Patrushev** – Head of Product Management AI/ML
- **Narek Tatevosyan** – Director of Product Management
- **Vinita Ananth** – Senior Product Director
- **Gleb Evstropov** – Director of Engineering
- **Mona Li** – Global Head / Director of Startup Ecosystem
- **Ilya Burkov** – Global Head of Life Science and Healthcare Sales/Growth
- **Evan Helda** – Head of Physical AI / GTM Physical AI
- **Joshua Liss** – Head of GTM Media & Entertainment
- **Alexey Tsishevski** – Head of GTM
- **Andrew Eliseev** – Head of GTM Retail & Commerce AI
- **Anubhav Maheshwari** – Head of Venture Partnerships
- **Peter Morley** – Head of Communications
- **Michael Rothman** – Chief of Staff and GM Israel
- **Olga Chasovskaya** – Head of Business Growth (AI Studio / Token Factory)

Other Named Personnel

- **Nathalie van Wigen** – Company Secretary
- **Mikhail Mokrushin** – Managed Schedulers Team Leader (Soperator)
- **Dylan Bristot** – Product Marketing Manager (Token Factory)
- **Danila Pavlov** – Former/earlier CFO (referenced in webinar context)

Customer Names, Locations, Industry/Type Of Entity, And Other Relevant Details

Consolidated Customer Summary for Nebius

Hyperscale / Strategic Partnerships

- **Microsoft** – Multi-billion dollar AI infrastructure agreement. Nebius to provide dedicated capacity from upcoming Vineland, New Jersey data center. Deal valued at approximately \$3 billion over 5 years (Q3 2025 announcement), plus an additional agreement for up to ~\$27 billion (March 2026 context).
- **Meta** – Long-term AI infrastructure supply agreement worth up to approximately \$27

billion over five years, including \$12 billion dedicated capacity across multiple locations and up to \$15 billion additional available compute capacity. Nebius is a partner of Meta, featured at LlamaCon 2025.

Named Customers with Details

Search & Browser

- **Brave Software** – Privacy-focused browser and Brave Search; 80M+ users, 1.3B search queries/month, 11M+ AI-generated answers daily. Uses Nebius for AI inference with 10–70B LLM parameters and nearly 100% compute utilization. Testimonial from Rémi Berson, Principal Engineer.

AI Frameworks & Open-Source Infrastructure

- **vLLM** – Open-source LLM inference framework under the Linux Foundation. Uses Nebius compute clusters for large-scale inference experiments including DeepSeek R1. Zero hardware-related issues reported; consistently accurate performance metrics.
- **SGLang** – LLM inference framework. Achieved $2\times$ throughput boost and lower latency for DeepSeek R1 on Nebius.
- **Hugging Face** – Collaboration on inference access / inference provider integration. Testimonial from Julien Chaumond, CTO.

Design & Creative AI

- **Recraft** – AI design tool for digital illustrations, vector art, icons. Trained a 20B-parameter generative AI model from scratch on Nebius using Kubeflow + PyTorch. Deployed NVIDIA HGX B200 for Recraft V4. 350M+ images generated on platform. Funded by Khosla Ventures and Nat Friedman. Testimonial from Anna Veronika Dorogush, Founder & CEO.
- **Photoroom** – AI-powered photo editing. Testimonial from Eliot Andres, Co-founder & CTO: "Nebius provided the reliable infrastructure we needed to scale training seamlessly."

AI Deployment & Inference Optimization

- **CentML** – AI deployment platform. Uses Nebius compute for inference optimization. Achieved $5\times$ lower costs compared to other major providers, EU compliance, 1 week to get cluster online. Testimonial from Shang Wang, CTO & Co-Founder.
- **TheStage AI** – Inference simulators and DNN optimization. Used NVIDIA H100 on Nebius for stable diffusion inference. $4\times$ speed leap, $\sim 500\text{ms}$ per image, 1B parameters. CEO Kirill Solodskih noted $3.5\times$ faster inference for diffusion models on B200.

Distributed Training & AI Infrastructure

- **Prime Intellect** – Go-to cloud for flexible on-demand compute. PoC with NVIDIA

GB200 NVL72. Trained INTELLECT-3 model.

AI Research

- **Decart** – AI research company. Testimonial from Menash Land, Partnerships Lead: "Nebius stands out due to their responsiveness, personal attention, and technical expertise."

Fintech & Financial Services

- **Revolut** – London-based fintech; 70M+ users. Rebuilt operations around AI with 200+ NVIDIA H100 GPUs on Nebius. Uses Token Factory for fraud prevention and customer support, handling 1.2M support tickets/month and up to 200B tokens/day. Also featured in Inference Frontier Program.

E-Commerce & Lifestyle

- **Prosus** – Power behind leading lifestyle and e-commerce brands. Achieved up to 26× cost reductions compared to proprietary models. Up to 200 billion tokens per day on dedicated endpoints. Testimonial from Zülküf Genç, Director of AI.
- **Shopify** – Uses Nebius large training clusters; objective is easy developer access to compute for rapid iteration. Harley Finkelstein, President, quoted regarding Toloka/Tendem partnership.

Video & Media AI

- **Higgsfield AI** (now Mirage, formerly Captions) – Leading AI video platform; \$130M Series A. Trained multi-billion parameter diffusion model on NVIDIA HGX B200. Scaled to \$200M run rate in 9 months, 15M+ users, 4.5M+ generations/day. Uses Nebius for on-demand and autoscaling inference; 40% cost efficiency gains with L40S. Testimonials from Alex Mashrakov, Founder & CEO, and Gaurav Misra, Co-Founder & CEO.
- **Dubformer** – AI dubbing and localization in 70+ languages. Uses NVIDIA A100 and H100 GPUs, Nebius Compute Cloud, and Object Storage for hundreds of TB.

Voice AI

- **Krisp** – AI-powered noise cancellation and accent localization. Used H100 GPUs; training 50–80% faster than A100.

Chatbot / Conversational AI

- **Chatfuel** – AI chatbot platform (founded 2015, Y Combinator). Used Nebius AI Studio for Llama 405B cascade. 24% improvement in routing over GPT-4o.

Mental Health AI

- **Slingshot AI** – New York-based mental health AI company (\$93M raised; a16z/Radical Ventures). Foundation LLM for psychology "Ash," launched July 2025.

Multimodal AI Research

- **Unum** – AI research lab developing compact multimodal models. Trained and open-sourced UForm-Gen2 and UForm-VL-English models on Nebius (training ~1 day on 8 H100s).

Genomics & Gene Editing

- **Stanford / Princeton / Google DeepMind (CRISPR-GPT)** – LLM-powered agent system for gene editing. Used Nebius GPU clusters for model screening and fine-tuning. Junior/novice researchers achieved 80–90% editing efficiency on first attempt.

Precision Medicine & Single-Cell Biology

- **Converge Bio** – Tel Aviv-based; precision medicine, single-cell RNA sequencing with LLMs. Trained Converge-SC model on 36M+ cells, ~2TB data, 7,000+ NVIDIA H100 GPU hours. Testimonials from Dov Gertz, CEO, and Noam Azoulay, ML engineer.

AI Neuroscience & Epigenetics

- **Prima Mente** – San Francisco/London-based AI neuroscience company. Uses 256 NVIDIA H200 GPUs on Nebius for epigenetic foundation models (Pleiades model: 90M, 600M, 7B parameters) for DNA methylation. Testimonial from Ravi Solanki, CEO.

Quantum Chemistry

- **Simulacra AI** – London-based; quantum chemistry with deep learning. Uses 16 NVIDIA H200 GPUs and NVIDIA BioNeMo to build large wavefunction models. Testimonial from Aleksei Ustimenko, Co-Founder & CEO.

Biotech & Drug Discovery

- **Helical** – Foundation model training for therapeutics. Trained Helix-mRNA model (500M parameters in 10 days on 16 GPUs). Testimonial from Rick Schneider, CEO.
- **Compugen** – Israel-based biotech/immuno-oncology. Trained spatial immune features prediction model. Lead drugs COM701, COM902; licensed to AstraZeneca and Gilead.
- **BostonGene** – Spatial immune features in tumors. Testimonial from Dr. Roy Granit.
- **Basecamp Research** – UK-based life sciences; AI for frontier challenges. Among first partners for UK deployment. Testimonial from Oliver Vince.
- **SieveStack** – Data generation and foundation model training for molecular simulations/drug discovery (Stanford/UCSF founders). 90%+ GPU utilization. Uses TractoAI. Testimonial from Slava Naprienko, CEO.
- **Quantori** – Cheminformatics research; developed EDM framework for molecular generation pipeline. Testimonial from Denis Sapegin.
- **YerevaNN** – Armenia-based non-profit ML research center. Three Nebius-based molecular models, 40B tokens; LLMs for molecular optimization. Testimonial from Hrant Khachatrian.

Medical Imaging

- **xAID** – AI for CT chest/abdomen imaging. Foundation model for CT scans, 580M parameters, 3 H100 GPUs, 5+ days per training cycle. Testimonial from Kirill Lopatin.

Clinical Trials

- **TrialHub** – Data intelligence platform for clinical trials. 250M vector database launched in days; 20× faster Standard of Care research. Testimonial from Ivaylo Yosifov, CTO.

Bioinformatics

- **ELIXIR BioHackathon / BioHackCloud** – Cross-border bioinformatics project. Testimonial from Jonathan Tedds.

Robotics & Physical AI

- **Positronic Robotics** – Startup training ML models on Nebius (H100 GPUs) for robotic cleaning.
- **RoboForce** – AI robots for outdoor environments. Uses NVIDIA Cosmos on Nebius; cut pipeline setup time by 70%+. Testimonial from Calvin Zhou.
- **Voxel51** – Physical AI data platform. Running FiftyOne workflows on Nebius GPU clusters. Works with Porsche. Testimonial from Brian Moore, CEO.
- **Porsche** – Autonomous driving data augmentation via Voxel51 + Nebius.

Autonomous Driving

- **Avride** – Sister company; autonomous driving. Uses Nebius infrastructure. Testimonial from Dmitry Polishchuk.

Video Management

- **Milestone Systems** – Video management software; fine-tuning VLMs on Nebius. Testimonial from Edward Mauser.

Graph AI & Analytics

- **Lynx Analytics** – Singapore-based Graph AI company. LynxKite 2000:MM platform; serves pharma, telecom, and finance verticals.

AI Post-Training & Data

- **SynthLabs** – AI post-training startup. Published Big Math dataset using TractoAI for distributed inference. Testimonial from Nathan Lile, Founder & CEO.

Distributed Computing / Other

- **Tilde** – Uses Soperator for distributed computing. Testimonial from Dhruv Pal.

- **Mithril** (formerly Foundry) – Strategic partnership integration. Uses Token Factory for real-time to batch inference. Testimonial from Raghav Kohli.
 - **Monday.com** – Featured in Inference Frontier Program. Testimonial from Assaf Elovic, Head of AI.
 - **FlowGPT** – Featured in Inference Frontier Program.
 - **Wubble** – AI music generation platform.
 - **London Institute for Mathematical Sciences (LIMS)** – Research on LLM abstraction capabilities using Nebius infrastructure via Research Grant program.
-

Industries Served

Healthcare & life sciences, precision medicine, genomics, drug discovery, biotech, immunology, clinical trials, bioinformatics, medical imaging, AI neuroscience, quantum chemistry, robotics & physical AI, autonomous driving, financial services (fintech), retail & e-commerce, media & entertainment, video AI, voice AI, mental health AI, web search, music generation, design & creative AI, distributed training infrastructure, cheminformatics, graph analytics, edtech, video management, and conversational AI.

Partner Names, Locations, Industry/Type Of Entity, And Other Relevant Details

Consolidated Summary of Nebius Partners

Strategic Investors & Financial Partners

- **NVIDIA** – Lead strategic partner. \$700M investment round led by NVIDIA and Accel; subsequently NVIDIA invested \$2 billion in Nebius (March 2026). Reference Platform NVIDIA Cloud Partner; one of the first to achieve NVIDIA Exemplar Status on H200 GPUs. Deep collaboration across hardware (Blackwell, Blackwell Ultra, GB200 NVL72, GB300, Vera Rubin) and software (NVIDIA AI Enterprise, NIM microservices, BioNeMo, Cosmos, OSMO, Isaac Sim, Run:ai, Dynamo, DGX Cloud Lepton). Joint work on inference, AI factory design, fleet management. Ecosystem partner for NVIDIA Dynamo. Collaboration with NVIDIA Inception for startups (AI Lift program offering up to \$150K credits). Joint PyTorch MoE training demonstration (up to 41% faster DeepSeek-V3 pre-training on Blackwell).
- **Accel** – Investor (part of \$700M round). Matthew Weigand serves on the board. Launch partner for Enterprise Readiness Initiative.
- **Insight Partners** – Launch partner for Enterprise Readiness Initiative.
- **Fellows Fund** – Launch partner for Enterprise Readiness Initiative.
- **Shanda Grab Ventures** – \$100K SAFE investment in Buildroid AI at Nebius Robotics Awards.

Major AI Infrastructure Customers / Agreements

- **Microsoft** – Multi-billion dollar AI infrastructure agreement. Dedicated capacity from Vineland, NJ data center.
- **Meta** – AI infrastructure agreement valued up to ~\$27B over 5 years (\$12B dedicated capacity + \$15B additional compute).

Data Center & Colocation Partners

- **DataOne** – AI hosting infrastructure company; building New Jersey data center (up to 300 MW).
- **Verne** – Sustainably powered data centers in the Nordics; Iceland colocation (10 MW, 100% renewable energy).
- **Patmos** – Kansas City data center partner (converted former Kansas City Star printing press facility; 5–40 MW capacity).
- **Equinix** – Paris colocation at PA10 campus, Saint-Denis.
- **Ark Data Centres** – UK facility at Longcross Park, Surrey.

Storage Partners

- **DDN** – Data intelligence platform; data architecture powers Nebius cloud.
- **WEKA** – High-performance data storage; GPU-as-a-Service (GPUaaS) partnership.
- **VAST Data** – Storage partner for enterprise AI.

AI / ML Platform & Technology Partners

- **DataRobot** – Validated AI Factory stack for agent workloads; co-announced at NVIDIA GTC 2026.
- **Nexla** – Production-ready data and agent stack; joint demo with Tripadvisor at NVIDIA GTC.
- **Eigen AI** – Inference optimization; co-developing optimized open-source models for Token Factory. Holds #1 output speed on multiple models per Artificial Analysis.
- **Anyscale (Ray)** – Managed Ray platform integration for multimodal and physical AI.
- **Outerbounds / Metaflow** – Strategic technology partnership; BYOC integration. Originally created at Netflix.
- **SkyPilot** – Open-source framework for AI workloads; integrated with Nebius AI Cloud.
- **dstack** – Open-source container orchestration platform; integration available.
- **Hugging Face** – Official inference provider integration.
- **Seqera Platform** (Nextflow creators) – Integration for bioinformatics workflows.
- **Saturn Cloud, Lightning AI** – Integration partners.
- **PyTorch** – Collaboration on frontier MoE training on NVIDIA Blackwell (demonstrated up to 41% faster DeepSeek-V3 pre-training).

Developer Tools & Ecosystem Integrations

- **Weights & Biases (W&B)** – Technology partnership; Launch agent on Marketplace.
- **DVC.ai** – Technology partnership for data version control.
- **Helicone, LlamaIndex, Postman, Agno, LangChain, CrewAI, Google ADK, LiteLLM** – Developer tool integrations.

Data & AI Services Partners

- **Toloka** – Data partner for AI development; Tendem MCP integration for human-in-the-loop AI agents. 10,000+ verified experts across 20+ domains. CEO: Olga Megorskaya.
- **Tavily** – Acquisition announced (agentic search). 3M+ monthly SDK downloads, 1M+ developer community. Serves IBM, Cohere, Groq.
- **Tripadvisor** – Joint "Inspiration to Trip" demo with Nexla and Nebius at NVIDIA GTC.

Distribution & Channel Partners

- **TD SYNnex** – Technology distributor; launched AI Infrastructure as a Service on Nebius.

Compliance, Benchmarking & Audit

- **Deloitte** – Independent auditor for SOC 2 Type II and HIPAA compliance.
 - **KRAI** – MLCommons founding member; supported MLPerf benchmark verification. Founder and CEO: Dr. Anton Lokhmotov.
-

Products

- **Nebius AI Cloud** – Full-stack cloud platform for AI workloads. Includes Compute Cloud (VMs with NVIDIA GPUs), GPU clusters with InfiniBand, managed Kubernetes, Slurm-based clusters (via Soperator), Object Storage (S3-compatible, Standard and Enhanced classes), Shared Filesystem, Block network storage, Container Registry, Virtual Private Cloud/networking, Identity and Access Management, Audit Logs, MysteryBox (secrets manager), monitoring/observability.
 - ****Neb**
-

Solutions

No relevant information found across any batch.

Services

No relevant information found across any batch.

Marketing Analysis

Source: nebius.md

Marketing – Strategic & Interpretive: <https://nebius.com>

Analyse a company's website for strategic positioning, brand identity, value propositions, and

competitive strategy. These topics require interpreting tone, patterns, and implied intent rather than extracting explicit facts.

Based on 748 stored pages analysed in 3 batch(es) with 8 AI calls (namespace: webscrape, model: claude-opus-4-6, profile: marketing analysis).

Positioning Statements

Nebius — Positioning Statements: Consolidated Summary

Primary Taglines & Hero Copy

- **"The ultimate cloud for AI innovators"** — primary hero tagline, used consistently across homepage and multiple landing pages.
- **"Built to democratize AI infrastructure and empower builders everywhere"** — supporting mission-level statement, paired with the primary tagline.
- **"The AI cloud company"** — core identity descriptor used in press releases, About page, and corporate communications.
- **"Purpose-built for AI. Powered for what's next."** — secondary tagline on AI Cloud page.
- **"Based in Amsterdam. Listed on Nasdaq. Operating worldwide."** — corporate credibility line.

Additional Positioning Statements

- **"Building the full-stack platform for developers and companies to take charge of their AI future — from data and model training to production deployment."**
- **"From building foundation models to scaling inference globally — ship AI faster, without having to manage infrastructure."**
- **"The combination of supercomputer performance with hyperscaler simplicity."** (press releases)
- **"Enterprise inference for open-source AI"** (Token Factory product)
- **"The ultimate AI cloud for startups"** (startup program)
- **"Don't just work in tech. Shape it."** (careers)

"We Are the X of Y" Positioning

- **Purpose-built AI cloud vs. adapted general-purpose clouds:** "Nebius has been built for AI since day one — not adapted from a general-purpose cloud, but designed for what developers actually need." — Arkady Volozh, CEO
- **Full-stack neocloud vs. hyperscalers AND bare-metal providers:** "Hyperscale clouds weren't built for the AI era — often complex, costly and inconsistent for performance at scale. On the other hand, bare-metal GPU providers may offer raw performance but lack the flexibility and native AI platform tools dev teams need."
- **Custom hardware leader among neoclouds:** "Nebius is the only neocloud that uses custom ODM chassis." — Dylan Patel, Chief Analyst, SemiAnalysis

Competitive Differentiators

Vertical Integration & Proprietary Stack

- "We design and operate our own custom NVIDIA-accelerated servers in energy-efficient data centers, giving us full control over quality assurance, performance tuning and delivery timelines."
- "A cloud computing platform rewritten from scratch in less than a year. We own every single part of the Nebius tech stack, and every line of code is all our own."
- Servers, racks, and greenfield data centers designed by in-house hardware R&D team.

Performance & Reliability

- **NVIDIA Exemplar Cloud Partner** status: demonstrated > 97% of performance benchmarks on NVIDIA H200 GPU clusters ($\geq 95\%$ of NVIDIA's reference architecture).
- **Gold medal** in SemiAnalysis GPU Cloud ClusterMAX™ Rating System — "officially on par with the hyperscalers — even surpassing some of our competitors."
- 167,000 GPU-hours MTBF on a 3,000-GPU cluster.
- Bare-metal-level performance: "We do not virtualize and share GPU and network cards"; minimizes virtualization overhead to maximize Model FLOPS Utilization (MFU).
- "Training that doesn't break" — fault-tolerant infrastructure with node health monitoring and auto-repair.
- Pre-configured NVIDIA GPU accelerators, high-performance InfiniBand, and Kubernetes/Slurm orchestration.

Speed of Delivery

- "AI Infrastructure — delivered fast. We provide access to NVIDIA accelerated computing clusters within hours, not weeks."

European & Global Firsts

- First cloud provider in Europe to deploy NVIDIA Blackwell Ultra systems in production.
- Europe's first GB300 NVL72 deployment on NVIDIA Quantum-X800 InfiniBand.
- ISEG2 supercomputer: **#13 on the Top500** of the world's fastest supercomputers (one batch cites #19; #13 appears to be the more recent ranking); "the most powerful commercially available system in Europe."

Cost & Value

- Published pricing: H100 at \$2.00/hr, H200 at \$2.30/hr, B200 at \$3.00/hr; Explorer Tier at \$1.50/hr for first 1,000 hours.
- Token Factory inference: **up to $3\times$ cost efficiency** (independently verified by Artificial Analysis); "more than $2.5\times$ cheaper than GPT-4o with comparable quality on Llama-405B."
- **Up to $4.5\times$ faster time-to-first-token** in Europe vs. other inference providers.
- Partner validation: "x5 lower costs compared to other major providers" (CentML case study).
- "By optimizing every layer of the stack, Nebius offers unparalleled efficiency, delivering substantial customer value over competitors."

Support

- 24/7 expert support and dedicated solution architects — "all free of charge."

Positioning Against Competitor Categories

Competitor type	Nebius's framing
Hyperscalers (AWS, Azure, GCP)	"Weren't built for the AI era — complex, costly, inconsistent at scale"; "quota roadblocks, cost volatility, hyperscaler friction."
Bare-metal GPU providers	"May offer raw performance but lack flexibility and native AI platform tools."
General-purpose clouds	Nebius is "built for AI since day one — not adapted."

Third-Party Validation Quotes

- "Nebius was the only provider that met our requirements." — Higgsfield AI
 - "x5 lower costs compared to other major providers." — CentML
-

Overall Brand Identity

Overall Brand Identity — Consolidated Summary

Mission & Vision

- **Core mission:** "Democratize AI infrastructure and empower builders everywhere."
- **Platform scope:** "A unified, full-stack platform spanning the complete AI journey — from data and model training and tuning to production runtime and deployment."
- **Positioning:** Founded on deep in-house technological expertise, operating at scale with a rapidly expanding global footprint, serving startups and enterprises building AI products, agents, and services worldwide.

Brand Name & Symbolism

- "Nebius" combines *nebula* (Latin for "cloud") with a reference to the infinite loop of the Möbius strip, embodying "the limitless possibilities created by our AI cloud."

Core Brand Values

- **Democratization:** Recurring commitment to making high-performance, scalable, cost-effective, and secure AI cloud infrastructure broadly accessible.

- **Builder-centric / developer-first:** "Built for the AI practitioners." Language throughout addresses "AI explorers," "AI innovators," and "AI builders." Developer-first simplicity is a product and cultural principle.
- **Engineering depth:** "Every line of code is all our own." In-house hardware design (custom server racks, ~20% more energy-efficient), proprietary software stack, ~400+ cloud engineers, own data center construction. Engineers "solve some of the hardest technical problems in AI."
- **Openness & transparency:** Open-source contributions (Soperator, Kvax, SWE-rebench on Hugging Face, contributions to PyTorch and vLLM). Transparent pricing. Public post-mortem incident reports, SEC filings, and sustainability reporting. "We are committed to openness and transparency."
- **Co-engineering partnership:** Positioned not as a vendor but a co-engineer — "Nebius felt like a co-engineer, not just a cloud provider" (Higgsfield AI). Free 24/7 support and dedicated solution architects.
- **Sustainability:** Free/natural cooling in Finland; waste heat reuse benefiting Mäntsälä households and a Paris urban farm; Iceland operations on 100% renewable energy; PUE of 1.13; 50+ GWh reused via district heating; published 2024 Sustainability Report; community benefits plan in Independence, MO.
- **Speed & momentum:** "At Nebius, the speed to impact really stands out." Fast-moving culture where bold ideas are turned into reality quickly.

Tone of Voice

- **Technical and authoritative but accessible:** Deep architecture deep-dives, benchmark analyses, and prompt-engineering tutorials sit alongside practical getting-started guides.
- **Direct and confident without being boastful:** Relies on concrete numbers, third-party validations (MLPerf, SemiAnalysis, Artificial Analysis, NVIDIA Exemplar), and engineering-focused customer quotes rather than abstract marketing claims.
- **Community-oriented:** Hackathons, Builder Days, research credits programs, AI Discovery Award (healthcare startups), Robotics & Physical AI Awards.
- **Occasionally playful:** April 1st digest with "hare rate" jokes; internal "Newbius" nickname.

Culture & Ethos

- **Strong engineering culture:** Dogfooding — the in-house AI R&D team is "the first client of Nebius."
- **Ownership and autonomy:** "New ideas are encouraged, experimentation is welcomed, and ownership is part of everyday work." Remote-friendly work culture built on trust and autonomy.
- **Rapid impact and continuous learning:** Emphasis on moving fast, trusting people, and enabling bold ideas. Mentorship-driven development; Early Talent Program with real product ownership, not simulations.
- **Global and multicultural:** HQ in Amsterdam; offices in San Francisco, Dallas, New York; R&D hubs across Europe, Israel, and North America; data centers in Finland, France, Iceland, and the US. Cultural markers include a sauna in the Finnish data center and cycling culture in Amsterdam.

Careers Messaging

- **Tagline:** "Don't just work in tech. Shape it."
 - **Positioning:** "At Nebius, opportunity spans both technology and business. We move fast, trust our people, and turn bold ideas into reality."
 - **Employee voice:** "Working at Nebius is intense and very rewarding." "A fast-moving culture where your voice can shape decisions."
-

Value Proposition Statements

Consolidated Summary: Nebius Value Proposition Statements

Core Value Propositions

- **Speed to production:** "We provide access to NVIDIA accelerated computing clusters within hours, not weeks, with pre-installed drivers, self-service access and engineering support every step of the way." "From zero to clusters in days." Provision training and inference environments rapidly.
 - **Operational simplicity:** "More AI. Less operations. Stay focused on what matters." Integrated observability, managed orchestrators (Kubernetes, Slurm), documented APIs, and managed services (MLflow, PostgreSQL, object storage) remove DevOps friction from the ML lifecycle. Full-stack platform covering data preparation, training, fine-tuning, and inference.
 - **Fault-tolerant training at scale:** "Training that doesn't break — Build foundation models easily with fault-tolerant infrastructure. With node health monitoring and auto-repair, Nebius ensures your training jobs continue to run — even at massive scale."
 - **Reliability/uptime:** "Best guaranteed uptime. Our platform provides a built-in self-healing system that allows VMs and hosts to restart within minutes instead of hours." Average MTTR of 12 minutes for multi-host training recovery.
 - **Bare-metal performance:** "Push your AI to the performance limit." Maximizes Model FLOPS Utilization (MFU) and delivers performance on par with leading industry benchmarks. Leading results in MLPerf® Training v5.1 benchmarks, including near-linear scaling ($1.97 \times$ speed increase when doubling from 512 to 1,024 GPUs).
 - **Cost efficiency:** "By optimizing every layer of the stack, Nebius offers unparalleled efficiency, delivering substantial customer value over competitors."
 - **Security by default:** "Scale safely in every regulated environment." HIPAA-, SOC 2-, GDPR-, ISO 27001-compliant, with privacy-focused architecture and tenant-level isolation as standard.
-

Token Factory (Inference API) Value Propositions

- "Run state-of-the-art open-source models with sub-second latency, predictable cost and zero-retention security, no MLOps required."
- **Up to $3 \times$ cost efficiency** — independently verified by Artificial Analysis; "prices are the best, according to AA" across all models offered.
- Handle **100M+ tokens per minute** with consistent throughput and **99.9% uptime SLAs**.

- **60+ open-source models** including DeepSeek, GPT OSS, Llama, Qwen, Mistral.
- Batch inference automatically billed at **50% of the base real-time model price**.
- Post-training (fine-tuning + distillation): **3–5× faster** inference, cutting inference costs and latency by up to 70%. Scales from 8 to 512 GPUs with no code changes.

Pricing

Resource	Price
H100 GPU	From \$2.00/hr (commitment); \$1.50/hr Explorer Tier for first 1,000 hrs/month
H200 GPU	From \$2.30/hr
B200 GPU	From \$3.00/hr
Commitment discounts	Save up to 35% on on-demand rates
Token Factory — Flux Schnell	\$0.0013/image
Token Factory — DeepSeek R1	\$0.80/1M input tokens
Training billing	"Pay only for the tokens processed during training, with no hidden costs"

Support Value

- **24/7 follow-the-sun support** across 3 global time zones — included at **no additional cost**.
- **85% of all issues resolved at first response**.
- **2.5 hr average resolution time** for non-escalated cases.
- **100% CSAT score** maintained since July 2025.
- Dedicated solution architects for multi-node cases, free of charge.

Quantified Customer Outcomes

Customer	Key Outcomes
Brave Search	Nearly 100% compute utilization; 11M+ AI-generated answers daily from 1.3B search queries/month

CRISPR-GPT (Stanford)	100% success rates for novice researchers; ~80% editing efficiency on first attempt; training time reduced from weeks/months to 1 day
CentML	5 × lower costs vs. other major providers; 1 week to get cluster online
Prosus	Up to 26 × cost reductions vs. proprietary models; up to 200B tokens/day
TheStage AI	4 × speed improvement over early framework version on A100
Higgsfield AI	40% cost efficiency gains with L40S; scaled from zero to \$200M run rate in 9 months; 15M+ users; 4.5M+ generations/day; time to first training run < 1 hour
Revolut	200+ H100 GPUs; 1.2M support tickets/month; 80% of chats resolved without human escalation; millions of fraudulent transactions prevented monthly
Chatfuel	24% routing performance improvement over GPT-4o; integration in 3 days
Recraft	Trained 20B-parameter model from scratch; V4 model: 350M+ images generated
Krisp	H100 training 50–80% faster than A100
SGLang + DeepSeek	2 × throughput boost and significantly lower latency on a single node
vLLM	Zero hardware-related issues; consistently accurate hardware performance metrics
RoboForce	Pipeline setup time cut by 70%+; iteration cycles reduced from weeks to days
Helical	500M-parameter model pre-trained in 10 days on 16 GPUs
TrialHub	250M-vector database built in days; 20 × faster Standard of Care research; 3 × faster patient recruitment
SieveStack	30–50% reduction in delays between test runs and production; up to 4 × training acceleration
Converge Bio	7,000+ hours of compute for training; processes 20,000+ genes per cell

Token Factory (general)	Sub-second latency, autoscaling throughput, 99.9% uptime
PyTorch collaboration	Up to 41% faster pre-training of DeepSeek-V3 on NVIDIA Blackwell GPUs

Startup & Developer Programs

- **AI Lift startup program:** Up to \$100K–\$150K in compute credits (sources vary between \$100K and \$150K; the higher figure appears on more recent pages).
 - **Introductory credits:** \$5,000 for new startups; \$1 in free Token Factory credits for all users to explore.
-

Ideal Customer Profile

Consolidated Ideal Customer Profile for Nebius

Explicitly Stated Targets

Nebius targets the full spectrum of AI compute consumers:

"From individual researchers, start-ups and scale-ups who want genuine on-demand, self-service access, through to enterprises and major AI developers who demand reserved capacity and large, interconnected superclusters."

Industries explicitly called out: **Healthcare & Life Sciences, Media & Entertainment, Physical AI & Robotics, Financial Services, Retail & E-commerce.**

Customer Segments

1. AI Natives / Startups

- **Named customers:** Recraft (20B-parameter model training), Higgsfield AI (video generation), Prime Intellect, Slingshot AI, Chatfuel, Potoroom, Wubble, TheStage AI, CentML, Decart, Helical, SieveStack, Converge Bio, Positronic Robotics, Simulacra AI, RoboForce
- **Startup program criteria:** VC/accelerator-backed, raised $\geq$ \$500K, building/scaling an AI-enabled product, officially incorporated < 7–10 years (sources vary between 7 and 10 years)
- **Incentives:** \$5K + up to \$100K in credits, co-engineering support, AI Lift program with NVIDIA Inception (22,000+ startups in ecosystem)

2. AI Enterprises

- **Named customers:** Meta (\$27B/5-year contract), Microsoft (multi-billion dollar agreement), Shopify, Monday.com, Prosus (Fortune 500), Revolut (70M+ users, £3.1B revenue), Brave Search (80M+ users), Milestone Systems
- **Key buying drivers:** Reserved capacity, 99.9% SLA, compliance (SOC 2 Type II, HIPAA, ISO 27001, GDPR, ISO 27799 for healthcare), RBAC, dedicated endpoints, enterprise SSO

3. Research & Academic

- **Named institutions:** Stanford (CRISPR-GPT), USC, London Institute for Mathematical Sciences, YerevaNN, ELIXIR BioHackathon
- **Engagement model:** Research credits program, Nebius Academy, AI Discovery Awards

4. Open-Source Ecosystem Builders

- **Named partners:** vLLM, SGLang — infrastructure framework developers testing and optimizing on Nebius

5. Hyperscale / Big Tech

- Microsoft and Meta as large-scale infrastructure customers, validating Nebius as a credible GPU cloud provider at massive scale

Industry Verticals (with evidence)

Vertical	Evidence
Healthcare & Life Sciences	12+ customer testimonials; NVIDIA BioNeMo NIM integration; customers include Prima Mente, Converge Bio, Helical, SimulacraAI, xAID, TrialHub, Quantori; Stanford CRISPR-GPT; AI Discovery Award
Physical AI & Robotics	Robotics Awards, Hannover Messe, ICRA presence; RoboForce, Positronic Robotics
Financial Services	Revolut (fraud prevention, FinCrime agents, credit scoring)
Media & Entertainment	Recraft, Higgsfield AI, Potoroom; content creation, localization, image/video generation
Retail & E-commerce	Shopify; explicitly listed vertical

Mapping to Buyer Personas

Alice (AI Platform Leader) — Strong fit

- **GPU scarcity:** Guaranteed capacity via Capacity Blocks, real-time Capacity Dashboard, self-service access to latest NVIDIA GPUs (H100 through GB300 NVL72)
- **Cost management:** Transparent per-hour/per-token pricing (\$2–5.50/GPU-hour), commitment discounts up to 35%, FOCUS-compliant billing exports, FinOps-ready
- **Vendor lock-in avoidance:** Open-source commitment (Soperator, SkyPilot, multi-cloud compatibility), Kubernetes/Slurm orchestration, Terraform provider
- **Compliance:** SOC 2 Type II, HIPAA, ISO 27001, GDPR
- **Developer experience:** Managed Kubernetes with GPU support, self-service console, API/CLI/SDK
- **Key quote (Milestone Systems):** *"We chose Nebius not just for their tech, but also for their commitment to data sovereignty."*

Carter (CTO / Technical Buyer) — Very strong fit

- **Performance validation:** MLPerf benchmarks, NVIDIA Exemplar Cloud status, no GPU virtualization
- **Capacity assurance:** Expanding to 5+ GW by 2030, strategic NVIDIA partnership (\$2B investment)
- **Vendor credibility:** Multi-billion dollar contracts with Microsoft and Meta; third-party validation from SemiAnalysis, Artificial Analysis
- **Cost predictability:** Transparent \$/token pricing, commitment discounts, build-vs-buy flexibility (managed vs. self-managed options)
- **Security/compliance:** Full certification suite

Cedric (CEO) — Moderate fit

- **Financial stability signals:** Nasdaq-listed (NBIS), \$700M+ investment round (NVIDIA + Accel), \$4.3B convertible notes, \$20B+ contract backlog, 479% YoY revenue growth
- **Board-ready information:** Quarterly earnings, Morgan Stanley/UBS conference appearances, sustainability/ESG report, economics whitepaper ("The economics of AI clusters")
- **Limitation:** Marketing is more technical than executive-focused; CEO Arkady Volozh quoted in major announcements but messaging skews toward practitioners

Diana (AI Leader / Head of AI) — Strong fit

- **R&D support:** Research credits, co-engineering, in-house AI R&D team publishing at ICML/ICLR, open-source contributions (SWE-rebench, Kvax)
- **Model flexibility:** Token Factory with 60+ open-source models, fine-tuning, distillation, post-training capabilities, OpenAI-compatible API for portability
- **Experiment tracking:** Managed MLflow
- **Regulatory compliance:** HIPAA, SOC 2, ISO 27799 (healthcare-specific)
- **Peer proof:** Case studies from Stanford, Revolut, vLLM

Malik (AI Researcher) — Strong fit

- **Compute access:** On-demand self-service up to 32 GPUs; reserved multi-node clusters scaling to 1,000+ GPUs
- **Research tools:** JAX-based Kvax training framework, Slurm/Kubernetes orchestration,

PyTorch pre-installed, Hugging Face integrations

- **Community:** Research credits program, Nebius Academy, open-source contributions (SWE-rebench, Soperator, Kvax)
- **Academic partnerships:** Stanford, USC, London Institute for Mathematical Sciences, YerevaNN, ELIXIR BioHackathon

Paul (Platform Engineer) — Strong fit

- **Infrastructure-as-code:** Terraform provider, CLI, API, SDK (Go and Python)
- **Orchestration:** Soperator (open-source Slurm-on-Kubernetes), managed Kubernetes with GPU support, SkyPilot and Ray integration
- **Observability:** Grafana dashboards, monitoring, audit logs, metrics/logs/tracing
- **Cost controls:** Capacity Blocks, FOCUS-compliant billing, real-time dashboard
- **Vendor independence:** Open-source tooling throughout the stack; container registry included

Sonia (AI Engineer) — Strong fit

- **Token Factory:** OpenAI-compatible APIs, 60+ open-source models, function calling, structured/guided JSON outputs, guardrails, batch inference (50% discount), LoRA hosting
- **Agentic AI:** Integration with LangChain, CrewAI, Google ADK, Tavily agentic search; tutorials for RAG, multi-agent systems, code generation
- **Developer onboarding:** Playground for testing, \$1 free credits, one-click fine-tuning and deployment
- **Security:** Zero data retention, SOC 2

Geographic Focus

- **Headquarters:** Amsterdam, Netherlands
- **Europe:** Data centers in Finland, France, Iceland, UK; strong data sovereignty positioning ("*guaranteeing that European customers' data can remain within Europe*")
- **North America:** Data centers in New Jersey and Missouri/Kansas City; expanding across the US
- **R&D hub:** Israel
- **Expanding into Asia-Pacific:** Singapore office, with stated expansion into Japan, South Korea, and India

Technical Maturity Spectrum Served

Nebius serves the **full range** of technical maturity:

- **Entry-level:** Self-service single-GPU access (Explorer Tier), \$1 free Token Factory credits
- **Growth-stage:** On-demand access up to 32 GPUs, startup credits up to \$100K, co-engineering
- **Enterprise-scale:** Thousand-GPU superclusters, reserved capacity, 99.9% SLA, dedicated endpoints, enterprise SSO, RBAC
- **Hyperscale:** Multi-billion dollar, multi-year infrastructure agreements (Microsoft,

Meta)

Both **AI Natives** (startups using open-source models, Hugging Face, vLLM) and **AI Enterprises** (compliance-first, managed services, data sovereignty) are explicitly and actively targeted.

Implied Marketing Strategy

Nebius Implied Marketing Strategy — Consolidated Summary

Core Positioning

Nebius is executing a **full-stack vertical integration strategy** to become the default AI-native cloud infrastructure provider, positioned between hyperscalers (AWS, Azure, GCP) and bare-metal GPU providers. The consistent framing is: hyperscalers are general-purpose clouds awkwardly adapted for AI, while Nebius is "purpose-built" — and specialized infrastructure outperforms general-purpose.

Strategic Pillars

1. Vertical Integration as Competitive Moat

- Nebius designs its own servers and racks (~20% more energy efficient), builds greenfield data centers, writes the entire cloud software stack from scratch, and operates its own facilities.
- This full-stack control enables performance optimization at every layer, competitive pricing, and rapid deployment of new NVIDIA hardware generations.
- Validated externally by SemiAnalysis: *"the only neocloud that uses custom ODM chassis, which translates to massively lower total cost of ownership."*
- Custom hardware R&D (patents portfolio, OCP Summit talks) creates durable differentiation that pure resellers or hyperscalers retrofitting general-purpose infrastructure cannot match.

2. NVIDIA Partnership as Structural Advantage

- Exceptionally deep, multi-dimensional relationship: \$2B NVIDIA investment, Exemplar Cloud Partner status, Reference Platform NCP, early access to Blackwell Ultra/GB200/GB300 (planning Vera Rubin), joint AI factory design engineering, NVIDIA Dynamo ecosystem partner, NVIDIA DGX Cloud Lepton provider, NVIDIA Inception collaboration, Elite sponsor at GTC, and featured in Jensen Huang's keynotes.
- Creates a **triple advantage**: supply chain security (guaranteed GPU access), time-to-market (among the first to deploy each new generation), and credibility (third-party validation positioning Nebius alongside hyperscalers).
- Generates a pull effect: customers seeking cutting-edge GPUs are drawn to Nebius as a preferred channel.

3. Anchor Tenant Strategy to Fund Growth

- Microsoft (~multi-billion) and Meta (~\$27B) infrastructure contracts serve dual purposes: predictable revenue to finance massive capex (gigawatt-scale data centers) and enterprise credibility validation.
- Meta press release explicitly states: *"the deal will also help us to accelerate the growth of our AI cloud business even further."*
- Classic "land with whales, expand to ecosystem" model.

4. Land-and-Expand from AI Natives to Enterprises

- **Bottom-up developer funnel:** Self-service model (\$25 minimum, instant access to up to 32 GPUs, no sales conversations), Explorer Tier (\$1.5/hr), startup programs (AI Lift with \$150K credits, \$5K credits + \$100K discounts), research grants, hackathons, and "Builder Hours."
- Mirrors the AWS/Snowflake playbook: make the initial experience frictionless, then capture growing workloads.
- As companies scale, Nebius captures increasing spend through reserved capacity, managed services, dedicated endpoints, and enterprise agreements — demonstrated by progression to Microsoft and Meta-scale contracts.

5. Dual Product Strategy — AI Cloud + Token Factory

- **AI Cloud** (training/fine-tuning): Captures high-spend, long-commitment workloads requiring technical depth.
- **Token Factory** (managed inference): Captures the rapidly growing inference market with per-token pricing and lower barriers to entry; supports 60+ open models with OpenAI-compatible APIs.
- Together they create a **natural lifecycle loop**: prototype on Token Factory → train/fine-tune on AI Cloud → deploy back on Token Factory.
- **Post-training pipeline** (fine-tuning → distillation → speculative decoding → deployment via Papyrax framework) creates stickiness across the full model lifecycle, reducing vendor sprawl and increasing switching costs.
- The **Tavily acquisition** (agentic search) extends the platform toward agentic AI workflows.

6. Open-Source Ecosystem Capture

- Heavy investment in open-source tools: Soperator (Slurm on K8s), Kvax (FlashAttention for JAX), SWE-rebench datasets.
- Deep integration with open frameworks: vLLM, PyTorch, Hugging Face, Anyscale/Ray, Outerbounds/Metaflow, SkyPilot, dstack.
- Token Factory exclusively serves open-source models (DeepSeek, Llama, Qwen, etc.), directly exploiting the shift from proprietary to open-weight models.
- Positions Nebius as aligned with the open-source AI community rather than proprietary lock-in — addressing vendor lock-in concerns across all buyer personas.
- Makes it easy to migrate from proprietary APIs (OpenAI, Anthropic) via OpenAI-compatible endpoints.

7. Geographic Expansion Riding Sovereign AI Trends

- Rapid buildout from Finland-only to: **Europe** (Finland owned DC, France, Iceland, UK/

Ark Data Centres), **US** (New Jersey 300MW, Independence MO gigawatt-scale/1.2 GW potential/\$650M+ PILOT payments, Kansas City), and **Asia-Pacific** (announced).

- Each new region opens access to local enterprises with data residency requirements and sovereign AI mandates.
- UK launch explicitly references the *"UK Government's AI Opportunities Action Plan"* and serving *"the NHS."*
- EU positioning leverages Amsterdam HQ, Finnish DC, and compliance certifications (GDPR, NIS2, DORA) to serve as the "European AI cloud."

8. Vertical Specialization in High-Value Industries

- **Healthcare & Life Sciences:** AI Discovery Awards, NVIDIA BioNeMo integration, multiple biotech customer stories, HIPAA compliance.
- **Physical AI & Robotics:** Dedicated awards program (\$1.5M in credits), presence at Hannover Messe and ICRA, partnership announcements.
- **Media & Entertainment:** Dedicated solutions positioning.
- Each vertical has specific compliance requirements and compute needs that Nebius's certifications and infrastructure directly address.

9. Community and Ecosystem Building

- Startup program, research credits, NVIDIA Inception partnership, Robotics & Physical AI Awards, Enterprise Readiness Initiative with VCs (Insight Partners, Accel, Fellows Fund), open-source contributions, hackathons, and Builder Hours.
 - Creates a **conversion funnel**: attract developers → prove value → scale to production → convert to enterprise contracts.
-

Market Dynamics Being Exploited

- **GPU scarcity:** Guaranteed capacity and availability is the #1 switching trigger; Nebius's capacity buildout and NVIDIA relationship directly address this.
- **Open-source model proliferation:** As open models approach proprietary quality, value shifts from model access to inference infrastructure — Token Factory's sweet spot.
- **Regulatory fragmentation:** EU data sovereignty (GDPR, NIS2, DORA) and sovereign AI mandates globally create demand for locally deployed, compliant AI compute.
- **AI-native vs. adapted cloud:** Growing recognition that specialized AI infrastructure outperforms general-purpose clouds retrofitted for AI.
- **Enterprise AI adoption:** Requires compliance (SOC 2, HIPAA, ISO 27001), governance, and cost predictability that general-purpose clouds lack.
- **Rising inference economics pressure:** As AI moves to production, predictable unit economics become critical.

Trends Being Ridden

- Explosive demand for AI training and inference compute (reasoning models like DeepSeek R1, Qwen QwQ driving more compute demand)
- Shift from proprietary to open-source AI models
- Rise of agentic AI systems requiring sustained inference, search grounding, and human verification (Toloka/Tendem integration for human-in-the-loop)

- Physical AI / robotics emergence requiring simulation-to-deployment pipelines
- AI in healthcare / drug discovery requiring specialized compliant infrastructure
- The "post-training" paradigm becoming essential for production AI economics
- Mixture-of-Experts architectures requiring optimized infrastructure

Composite Competitive Moat

The deepest moat is the **combination** of owned infrastructure (custom hardware, owned data centers) + NVIDIA strategic partnership (early hardware access, joint engineering, \$2B investment) + full-stack proprietary software (purpose-built cloud stack) + open-source ecosystem integration + geographic/compliance positioning. Any individual element is replicable; the simultaneous combination across all layers creates compounding advantages in performance, cost, and time-to-market that is extremely difficult to replicate.

Operational Metrics

Source: nebius.md

Marketing – Factual & Quantitative: <https://nebius.com>

Extract concrete, verifiable facts from a company's website: corporate details, infrastructure specs, customer names, product listings, and quantitative claims.

Based on 748 stored pages analysed in 2 batch(es) with 7 AI calls (namespace: webscrape, model: claude-opus-4-6, profile: operational_metrics).

Number Of Data Centers, Their Locations, Gpus, And Power Capacity

Nebius: Data Centers, Locations, GPUs, and Power Capacity — Consolidated Summary

Operational Data Centers

- Finland (Mäntsälä, eu-north1):
 - Own data center, ~60 km from Helsinki.
 - Houses the **ISEG supercomputer** (ranked #16 → later #19 on Top500) and **ISEG2** (#13 on Top500; described as "the most powerful commercially available system in Europe and the second most powerful commercially available supercomputer anywhere in the world").

- ISEG2 energy efficiency: **38.189 GFLOPs/watt** (#67 on Green500). PUE of **1.13** at full capacity.
 - Capacity being **tripled to 75 MW**, enabling placement of up to **60,000 graphic cards**.
 - **Europe's first operational NVIDIA GB300 NVL72 deployment** and first GB300 NVL72 on **NVIDIA Quantum-X800 InfiniBand** (800 Gbps).
 - Liquid cooling being installed for GB200 NVL72.
 - GPU types on-site: NVIDIA GB300 NVL72, GB200 NVL72, HGX B300, HGX B200, HGX H200, HGX H100, L40S.
- **France (Paris/Saint-Denis, eu-west1):**
 - Colocation at **Equinix PA10** campus in Saint-Denis.
 - First equipped solely with Nebius-designed servers; among the first in Europe/world to offer **NVIDIA H200 Tensor Core GPUs**.
 - Waste heat warms a rooftop urban farm.
- **Kansas City, Missouri (us-central1):**
 - Colocation with **Patmos** (converted Kansas City Star printing press).
 - Initial **5 MW**, expandable up to **40 MW** (~**35,000 GPUs** at full capacity).
 - Houses **NVIDIA Blackwell and Hopper H200 GPUs**.
- **Iceland (Keflavík):**
 - Colocation with **Verne**. **10 MW** compute cluster.
 - Runs entirely on **100% renewable** hydroelectric and geothermal energy.
 - NVIDIA H200 GPUs.
- **United Kingdom (Longcross Park, Surrey):**
 - Colocation at **Ark Data Centres**.
 - Initial deployment (expected Q4 2025) of **4,000 NVIDIA Blackwell Ultra GPUs** — one of the first in Europe.
 - Liquid-ready data-hall capacity; **NVIDIA Quantum-X800 InfiniBand** networking.
-

Planned / Under Construction

- **New Jersey, USA (Vineland):**
 - Greenfield data center built with **DataOne**. Expandable up to **300 MW** total.
 - First phase went live **summer 2025**. Nebius' first major US data center.
 - Behind-the-meter electricity and advanced energy technology. NVIDIA Blackwell GPUs.
 - Dedicated capacity for **Microsoft** starting later in 2025.
- **Independence, Missouri, USA:**
 - Approved (City Council) **gigawatt-scale AI factory** with potential capacity of **up to 1.2 GW**.
 - Multi-building campus on ~**400 acres**; connects to **Independence Power & Light (IPL)**.
 - ~1,200 construction jobs, ~130 permanent high-tech positions.

○ PILOT payments projected at over \$650M over 20 years.

GPU Portfolio Available on Nebius Cloud

GPU	Key Specs
NVIDIA GB300 NVL72	Blackwell Ultra, liquid-cooled, 72 GPUs per rack
NVIDIA GB200 NVL72	Blackwell, liquid-cooled, 72 GPUs per rack
NVIDIA HGX B300	Blackwell Ultra, 8 GPUs, 270 GB memory per GPU
NVIDIA HGX B200	Blackwell, 8 GPUs, 180 GB SXM per GPU
NVIDIA HGX H200	141 GB SXM per GPU
NVIDIA HGX H100	80 GB SXM per GPU
NVIDIA L40S	48 GB PCIe per GPU
NVIDIA RTX PRO 6000 Blackwell Server Edition	New addition (Aether 3.5 platform)

Aggregate Scale & Targets

- Over 22,000 NVIDIA Blackwell GPUs announced for deployment on Nebius AI-native cloud (Dec 2024).
- Thousand-GPU clusters available in both **Europe and the US**.
- Site-level maximums: Finland up to **60,000 GPUs**; Kansas City up to **35,000 GPUs**.
- Strategic target: deploy **more than 5 gigawatts** of NVIDIA systems by **end of 2030** (per NVIDIA partnership).

Active And Contracted Megawatts, Terawatts, And Other Measures Of Power

Consolidated Summary: Active and Contracted Power & Related Measures

Active / Planned Facility Capacities

- **Finland (Mäntsälä):** Tripling capacity to **75 MW**; operates at a PUE of **1.13** at full capacity
- **New Jersey (DataOne, Vineland):** Expandable up to **300 MW**
- **Kansas City (Patmos):** Initial **5 MW**, expandable up to **40 MW**
- **Iceland (Verne):** **10 MW** compute cluster, powered by 100% renewable hydroelectric and geothermal energy
- **Independence, Missouri:** Up to **1.2 GW** potential capacity (gigawatt-scale AI factory)

Strategic Deployment Target

- **By end of 2030:** More than **5 gigawatts** of NVIDIA systems deployed (per NVIDIA partnership announcement)
- **NVIDIA investment:** \$2 billion from NVIDIA as part of strategic partnership

Contract Backlog & Major Agreements

- **Total contract backlog:** Exceeding **\$20 billion**, including multi-year agreements with Microsoft and Meta
- **Microsoft deal:** Multi-year, multi-billion dollar agreement for dedicated AI infrastructure capacity from the Vineland, NJ data center
- **Meta deal:** \$12 billion of dedicated capacity across multiple locations over five years, plus up to \$15 billion in additional capacity — total contract value up to approximately **\$27 billion** over five years

Energy Efficiency & Reuse (Finland Operations)

- Servers consume **~20% less energy** vs. off-the-shelf servers and run reliably at temperatures up to **40°C**
- Over **50 GWh** of server waste heat reused for city district heating (2022–2024), supplying over **30 GWh** of heat to Mäntsälä and covering up to **72%** of the town's heating demand
- Nearly **10 GWh** of energy saved in 2024 from server efficiency improvements

Broader Industry Energy Context

- **GPT-3 training** consumed approximately **1,287 MWh**
 - Data centers accounted for roughly **1.5% of global electricity use** by 2024
 - U.S. data center electricity consumption projected at **325–580 TWh/year** by 2028
-

Installed, Active And Future Number Of Gpus

Consolidated Summary: Installed, Active, and Future Number of GPUs

Major Announced Deployments

- **Over 22,000 NVIDIA Blackwell GPUs** announced (Dec 2024) for deployment across data centers in the US and Finland.
- **Finland (Mäntsälä):** Up to **60,000 graphic cards** at the Mäntsälä location.
- **Kansas City:** Up to **~35,000 GPUs** at full potential capacity (40 MW).
- **UK:** Initial **4,000 NVIDIA Blackwell Ultra GPUs** (targeted Q4 2025).
- **10,000+ GPU cluster:** Decision to build a 10K GPU cluster mentioned as a natural progression of scaling efforts.

Currently Active / Demonstrated Clusters

- **3,000-GPU production cluster:** Achieved 169,800 GPU hours (56.6 hours) MTBF, as recorded by a customer (Recraft); separately cited as **167,000 GPU hours MTBF** for NVIDIA Exemplar Status.
- **MLPerf Training benchmarks** run on **512 NVIDIA H200 GPUs** (64-node cluster) and **1,024 NVIDIA H200 GPUs** (128-node cluster).
- **256 NVIDIA H200 GPUs** used by Prima Mente (interconnected via NVIDIA Quantum-2 InfiniBand).
- **256 NVIDIA HGX B200 GPUs** (32 nodes) used in Nebius–PyTorch collaboration.
- **ISEG supercomputer:** Part of the platform; capacity multiplied **5× since Nov 2023** ranking.

Self-Service / On-Demand Access

- Up to **32 NVIDIA H200 GPUs** on demand (US and Europe).
- Up to **16 NVIDIA H100 GPUs**, **8 NVIDIA H200 GPUs**, and **2 NVIDIA L40S GPUs** available via console per user.

GPU Models Available (Current & Planned)

- **Current:** NVIDIA GB300 NVL72, GB200 NVL72, HGX B300, HGX B200, HGX H200, HGX H100, L40S (Intel and AMD), NVIDIA RTX PRO 6000 Blackwell Server Edition.
- **Future:** NVIDIA Vera Rubin NVL72 planned from **H2 2026** in the US and Europe — among the first NVIDIA Cloud Partners to deploy.

Long-Term Target

- **More than 5 GW of NVIDIA systems by end of 2030**, implying tens of thousands of additional GPUs across future expansions.

Teraflops And Exaflops

Consolidated Summary: Teraflops and Exaflops

Supercomputer Rankings (ISEG / ISEG2)

- **ISEG2:** Ranked **#13** on the Top500 list of the world's fastest supercomputers; energy efficiency of **38.189 GFLOPs/watt** (#67 on the Green500).
- **ISEG:** Ranked **#16** (Nov 2023), later dropped to **#19** on the Top500. Exact FLOPS figures were referenced but not provided in the extracted sources.

NVIDIA Hardware Performance

- **GB200 NVL72:**
 - Delivers **1,000 tokens per second** per user, a **50 × boost in AI Factory output performance**, and **30 × improvement for real-time video generation**.
 - **130 TB/s** of GPU-to-GPU communication via fifth-generation NVLink.
 - Leading results in MLPerf Inference v5.1 (e.g., **855.82 tokens/s** for Llama 3.1 405B offline).
- **HGX B200:**
 - Up to **245,000 tokens per second** total throughput (8 × HGX B200 system).
 - **101,611 tokens/s** for Llama 2 70B in server mode.
- **HGX B300:**
 - Average **12.6% reduction in training time** vs. HGX B200.
- **H100 (Hopper):** Fourth-generation NVLink at **900 GB/s** per GPU; HBM3 memory at 81,559 MiB.

MLPerf Benchmark Results

Training (v5.1)

- **Seven first-place results** out of nine submissions (NVIDIA):
 - Llama-2-70B LoRA on 8 × HGX B300: **8.48 min** (1st place)
 - Llama-3.1-8B on 8 × HGX B300: **75.84 min** (1st place)
 - FLUX.1 on 32 × HGX B200: **93.17 min** (1st place)
 - Llama 3.1 405B: **124.5 min** (128-node cluster) and **244.6 min** (64-node cluster), demonstrating **1.97 × near-linear scaling**.
- **~3.1 × speed-up** scaling from 8 to 32 GPUs on HGX B200.

Inference (v5.1)

- Leading results on GB200 NVL72 and HGX B200 systems (see throughput figures above).

LLM Training Compute Reference

- A 175-billion-parameter LLM requires approximately **3.14 × 10²³ FLOPs** for full

training — illustrating the exaflop-scale compute demands of modern large language models.

Note: Neither batch provided direct teraflop or exaflop peak-performance specifications (e.g., exact TFLOPS per GPU or system-level exaflops ratings). The data primarily reflects benchmark throughput, training times, energy efficiency, and interconnect bandwidth rather than raw FLOPS ratings.

Other Capacity And Performance

Consolidated Summary: Other Capacity and Performance

Storage Performance

- **Shared filesystem (all-flash NVMe):** Up to **1 TB/s** read throughput; can scale performance linearly up to **4 PB**; up to **100 GBps** and **1M IOPS** for aggregated read operations
- **Object Storage:** Up to **2 GB/s per GPU**; Enhanced class offers up to **10 GB/s write throughput per client** and **~5x** more read/write throughput per TB
- **File storage per 8 GPU VM:** Write improved from 4 GBps to **8 GBps**; read from 8 GBps to **12 GBps**
- **WEKA-based storage:** **~20 GBps read / 18 GBps write** for an 8 GPU VM
- **Burst performance:** **500+ GB/s** aggregate read from a cluster with 254 GPU + 254 CPU hosts (measured during multi-modal training)
- **Expected next-gen:** **180 GBps** per NVIDIA GB200 NVL72 rack for filesystem read operations

Networking

- Up to **3.2 Tbit/s** InfiniBand per 8-GPU host (NVIDIA Quantum-2), non-blocking fabric
- **800 Gbps** end-to-end connectivity with NVIDIA Quantum-X800 InfiniBand (doubled throughput vs. prior generation)
- NCCL AllReduce bus bandwidth: **488 GB/s** in a two-node setup

Inference Performance (Token Factory / AI Studio)

- **Aggregate throughput:** Over **100M+ tokens per minute** capacity
- **DeepSeek R1:** **\$0.80 per 1M input tokens**; SGLang achieved **2x throughput boost** on a single node; DeepSeek R1 0528 reaches **248 output tokens/sec** (Top-2 worldwide per Artificial Analysis)
- **Qwen2.5 72B Fast:** **\$0.38/million tokens**, **70+ tokens/second** output
- **Eigen AI (GPT-OSS-120B):** Up to **911 output tokens/sec**
- **Prosus:** Up to **200 billion tokens per day** with autoscaling on dedicated endpoints
- **Brave Search:** **11M+ AI-generated answers daily** from **1.3B search queries/month**, nearly 100% compute utilization

- **Default rate limits:** Up to **400K tokens per minute** (200K–100K TPM per model)
- **Batch API:** Up to **5 million requests per file**, files up to **10 GB**, 500 concurrent batches
- **Token Factory SLA:** **99.9% uptime** for dedicated endpoints

GPU Pricing (On-Demand)

GPU	Managed/Contract	Self-Service
NVIDIA B200	\$3.00/hr	\$5.50/hr
NVIDIA H200	\$2.30/hr	\$3.50/hr
NVIDIA H100	\$2.00/hr	\$2.95/hr
NVIDIA L40S	From \$1.55/hr	—

- **3-year reserved H100:** \$2.12/hr
- **Commitment discounts:** Up to **35%** savings
- **Explorer Tier:** **\$1.50/hr** for first 1,000 H100 GPU-hours/month
- **JupyterLab:** Starting from \$2.95/hr (H100 configuration)

Reliability & Support

- **169,800 GPU hours** (56.6 hours) stable operation for a 3,000-GPU cluster; **167,000 GPU hours MTBF**
- More than **97%** of NVIDIA performance benchmark on H200 clusters (Exemplar Status); optimized to within **95%** of NVIDIA reference architecture
- Average incident response time: less than **10 minutes**; average MTTR for multi-host training: **12 minutes**
- Average support resolution time: **2.5 hours**; **85%** first-response resolution rate

Financial Metrics

- **479%** year-over-year revenue growth (2025)
- Targeting annualized run-rate revenue of **\$7–9 billion** by end of 2026
- Contract backlog exceeding **\$20 billion**
- **\$700 million** investment round (led by NVIDIA and Accel); **\$2 billion** NVIDIA investment
- **~\$4.3 billion** in convertible senior notes; **\$3.7 billion** cash position
- Listed on **NASDAQ: NBIS**
- **~400** cloud engineers; **500+** **AI & ML experts**

Compliance

SOC 2 Type II (with HIPAA), SOC 3, ISO 27001, ISO 27701, ISO 27018, ISO 27799, ISO 27032, ISO 22301, NIS2, DORA alignment

Named Customers

Revolut (200+ H100 GPUs, 70M users), Microsoft, Meta, Brave Software (80M+ users), Prosus, Hugging Face, Shopify, Recraft, Potoroom, Decart, Higgsfield AI/Mirage/Captions (15M+ users, 4.5M+ generations/day), Stanford (CRISPR-GPT), vLLM, SGLang, CentML, TheStage AI, Converge Bio, Basecamp Research, Compugen, Helical, Prime Intellect, Prima Mente, Simulacra AI, Positronic Robotics, Chatfuel, Quantori, TrialHub, RoboForce, Milestone Systems, Voxel51, Porsche, Monday.com, Tripadvisor, FlowGPT, Krisp, Dubformer, Unum, DataRobot, Nexla, Anyscale, Outerbounds, Seqera, SynthLabs, YerevaNN, Slingshot AI, SieveStack, xAID, Lynx Analytics, London Institute for Mathematical Sciences, and others.
