

Consolidated Report: nscale

Generated 2026-03-29 04:11 UTC from 3 analysis profile(s): company_profile, marketing_analysis, operational_metrics.

Note: The primary source of information in this report is the company's own website. The content reflects the company's self-reported claims and marketing materials and should not be treated as independently verified fact.

Executive Summary

Executive Summary: nscale

Company Overview

Nscale (NScale Global Holdings Limited) is a U.K.-headquartered AI infrastructure hyperscaler founded by CEO Josh Payne, who began building the company in Q3 2023 before launching from stealth in May 2024. Positioning itself as "the engine of superintelligence," Nscale offers a vertically integrated, full-stack AI cloud platform spanning power generation, data center operations, GPU compute, orchestration software, and managed AI services. The company has raised over \$3.7 billion across equity rounds—including a \$2 billion Series C at a \$14.6 billion valuation (the largest in European history)—plus a \$1.4 billion GPU-backed debt facility, and counts Microsoft, OpenAI, NVIDIA, and Aker ASA among its anchor partners and investors.

Key Findings

Infrastructure and Scale. Nscale operates or has contracted a rapidly expanding global footprint. Active and planned facilities include sites in Glomfjord and Narvik (Norway), Loughton (UK), Ward County (Texas), Sines (Portugal), and Keflavík (Iceland), with the flagship Monarch Compute Campus in West Virginia offering a pathway to over 8 GW of power by 2031. Near-term targets include 250 MW commissioned by Q4 2026 and scaling beyond 1 GW by 2029. GPU deployments are equally ambitious: 10,000 NVIDIA Blackwell GPUs to the UK by 2026, 100,000+ Vera Rubin GPUs to Europe in 2027, and plans to reach several hundred thousand GPUs within three years.

Strategic Positioning and Go-to-Market. Nscale explicitly differentiates itself from both legacy hyperscalers and neoclouds through vertical integration—owning power, facilities, hardware, and software—which it claims yields at least a 10% structural cost-of-production advantage and enables 80% lower pricing versus hyperscalers. The company pursues a dual go-to-market motion: top-down anchor deals (Microsoft's 1.35 GW letter of intent, OpenAI's Stargate Norway gigafactory) and bottom-up developer acquisition through serverless inference with pay-per-token pricing and free credits. A dedicated telco channel strategy (Singtel partnership, AI Grid playbook) targets distributed inference at scale. European data sovereignty—backed by UK/EU legal governance, GDPR/DORA compliance positioning, and the Trusted Tech Alliance co-founded with Anthropic, AWS, Google Cloud, and Microsoft—is a central selling point for government and regulated-industry customers.

Operational and Financial Metrics. ABI Research ranked Nscale the #1 neocloud overall across 14 providers, with perfect scores for distributed cluster scale and near-perfect marks for GPU availability and interconnect bandwidth. Performance claims include $7.2\times$ throughput/latency improvement via GEMM tuning, 30% faster time-to-insights, and a PUE target of ~ 1.1 at Glomfjord with 85% waste heat reuse. The inference platform has grown $148\times$ and serves 5,000+ users. The leadership team is notably senior, featuring former executives from Microsoft, Intel, Uber, AWS, Google, JP Morgan, and IBM, with a board that includes Sheryl Sandberg, Susan Decker, and Nick Clegg.

Cross-Cutting Themes, Strengths, and Risks

Several strengths emerge across all three dimensions: the depth of vertical integration from energy assets to application-layer services; the caliber and credibility of anchor partnerships with NVIDIA, Microsoft, and OpenAI; a sustainability narrative grounded in verifiable renewable-energy operations; and an experienced leadership bench recruited from tier-one technology and financial institutions. The sovereign AI positioning is particularly well-timed given accelerating European regulatory momentum.

Key risks warrant attention. Nscale's ambitious capacity targets (8 GW by 2031) require flawless execution across construction, permitting, power procurement, and capital deployment—any of which could slip. The Monarch campus relies on natural gas generation, creating tension with the company's renewable-energy brand. Many headline metrics (GPU counts, megawatt targets, cost-savings claims) are forward-looking or self-reported and await independent verification at scale. Concentration risk around a small number of anchor customers—particularly Microsoft—is significant. Finally, the competitive landscape is intensifying as both legacy hyperscalers and well-funded neoclouds expand AI-specific capacity.

Strategic Outlook

Nscale has assembled a compelling combination of capital, partnerships, physical infrastructure pipeline, and strategic narrative at a pivotal moment in AI infrastructure demand. If the company can execute on its construction and deployment roadmap—particularly the Monarch campus and Vera Rubin rollouts—while converting its anchor relationships into durable, recurring revenue, it is well-positioned to establish itself as a top-tier AI infrastructure provider bridging the gap between hyperscalers and neoclouds. The next 18–24 months will be decisive: delivering commissioned megawatts on schedule, demonstrating sustained financial performance behind the headline valuations, and proving that vertical integration translates into the reliable, cost-effective service that enterprise and sovereign customers require at scale.

Company Profile

Source: nscale.md

Marketing – Factual & Quantitative:
<https://nscale.com>

Extract concrete, verifiable facts from a company's website: corporate details, infrastructure specs, customer names, product listings, and quantitative claims.

Based on 93 stored pages analysed in 1 batch(es) with 1 AI calls (namespace: webscrape, model: claude-opus-4-6, profile: company_profile).

Headquarters

Nscale appears to be U.K.-based/headquartered. The Series C press release (March 9, 2026) describes Nscale as a "U.K.-based AI infrastructure hyperscaler." The legal entity is "NScale Global Holdings Limited." The company opened a London office at 16 New Burlington Place, Mayfair, London. The company also has a significant presence in Norway (Nscale AS mentioned in Terms & Conditions). The Nscale Energy & Power division is headquartered in Houston, Texas.

Years In Business

Josh Payne stated in the Series A blog (December 2024) that he began building Nscale in Q3 2023 and that the company had "spent the last 3 years consolidating a pipeline of 1.3 GW of greenfield data centre sites." Nscale launched from stealth in May 2024. The Glomfjord hydropower plant was "built at the end of the fjord over 100 years ago." The company's earliest referenced blog post dates to February 10, 2024.

Company Executive Names And Titles

- **Josh Payne** — Founder & CEO. Entrepreneur with companies across recruitment, capital markets, and supercomputing infrastructure. Deep experience in energy, data centres, and hardware supply chains.
- **Alex Sharp** — President, Data Center Group. Chartered Electrical and Electronic Engineer with 30+ years of experience. Led design/construction of 54 data centres including some of Europe's largest. Contributed to London Stock Exchange data infrastructure and European DC standards.
- **Philippe Sachs** — President EMEA & Global Affairs. 25 years in banking (Standard Chartered Bank Global Public Sector), CEO of educational company with 500+ UK employees, blockchain/deep tech VC who led early investment rounds in Nscale.
- **Nidhi Chappell** — President AI Infrastructure. Former Head of AI Infrastructure at Microsoft (managed 2M GPUs), former Senior Director of Data Centres at Intel.
- **Jing Yin** — President APAC. Leadership roles at Alibaba Cloud, Ant Group, and Lazada across cloud, payments, and e-commerce in Asia.
- **Phoebee Gahan** — Chief Legal Officer. Experience in corporate law, capital markets, M&A, finance, technology, and corporate governance.
- **David Power** — Chief Technology Officer. 20+ years in HPC. Designed/delivered supercomputing systems. Decade on Irish Supercomputing List advisory board.
- **Daniel Bathurst** — Chief Product Officer. 10+ years launching platforms across cloud, finance, AI, and Web3.
- **Ron Huisman** — Chief Accounting & Financial Ops Officer. 25+ years in finance, capital markets, M&A. Previously senior finance at Liberty Global and CFO at AtlasEdge Data Centres.

- **Tom Burke** — Chief Revenue Officer. 15+ years in global enterprise sales. Previously CRO at Applied Digital (NASDAQ: APLD), 15+ years at IBM as VP and Head of Global Sales for Data & AI.
- **Alice Takhtajan** — Chief Financial Officer. 25+ years at JP Morgan, led 150+ equity transactions in Tech sector.
- **Naeem Hatimi** — Chief People Officer. Previously at SharkNinja, Netflix, Vodafone.
- **Karl Havard** — Chief Commercial Officer. 25 years in cloud and AI infrastructure, leadership at AWS, Google, and startups.
- **Latha Maripuri** — Chief Information Security Officer. 25+ years in cybersecurity. Former CISO at Uber, former Global CISO and Deputy CTO at News Corp, product/technical leadership at IBM.
- **Daniel Shapiro** — Chief Power Officer (Nscale Energy & Power).
- **Bengt Jarlsjo** — Deputy Chief Power Officer.
- **Steve Relf** — VP of Operational Engineering.
- **Brian Dervan** — VP of Alliances and Partners.
- **Daniel Bergli** — Director, Industry Partnerships.

Board of Directors: Josh Payne, Rael Nurick, Jacob Leschly, Øyvind Eriksen (CEO of Aker ASA), Sheryl Sandberg (co-founder Sandberg Bernthal Venture Partners, former COO Meta), Susan Decker (CEO of Raft, former President Yahoo, board member Costco/Berkshire Hathaway/Vail Resorts), Nick Clegg (General Partner Hiro Capital, former UK Deputy PM, former President Global Affairs at Meta).

Customer Names, Locations, Industry/Type Of Entity, And Other Relevant Details

- **Microsoft** — Major customer/collaborator. Letter of intent for up to 1.35GW of AI compute using NVIDIA Vera Rubin NVL72 GPUs at West Virginia campus. Quote from Jon Tinter, President, Business Development and Ventures at Microsoft: "Together with Nscale, Microsoft is delivering cutting-edge AI infrastructure for our customers."
 - **OpenAI** — Partner on Stargate Norway (Narvik), the first AI gigafactory from OpenAI in Europe under its "For Countries" programme. Sam Altman quoted praising the partnership.
 - **Singtel** — Strategic partnership to unlock GPU capacity across Europe and Southeast Asia. Singtel leverages Nscale's AMD and NVIDIA GPU capacity in Europe; Nscale taps Singtel's NVIDIA H100 capacity in Southeast Asia.
 - **NORSAR** — Referenced as a customer/user; Birger Steen, CEO of NORSAR (research foundation in seismology), quoted in survey about sovereignty.
 - **Serverless inference platform** serves 5,000+ users.
 - Industries mentioned as target customers: finance, healthcare, biotech, telecommunications, public sector/government, software & technology, insurance, retail.
 - ABI Research ranked Nscale #1 neocloud (Overall Leader, Top Innovator, Top Implementer) among 14 providers.
 - Nscale states it runs "tens of thousands of GPUs globally for enterprise customers."
-

Partner Names, Locations, Industry/Type Of Entity, And Other Relevant Details

- **NVIDIA** — GPU partner. Nscale will deploy NVIDIA Vera Rubin platform in 2027 (100,000+ GPUs to Europe). Uses NVIDIA H100, H200, GB200, Blackwell Ultra GPUs. NVIDIA invested in Series C. Quote from Ian Buck, VP and GM Hyperscale and HPC, NVIDIA.
- **Microsoft** — Collaboration partner for West Virginia AI Factory and Vera Rubin deployment across UK/Norway.
- **Aker ASA** — Norwegian industrial conglomerate. Joint venture partner (later rolled into Nscale). Leading shareholder. Partner on Stargate Norway (Narvik). CEO Øyvind Eriksen is on Nscale board.
- **OpenAI** — Partner on Stargate Norway under "OpenAI For Countries" programme.
- **Caterpillar** — Deploying G3500 series natural gas generator sets at Monarch Compute Campus (West Virginia) for 2GW power generation.
- **Dell** — Partner (also Series C investor).
- **Nokia** — Networking partner. Joint innovation across switching and optical layers. Series C investor.
- **VAST Data** — Storage partner.
- **Verne** — Iceland data center partner. Keflavik site hosts 4,600+ NVIDIA Blackwell Ultra GPUs across Verne's Icelandic campus.
- **Lightning AI** — Partnership to launch enterprise-grade AI Studio on Nscale cloud. Access planned Q3 2025.
- **Hugging Face** — Nscale is an official Inference Provider on Hugging Face.
- **Ionic Digital** — Partner for Ward County, Texas data center (~240MW, plans to 1.2GW).
- **Start Campus** — Partner for Sines, Portugal hyperscale site.
- **Armada** — Signed LOI for global hyperscale and edge AI deployments. San Francisco-based modular edge infrastructure company.
- **Singtel / Nxera** — Strategic GPU capacity partnership across Europe and Southeast Asia.
- **Lenovo** — Series C investor. Presented jointly at SuperComputing 2024.
- **AMD** — GPU partner (MI250X, MI300X). Sponsored Nscale at SuperComputing 2024.
- **Fidelis New Energy** — Sponsor of AIPCorp (acquired by Nscale).
- **8090 Industries** — Sponsor of AIPCorp; investor in Series C.
- **Broadcom** — Networking (Ethernet fabric at Glomfjord).

Trusted Tech Alliance members (co-founded with Nscale): Anthropic, AWS, Cassava Technologies, Cohere, Ericsson, Google Cloud, Hanwha, Jio Platforms, Microsoft, Nokia, NTT, Rapidus, Saab, SAP.

Series C investors: Aker ASA, 8090 Industries, Astra Capital Management, Citadel, Dell, Jane Street, Lenovo, Linden Advisors, Nokia, NVIDIA, Point72. Goldman Sachs & Co. LLC and J.P. Morgan acted as joint placement agents.

Series A investors: Sandton Capital Partners (lead), Kestrel 0x1, Blue Sky Capital Managers Ltd, Florence Capital.

GPU DDTL (\$1.4B) lenders: PIMCO, Blue Owl, LuminArx Capital Management. Goldman Sachs as Sole Structuring/Placement Agent.

Advisors on AIPCorp acquisition: Morgan Stanley (financial advisor to Nscale), Latham & Watkins (legal to Nscale), Greenhill/Mizuho (financial advisor to AIPCorp), Vinson & Elkins (legal to Fidelis).

Products

AI Services:

- **Inference Endpoints** — Run models via API. Autoscaling inference layer backed by Nscale-managed GPU clusters. Supports 100+ open-source models. OpenAI API compatible.
- **Prompt Workbench** — Browser-based tool for testing and iterating prompts with versioned prompts and real-time model feedback.
- **Fine-tuning** — Serverless, API-driven fine-tuning. Pay-as-you-train, token-based pricing (\$0.50/1M tokens for models up to 20B). Supports models including Mistral 7B, Qwen 2/2.5, DeepSeek R1 distill variants, Meta-Llama-3-8B. Export to PyTorch, ONNX, or Hugging Face.
- **Serverless Inference** — Pay-per-request model inference with no infrastructure management. Models include GPT OSS 120B/20B, Devstral Small, Qwen 3 235B, Qwen 2.5 Coder series, Llama 4 Scout, Llama 3.x, DeepSeek R1 distills, Mixtral 8x22B, Flux.1, Stable Diffusion XL, SDXL Lightning. Pricing from \$0.01 to \$1.2 per 1M tokens (varies by model). \$5 free credits for new users. Zero rate limits, no cold starts.

Platform Services:

- **Managed Slurm** — HPC-grade batch scheduling (NVIDIA Slinky). Slurm on Kubernetes.
- **Nscale Kubernetes Service (NKS)** — Provision isolated Kubernetes environments in under 2 minutes. GPU-aware scheduling, autoscaling, enterprise-grade security.
- **Instances** — Virtual Machines or Bare Metal nodes with prebuilt AI images and optional VPC isolation.

Infrastructure Services:

- **Compute** — On-demand GPU & CPU. Bare-metal nodes on NVIDIA H100, H200, GB200 NVL72, AMD MI250X, MI300X GPUs.
- **Storage** — Parallel, AI-optimised storage tiers with GPU-tuned distributed file systems.
- **Networking** — RDMA/InfiniBand/NVLink fabrics, multi-rack topology, low-latency interconnects.

Fleet Operations:

- **Control Center** — Unified lifecycle manager for provisioning, scaling, patching, node health tracking, remediation.
- **Observability** — Telemetry across compute, storage, networking with dashboards, alerts, integrated reporting.
- **Radar API** — Real-time GPU resource governance, repair visibility, availability/maintenance metrics via unified API.

Other Products (listed in footer):

- GPU Nodes
 - Training Clusters
 - AI Marketplace (curated tools: PyTorch, TensorFlow, pre-trained models, preconfigured hardware)
-

Solutions

- **Model Training** — GPU clusters for LLM training with Managed Kubernetes and Slurm orchestration.
 - **Model Fine-Tuning** — GPU cloud solutions for fine-tuning AI models for peak performance. Serverless fine-tuning service.
 - **AI & ML Inference** — GPU-accelerated nodes for efficient inference. Batch and streaming workloads. Claims 80% lower cost vs. hyperscalers, 7.2x faster throughput/latency with GEMM tuning, 30% faster time to insights.
 - **AI Development** — End-to-end AI development platform from Jupyter notebooks to multi-GPU cluster deployment.
 - **Sovereign Cloud** — European-domiciled full-stack AI cloud with data sovereignty, UK/European law governance, sovereign controls.
 - **AI Grid** — Distributed AI infrastructure fabric for Telcos connecting centralized factories, regional hubs, and edge nodes. Described as a "2026 playbook for Telco AI" targeting the "\$118B opportunity."
 - **Telco Solutions** — Telco-branded AI services including GPU-as-a-Service, real-time inference, sovereign AI. Use cases: finance/regulated enterprise, public safety/government, healthcare, retail.
 - **Healthcare & Biotech** — Drug discovery, genomics, life sciences simulations, bioinformatics.
 - **Finance & Insurance** — Financial modelling, Monte Carlo simulations, fraud detection, AI-powered customer service.
 - **Software & Technology** — LLM building, computer vision, drug discovery acceleration, cybersecurity.
-

Services

Professional/Managed Services:

- Managed Slurm (HPC-grade batch scheduling as a service)
- Managed Kubernetes Service (NKS)
- Fleet Operations (automated GPU lifecycle management, health monitoring, remediation)
- Serverless Inference (fully managed inference platform)
- Serverless Fine-tuning (fully managed fine-tuning platform)

Support:

- Contact support available (help@nscale.com)
- Expert team for onboarding and questions
- Detailed documentation at docs.nscale.com
- Trial period available before commitment

Consulting/Expert Services:

- "Expert Assistance — Leverage our team of experts to ensure a smooth and successful transition"
- GPU optimization expertise ("access to GPU experts who can optimise AI workloads, maximise utilisation and ensure scalability")

SLAs/Performance Claims:

- 80% lower cost vs. hyperscalers
- 30% faster time to insights
- 40% improvement in resource utilisation
- 7.2x improvement in throughput and latency with GEMM tuning
- Kubernetes environments provisioned in under 2 minutes
- ABI Research scored Nscale 9/10 for GPU availability ("most on-demand jobs typically start within minutes")
- ABI Research scored 10/10 for maximum distributed cluster scale
- ABI Research noted Nscale's optimization stack "can deliver 20 to 40 percent better throughput without changing hardware"
- Inference service grown 148x, serving 5,000+ users
- Glomfjord PUE target: ~1.1
- 85% waste heat reuse target at Glomfjord

Training/Education:

- Approved apprenticeship company in Nordland County, Norway
- Student programs (16-18 year olds) at Glomfjord data center
- Plans to hire and train 100 AI specialists in the UK over 12 months

Pricing:

- Serverless Inference: Pay-per-token (varies by model, e.g., \$0.01-\$1.2/1M tokens). \$5 free credits for new users.
- Fine-tuning: \$0.50/1M tokens for models up to 20B parameters. \$5 free credits.
- GPU Nodes: Contact sales for pricing.

Marketing Analysis

Source: nscale.md

Marketing – Strategic & Interpretive: <https://nscale.com>

Analyse a company's website for strategic positioning, brand identity, value propositions, and competitive strategy. These topics require interpreting tone, patterns, and implied intent rather than extracting explicit facts.

Based on 93 stored pages analysed in 1 batch(es) with 1 AI calls (namespace: webscrape, model: claude-opus-4-6, profile: marketing analysis).

Positioning Statements

Primary tagline: "The engine of superintelligence" — used as the hero headline across the website.

Supporting tagline: "Full-stack AI infrastructure powering the world's most powerful systems, from ground to cloud."

Key positioning statements:

- "The hyperscaler engineered for AI" — used repeatedly across blog posts, press releases, and CEO communications.
- "The antithesis to the Neoclouds" — CEO Josh Payne explicitly contrasted Nscale's model against neocloud competitors in his Series A blog post.
- "The first AI-native hyperscaler" — used on the careers page as the company mission.
- "The first true hyperscaler, purposely engineered for AI" — Karl Havard, COO.
- "The only full stack sovereign AI infrastructure provider in the UK" — Josh Payne, CEO, in response to UK government compute strategy.
- ABI Research ranked Nscale as "#1 Overall Leader, Top Innovator, and Top Implementer" in their neocloud provider assessment of 14 companies.

Competitive differentiators explicitly stated:

- **Vertical integration:** "Nscale is the hyperscaler engineered for AI, the antithesis to the Neoclouds... ONLY way to do so is to build a fully vertically integrated platform. One that owns the Data Centre, the GPU cluster and the software."
 - **Cost advantage:** "Our cost of production for compute is at least 10% lower than our competitors. Enabling Nscale to always offer a lower-cost service at any scale." Also: "80% cost-saving in comparison to hyperscalers."
 - **Sustainability:** "Powered by 100% renewable energy" (Glomfjord), PUE targeting ~1.1.
 - **Sovereignty:** European-domiciled, UK/European law governance, data sovereignty by design.
 - **Scale:** "ABI Research scored Nscale a perfect 10 for maximum distributed cluster scale."
 - **Against neoclouds:** CEO described average neoclouds as "80-90% leveraged with an average cost of capital of 15 + %," not owning data centers, licensing third-party software, relying on spot instances — "thin margin, high risk."
-

Overall Brand Identity

Mission: "Build the first AI-native hyperscaler, empowering innovators with high-performance, scalable infrastructure." Also: "To power the next generation of intelligence with unprecedented performance and efficiency for people everywhere."

Vision: CEO Josh Payne: "This is the fourth industrial revolution... Nscale is leading this buildout. We are building this foundation that the market sits on, the engine of superintelligence."

Brand values (explicitly stated on About and Careers pages):

- **Invent Boldly:** "We push boundaries without fear, creating products that move the world forward."
- **Own the Outcome:** "We take full responsibility for our work, holding the bar high and delivering with urgency."
- **Move Fast, Stay Humble:** "We value speed paired with openness to learn—no egos, only progress."
- **Build with Trust:** "We operate transparently, earning trust through clear communication and consistent delivery."

Additional cultural values from Careers page:

- Relentless Innovation
- Openness and Transparency ("Our open-source approach allows customers to explore our technology")
- Sustainability ("integrating sustainability into our operations and product development")
- Ownership and Accountability
- Customer-Centric Focus
- Full-Speed Collaboration

Tone of voice patterns: Technical authority combined with bold ambition. Language is infrastructure-focused, engineering-first, but elevated with aspirational framing ("superintelligence," "fourth industrial revolution," "engine of superintelligence"). Community-oriented when discussing Glomfjord and Norway operations. Direct and confident when discussing competitive positioning — CEO openly critiques neocloud model economics.

Cultural ethos: "We're builders shaping the future of AI — and how we build defines us." Strong emphasis on community impact (Glomfjord apprenticeships, heat reuse for fish farms and swimming pools, local job creation). Norwegian Transparency Act compliance highlighted. Joined the Trusted Tech Alliance alongside Anthropic, AWS, Google Cloud, Microsoft, Nokia, and others.

Value Proposition Statements

Primary value proposition: "Deploy AI on infrastructure designed for scale, resilience, and speed." Also: "Reduce costs, grow revenue, and run your AI workloads more efficiently on a fully integrated platform."

Quantified outcomes and claims:

- "80% LOWER COST — Nscale delivers on average 80% cost-saving in comparison to hyperscalers."
- "30% FASTER — Nscale Cloud accelerates time to insights by up to 30% thanks to its AI-optimised stack."
- "+ 40% EFFICIENCY — Up to 40% improvement on efficiency."
- "7.2X FASTER — AMD MI300X GPUs with GEMM tuning improves throughput and latency by up to 7.2x."
- "4 out of 5 developers ranked us as the most cost-effective GenAI inferencing provider."
- ABI Research noted mature optimization stack "can deliver 20 to 40 percent better throughput without changing hardware."
- NVIDIA GB200 NVL72: "4x faster training," "25x more energy efficiency," "30x faster real-time LLM inference," "18x faster data processing."
- "Inference service has grown 148x, now serving more than 5,000 users."
- "Our cost of production for compute is at least 10% lower than our competitors."

Service-level value propositions:

- **Inference Endpoints:** "Launch inference in minutes while eliminating cluster management overhead."
- **Fine-Tuning:** "Customize frontier models to your domain quickly and efficiently, reduce reliance on generic models, and accelerate the path from POC to production."
- **Prompt Workbench:** "Experiment and optimize prompts quickly without burning GPU"

hours."

- **Compute:** "Maximise GPU efficiency and utilisation to lower cost per run and accelerate experiments."
- **Storage:** "Prevent slowdowns that delay product launches by ensuring predictable throughput."
- **Networking:** "Scale training from dozens to thousands of GPUs with no network bottlenecks."
- **Control Center:** "Cut operational overhead and maximize GPU efficiency and utilization."
- **Observability:** "Gain end-to-end visibility into workloads to ensure predictable performance, cost accountability, and regulatory compliance."
- **Kubernetes Service:** "Provisioning isolated Kubernetes environments in under two minutes."
- **Serverless Inference:** "No cold starts, no rate limits, no waiting," "Your data stays yours — always," "We never log, repurpose or train on your request or response content."

Industry-specific value propositions:

- Healthcare: "Accelerated analytics... improve diagnostic accuracy and reduce patient wait times."
 - Finance: "Real-time fraud detection and transaction monitoring with sub-30ms inference."
 - Telcos: "Nscale gives Telcos a fast, low-risk path to monetize AI by activating existing infrastructure as a distributed AI inference network."
-

Ideal Customer Profile

Explicit target segments mentioned across the website:

1. **Global hyperscalers and frontier AI labs:** Microsoft (letter of intent for 1.35GW), OpenAI (Stargate Norway partner), and references to "teams from Microsoft, OpenAI, and others who train frontier models on our clusters." This is the highest-profile segment.
2. **LLM platforms / AI model builders ("Model Makers"):** Companies needing large GPU clusters (10,000+) for training foundation models. CEO blog: "the larger customers... increasing exponentially... large contiguous power availability."
3. **Enterprises deploying AI:** Healthcare, finance/insurance, software & technology companies. Explicit solution pages for these industries. Language emphasizes governance, compliance, DORA, EU AI Act, GDPR.
4. **Governments requiring sovereign compute:** "Governments requiring large-scale sovereign compute capacity." UK sovereign AI forum, Stargate Norway as "sovereign AI" project, Trusted Tech Alliance membership.
5. **Telcos:** Dedicated Telco page and AI Grid playbook. Partnership with Singtel. "Telcos are key to a global AI distribution network." Targeting telco operators to become AI infrastructure providers.
6. **Developers and startups:** Serverless inference with \$5 free credits, pay-as-you-go, OpenAI API compatibility. Blog tutorials on Cursor IDE integration, fine-tuning

walkthroughs.

Best persona matches from provided profiles:

- **Alice (AI Platform Leader):** Strong match. Nscale addresses her core concerns: GPU scarcity (guaranteed capacity), cost management (80% lower cost claims, observability, cost accountability tools), vendor lock-in avoidance (open-source model support, OpenAI API compatibility, chip-agnostic infrastructure), security/compliance (sovereign cloud, GDPR, DORA content), and developer experience (Kubernetes in 2 minutes, serverless inference).
- **Carter (CTO/Technical Buyer):** Strong match. Nscale speaks directly to infrastructure strategy, build-vs-buy decisions, GPU capacity guarantees, platform maturity (ABI #1 ranking), transparent pricing, and partner ecosystem depth (NVIDIA, Dell, Nokia, VAST, Microsoft). The full-stack pitch addresses his need for system integration and reducing vendor sprawl.
- **Cedric (CEO):** Moderate-strong match. Board-ready scale (Stargate Norway, \$2B Series C, \$14.6B valuation), sustainability messaging for ESG, cost management tools, and security/compliance. The "engine of superintelligence" framing and "fourth industrial revolution" language is CEO-level strategic narrative.
- **Paul (Platform Engineer):** Strong match. Bare metal GPU nodes, Slurm/Kubernetes orchestration, Fleet Operations (observability, automated lifecycle management), open-source compatibility, cost transparency, and detailed engineering blog content (distributed fine-tuning, benchmark results).
- **Diana (AI Leader):** Good match. Fine-tuning service, inference endpoints, model customization, sovereign cloud compliance, and vendor reputation (ABI ranking, Microsoft/OpenAI partnerships).
- **Malik (AI Researcher):** Moderate match. Training clusters, Slurm scheduling, distributed fine-tuning (DeepSeek-R1 walkthrough), compute availability, but less content on research-specific evaluation suites or model alignment tooling.
- **Sonia (AI Engineer):** Good match. Serverless inference, OpenAI API compatibility, prompt workbench, fine-tuning, pay-as-you-go, Hugging Face integration, low-latency inference.

Geographic focus: Europe (Norway, UK, Portugal, Iceland) and North America (Texas, West Virginia, North Carolina). Expanding into Asia-Pacific (Singtel partnership, Jing Yin as President APAC). European sovereign AI is a primary positioning thrust, with US expansion accelerating rapidly.

Company size implied: Ranges from individual developers (\$5 free credit) to the world's largest tech companies (Microsoft, OpenAI). Core enterprise customers appear to be large-scale AI deployers needing 10,000+ GPUs and multi-megawatt capacity.

Implied Marketing Strategy

How Nscale intends to win:

Nscale's implied strategy is to become the **dominant AI-purpose-built hyperscaler** by

exploiting a structural market gap: the massive demand for AI compute infrastructure that traditional hyperscalers (AWS, Azure, GCP) cannot fully satisfy due to capacity constraints, and that neoclouds cannot reliably serve due to fragile economics and lack of vertical integration.

Competitive moat being built:

1. **Vertical integration from energy to compute:** Nscale is assembling a moat by owning the entire value chain — power assets, data centers, GPU clusters, networking, orchestration software, and AI services. The acquisition of AIPCorp (8GW + microgrid in West Virginia) and establishment of Nscale Energy & Power represent a strategy to control the scarcest resource in AI: power. CEO explicitly stated this is the "ONLY way" to build a resilient, profitable AI cloud. This creates structural cost advantages ("at least 10% lower cost of production") that compound at scale.
2. **Capacity as a moat:** Nscale is aggressively securing physical capacity — 1.3GW pipeline, 1.35GW Microsoft LOI, 8GW + West Virginia potential, Stargate Norway (230MW + 290MW expansion). In a supply-constrained market, the ability to guarantee capacity is a decisive competitive advantage.
3. **Anchor customer relationships:** The Microsoft LOI (1.35GW), OpenAI (Stargate Norway), and NVIDIA (first-to-deploy Vera Rubin in Europe) create co-dependency relationships that are difficult for competitors to replicate. These marquee partnerships serve as both revenue anchors and credibility signals.
4. **Sovereign positioning:** Nscale is riding the European digital sovereignty trend aggressively — "European-domiciled," UK/European law governed, sovereign cloud messaging, DORA/EU AI Act content, Trusted Tech Alliance membership, and "only full stack sovereign AI infrastructure provider in the UK." This positions Nscale as the default choice for European governments and regulated enterprises.

Market dynamics being exploited:

- **GPU and power scarcity:** The core constraint in AI infrastructure is not demand but the ability to deploy capacity. Nscale's early moves to secure low-cost renewable power sites and preferential GPU allocations (NVIDIA relationship) exploit this bottleneck.
- **Neocloud fragility:** CEO explicitly called out neoclouds' structural weaknesses (high leverage, no owned DCs, third-party software, spot compute). Nscale positions as the "adult in the room" — a financially stable, vertically integrated alternative.
- **Rapid GPU depreciation cycles:** "Every twelve months, a new generation of GPUs leapfrogs the last... the previous generation can lose 20 to 30 percent of its value within a year." Nscale's modular, chip-agnostic infrastructure is designed to ride this cycle rather than be damaged by it.
- **Regulatory momentum:** GDPR, DORA, EU AI Act, UK national compute strategy — Nscale converts regulatory pressure into a selling point for its sovereign, compliant infrastructure.

Trends being ridden:

- **AI infrastructure as the "fourth industrial revolution"** — CEO's explicit framing.
- **Open-source AI models:** Nscale supports Meta Llama, DeepSeek, Qwen, Mistral, etc., and integrates with Hugging Face. This aligns with the trend away from proprietary model lock-in.
- **Inference-at-the-edge / AI Grid:** The Telco strategy and AI Grid playbook position

Nscale for the next wave — distributed inference beyond centralized cloud, with telcos as channel partners.

- **Sustainability as a competitive requirement:** 100% renewable energy, PUE ~1.1, waste heat reuse, no diesel generators — converts ESG pressure into differentiation.

Implied GTM strategy:

- **Top-down anchor deals** (Microsoft, OpenAI, governments) create credibility and revenue certainty.
- **Bottom-up developer adoption** via serverless inference (\$5 free credits, pay-as-you-go, OpenAI API compatibility) builds a long-tail user base.
- **Channel through telcos** (Singtel partnership, AI Grid playbook, MWC presence) opens enterprise and government markets globally without direct salesforce in every country.
- **Thought leadership and third-party validation** (ABI Research #1 ranking, Top500 list, BBC Storyworks feature, Europas 100, Trusted Tech Alliance) builds trust among technical buyers and C-suite.
- **Board-level signaling:** Addition of Sheryl Sandberg, Susan Decker, and Nick Clegg to the board — names chosen for operational scaling expertise, governance credibility, and policy influence, signaling Nscale's ambition to become a Fortune-level infrastructure company.

Pricing strategy: Transparent, pay-as-you-go for serverless (per-token); reserved capacity for enterprise training clusters; token-based fine-tuning pricing. The explicit "80% lower cost than hyperscalers" claim and "10% lower cost of production than competitors" suggest a penetration pricing strategy enabled by structural cost advantages.

Operational Metrics

Source: *nscale.md*

Marketing – Factual & Quantitative: <https://nscale.com>

Extract concrete, verifiable facts from a company's website: corporate details, infrastructure specs, customer names, product listings, and quantitative claims.

Based on 93 stored pages analysed in 1 batch(es) with 1 AI calls (namespace: webscrape, model: claude-opus-4-6, profile: operational_metrics).

Number Of Data Centers, Their Locations, Gpus, And Power Capacity

Nscale-operated data centers:

- **Glomfjord, Norway** – Located in the Arctic Circle. 100% renewable hydropower. 30MW operational capacity, expandable to 60MW. Powered by a 120MW hydro power plant. Uses adiabatic cooling and direct-to-chip liquid cooling. AMD MI250X GPU

cluster (used for Top500 submission). Town of ~1,000 people.

- **Narvik, Norway** – Location of "Stargate Norway" (partnership with Aker ASA and OpenAI). 230MW initial capacity with plans to expand by an additional 290MW. Powered entirely by renewable hydropower. Plans to deliver 100,000 NVIDIA GPUs. Closed-loop liquid cooling.
- **Loughton, United Kingdom** – London-proximate AI campus. High-density GPU deployments. Ability to scale power allocation up to 90MW. Direct-liquid cooling, waste heat and water reuse.
- **Ward County, Texas, USA** – ~240MW AI data center in West Texas, developed in partnership with Ionic Digital. Plans to expand to 1.2GW. Closed-loop direct-liquid cooling and rear-door heat-exchangers.

Partner-run data centers:

- **Sines, Portugal** – Located at Start Campus data center on Portugal's Atlantic coast. Designed for GW-scale AI deployments. Powered entirely by renewable energy with seawater cooling. Low-latency, high-bandwidth hub for European and trans-Atlantic AI traffic.
- **Keflavik, Iceland** – One of Iceland's largest liquid-cooled GPU installations. Set to host more than 4,600 NVIDIA Blackwell Ultra GPUs across Verne's Icelandic campus in 2026. 100% renewable energy from geothermal and hydropower.

Available data centers (additional locations listed):

- Stavanger, Norway
- Oslo, Norway
- Blönduós, Iceland
- Slough, United Kingdom
- North Carolina, USA

Monarch Compute Campus (acquired via AIPCorp):

- Mason County, West Virginia, USA – Up to 2,250 acres. First U.S. state-certified AI microgrid. Power runway scalable to over 8 gigawatts. Initial 2GW expected online by first half of 2028, expansion to ~8GW planned for 2031. Natural gas powered (Caterpillar G3500 series). LOI with Microsoft for 1.35GW of AI compute using NVIDIA Vera Rubin NVL72 GPUs.

GPU types mentioned:

- AMD Instinct MI250X (operational at Glomfjord, used for Top500)
- AMD Instinct MI300X (standing up at Glomfjord; also available for customers)
- AMD Instinct MI210 (referenced in CPU inferencing blog benchmarks)
- NVIDIA H100, H200, GB200 NVL72 (offered as GPU node options)
- NVIDIA Blackwell Ultra GPUs (4,600+ planned for Keflavik in 2026)
- NVIDIA Vera Rubin NVL72 (planned deployment in 2027, 100,000+ GPUs to Europe)
- NVIDIA A100, V100 (mentioned as available)

Other infrastructure stats:

- Alex Sharp (President, Data Center Group) has led design/construction of 54 data centers.
- Nscale expanded pipeline from 300MW to 1.3GW of greenfield data center sites as of Series A (Dec 2024).
- Expects 250MW commissioned by Q4 2026, with potential to expand to 1GW+ by

2029.

- Nscale's current capacity described as "over 1GW" (in Monarch press release context).
 - UK commitment: 10,000 NVIDIA Blackwell GPUs by 2026.
-

Active And Contracted Megawatts, Terawatts, And Other Measures Of Power

- **Glomfjord:** 30MW operational, expandable to 60MW. Powered by 120MW hydropower plant. 100% renewable.
 - **Narvik (Stargate Norway):** 230MW initial capacity, plans to expand by additional 290MW (total potential ~520MW). 100% renewable hydropower.
 - **Loughton, UK:** Up to 90MW.
 - **Ward County, Texas:** ~240MW, plans to expand to 1.2GW.
 - **Sines, Portugal:** Designed for GW-scale AI deployments. 100% renewable energy.
 - **Monarch Compute Campus, West Virginia:** 2GW initial power expected online by H1 2028, scalable to over 8GW by 2031. LOI with Microsoft for 1.35GW of compute. Powered by Caterpillar G3500 natural gas generator sets. Microgrid designed to enable future tie-in to the grid. On-site powered microgrid (not connected to local grid).
 - **Pipeline:** 1.3GW of greenfield data center sites across Europe and North America (as of December 2024 Series A announcement).
 - **250MW expected commissioned by Q4 2026**, potential to expand to 1GW+ by 2029.
 - **\$1.4 billion GPU-backed Delayed Draw Term Loan** to finance multiple cluster deployments across Norway, Portugal, Iceland, and UK.
 - Glomfjord PUE target: approximately 1.1.
 - 85% of waste heat at Glomfjord targeted to be repurposed (e.g., fish farms, swimming pools).
 - No diesel generators at Glomfjord (zero direct diesel emissions).
-

Installed, Active And Future Number Of Gpus

- **Keflavik, Iceland:** More than 4,600 NVIDIA Blackwell Ultra GPUs planned for deployment in 2026.
- **Narvik (Stargate Norway):** Plans to deliver 100,000 NVIDIA GPUs.
- **Europe (Vera Rubin deployment):** 100,000+ NVIDIA Vera Rubin GPUs to Europe in 2027.
- **UK:** Commitment to deploy 10,000 NVIDIA Blackwell GPUs by 2026.
- **Nidhi Chappell** (President AI Infrastructure): Former Head of AI Infrastructure at Microsoft where she "managed 2m GPUs."
- Blog references indicate scaling "toward several hundred thousand GPUs over the next three years" (from December 2025).
- ABI Research scored Nscale a "perfect 10 for maximum distributed cluster scale" and "running tens of thousands of GPUs globally for enterprise customers."
- GPU availability scored 9/10 by ABI Research: "most on-demand jobs typically start within minutes."
- Inference service has grown 148x, serving more than 5,000 users.
- Glomfjord Top500 cluster: AMD MI250X nodes, each with 4x AMD Instinct MI250X OAMs (8x GPUs per node), 2x AMD EPYC 7713 CPUs per node.
- Singtel partnership references NVIDIA H100 Tensor Core GPU capacity in Southeast Asia.

- GPU node offerings: NVIDIA A100, H100, H200, GB200 NVL72, V100.
-

Teraflops And Exaflops

- GPT-1 training used approximately **0.96 petaflop/s-days (pfs-days)**; GPT-3 used **3,630 pfs-days** (cited in training/inference blog).
 - **NVIDIA GB200 NVL72** performance claims: 4x faster LLM training, 25x more energy efficiency than legacy solutions, 30x faster real-time trillion-parameter LLM inference, 18x faster data processing than Intel Xeon 8480+.
 - Nscale's Svartisen Cluster made the **Top500 list** (24th official Top500 list, November 2024). HPL benchmark run in FP64 mode; also ran mixed precision (Linpack MXP). Specific TFLOPS/PFLOPS ranking number not stated on the website.
 - **GEMM tuning benchmarks:** Up to 7.2x improvement in throughput and latency (LLaMA-2-70B model). ~33% latency reduction with Gradlib GEMM tuning.
-

Other Capacity And Performance

Funding & Valuation:

- **Series C:** \$2 billion (March 2026) – largest in European history. Valuation: \$14.6 billion. Led by Aker ASA and 8090 Industries, with support from Astra Capital Management, Citadel, Dell, Jane Street, Lenovo, Linden Advisors, Nokia, NVIDIA, Point72.
- **Pre-Series C SAFE:** \$433 million.
- **Series B:** \$1.1 billion – largest Series B in European history.
- **Series A:** \$155 million (December 2024). Led by Sandton Capital Partners, with Kestrel 0x1, Blue Sky Capital Managers, Florence Capital.
- **GPU DDTL:** \$1.4 billion delayed draw term loan backed by GPUs. Led by PIMCO, Blue Owl, LuminArx Capital Management.

Board members (added Series C): Sheryl Sandberg, Susan Decker, Nick Clegg. Existing: Josh Payne, Rael Nurick, Jacob Leschly, Øyvind Eriksen.

Performance claims:

- 80% lower cost compared to hyperscalers (stated multiple times).
- 30% faster time to insights with AI-optimised stack.
- Up to 40% improvement in resource utilization/efficiency.
- ABI Research: "mature optimization stack... can deliver 20 to 40 percent better throughput without changing hardware."
- Interconnect Bandwidth and Topology scored 9/10 by ABI Research (InfiniBand, high-bandwidth Ethernet with RDMA, NVLink/NVSwitch for NVL72 systems, Nokia partnership).
- ISV Partner and Co-Development Innovation scored 9/10 by ABI Research.
- Ranked **#1 overall neocloud** by ABI Research (Overall Leader, Top Innovator, Top Implementer) across 14 companies assessed.

Partners & customers named:

- Microsoft, NVIDIA, OpenAI, Aker ASA, Caterpillar, Dell, Nokia, VAST (storage), Verne (Iceland), Ionic Digital (Texas), Singtel, Lightning AI, Hugging Face, Armada, Lenovo,

AMD, Broadcom, Fidelis New Energy, 8090 Industries.

- Trusted Tech Alliance members (co-founded): Anthropic, AWS, Cassava Technologies, Cohere, Ericsson, Google Cloud, Hanwha, Jio Platforms, Microsoft, Nokia, Nscale, NTT, Rapidus, Saab, SAP.

Rack density evolution: Blog mentions power demands climbing from 7-10 kW/rack toward 150 kW/rack, with predictions for 1 MW+ racks.

Networking: RDMA/InfiniBand/NVLink fabrics; Ethernet fabric with RoCE at Glomfjord (Broadcom); joint innovation with Nokia across switching and optical layers.

Staffing commitment: Plan to hire and train 100 AI specialists in the UK over next 12 months (as of June 2025 announcement).

Acquisition of Future-tech: European data center engineering consultancy (2025).

Acquisition of AIPCorp: Creates Nscale Energy & Power division headquartered in Houston, Texas.

Cost of production: "At least 10% lower than competitors" due to vertical integration.

Leverage: Nscale targets 60-70% leverage vs. industry average of 80-90%.

Corporate entity: NScale Global Holdings Limited (registered name per footer).
